

Software

Feature implementation and computation

Andrea Vedaldi, University of Oxford

<https://sites.google.com/site/eccv12features/>

goo.gl/jKC2J

VLFeat

MATLAB ease of use

Covariant detectors, MSER, SIFT, etc ...

OpenCV

One-shop experience

The most popular vision library

Others

There is certainly no lack of software!

VLFeat

MATLAB ease of use

Covariant detectors, MSER, SIFT, etc ...

OpenCV

One-shop experience

The most popular vision library

Others

There is certainly no lack of software!

VLFeat

An open and portable library of CV software

[**http://www.vlfeat.org/benchmarks/index.html**](http://www.vlfeat.org/benchmarks/index.html)

Quick Start with MATLAB

```
>> untar http://www.vlfeat.org/download/vlfeat-0.9.16-bin.tar.gz
>> run vlfeat-0.9.16/toolbox/vl_setup
```



Mac



Linux



Windows

Since 2008
ACM OSS award

SIFT detector and descriptor

1. load an image

```
imPath = fullfile('oxford.jpg') ;
im = imread(imPath) ;
```

2. convert it to single precision gray scale

```
imgs = im2single(rgb2gray(im)) ;
```

3. run SIFT

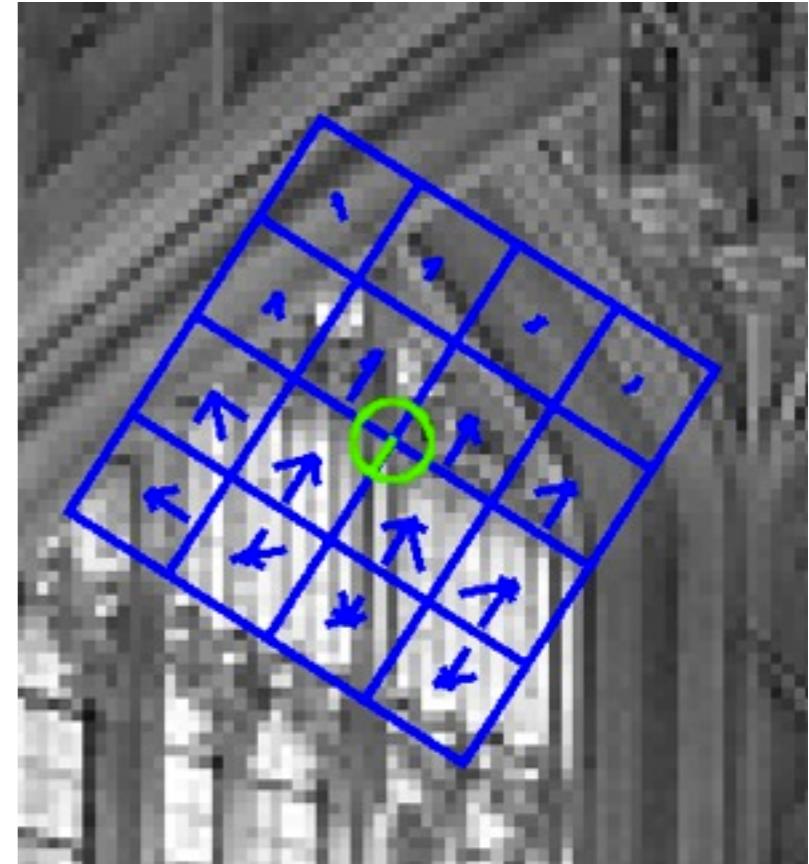
```
[frames, descrs] = vl_sift(imgs) ;
```

4. visualise keypoints

```
imagesc(im) ; colormap gray; hold on ;
vl_plotframe(frames) ;
```

5. visualise descriptors

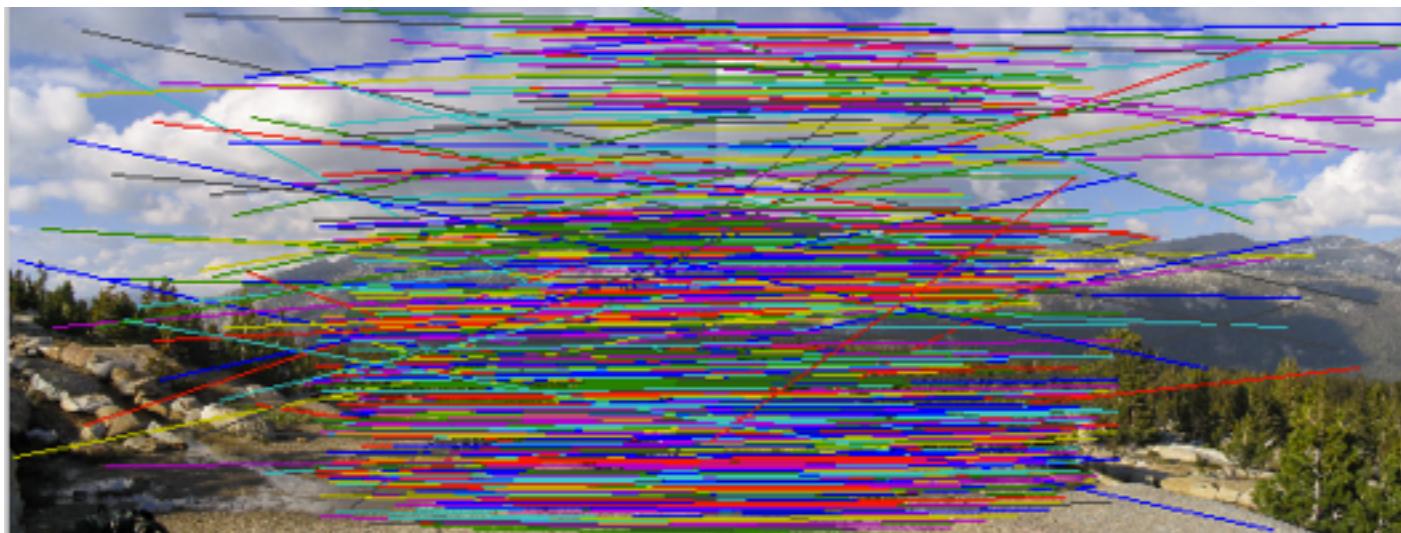
```
vl_plotsiftdescriptor(...  
descrs(:,432), frames(:,432)) ;
```



Example: matching features



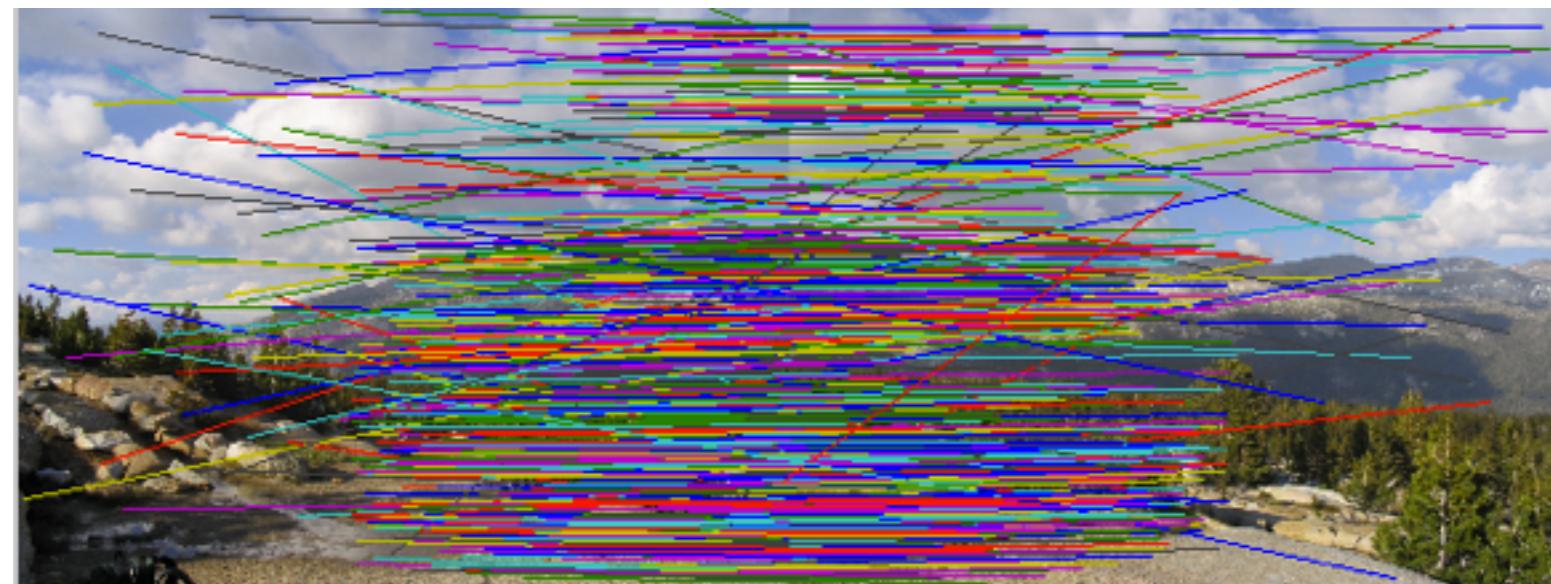
684 tentative matches



- `[f1,d1] = vl_sift(im1) ;`
`[f2,d2] = vl_sift(im2) ;`
`[matches, scores] = vl_ubcmatch(d1,d2) ;`

Example: stitching

684 tentative matches



492 (76%) inliers



mosaic



- Live demonstration

Other covariant detectors

```
frames = vl_covdet(imgs, 'verbose') ;
```



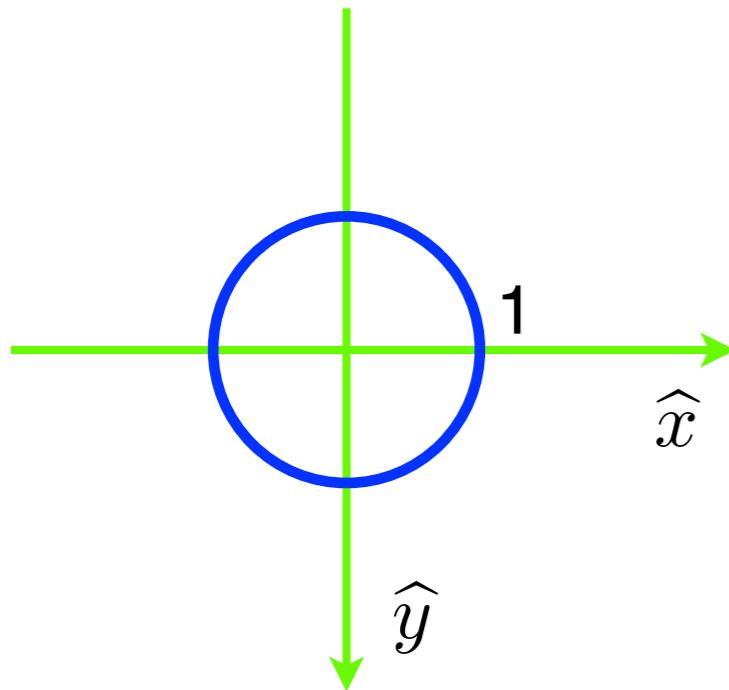
```
imagesc(im) ;
```



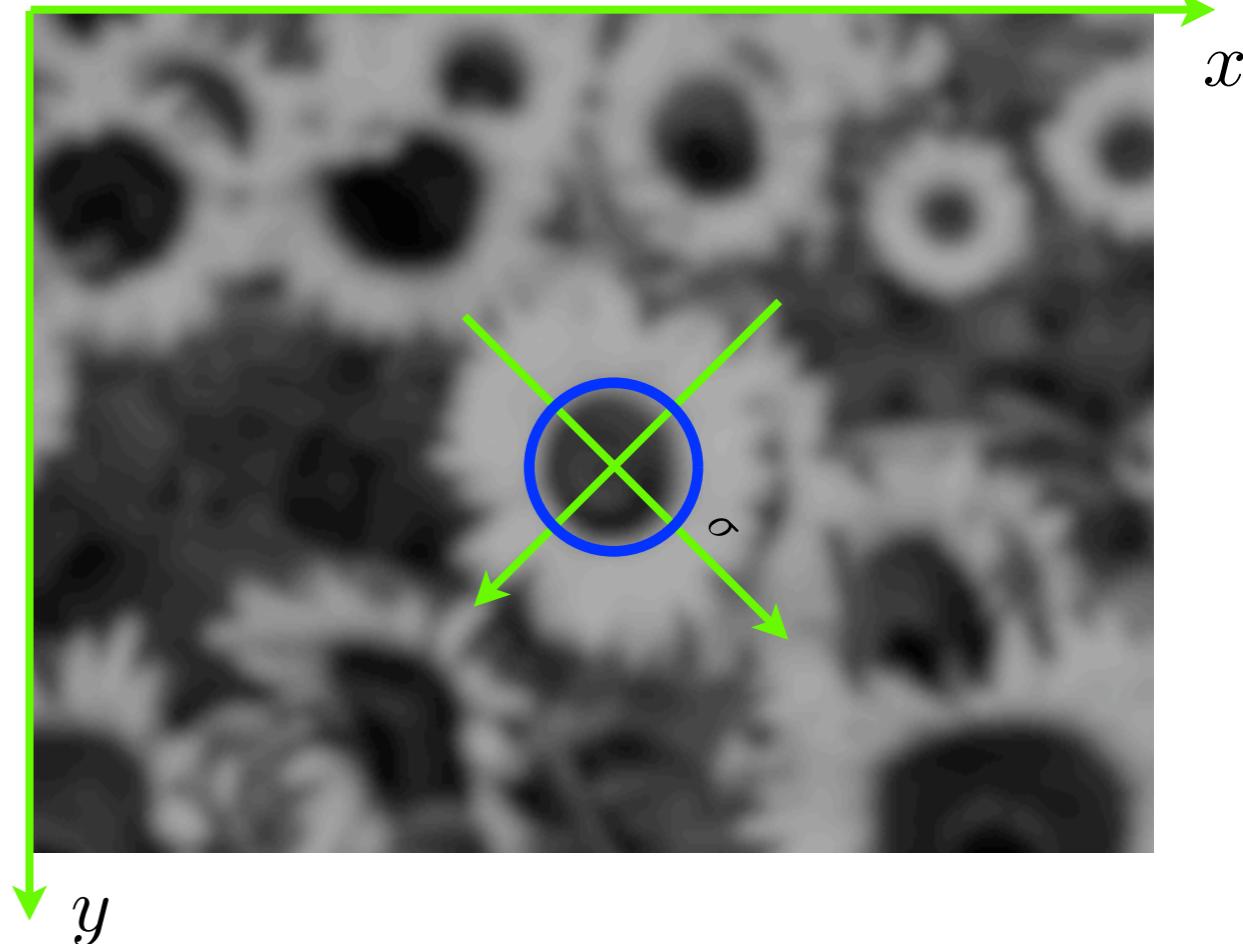
```
vl_plotframes(frames) ;
```

Do not forget to convert the images to
grayscale in single precision. Use
`im2single(im)` and not `single(im)` !

Understanding feature frames



affine tf.

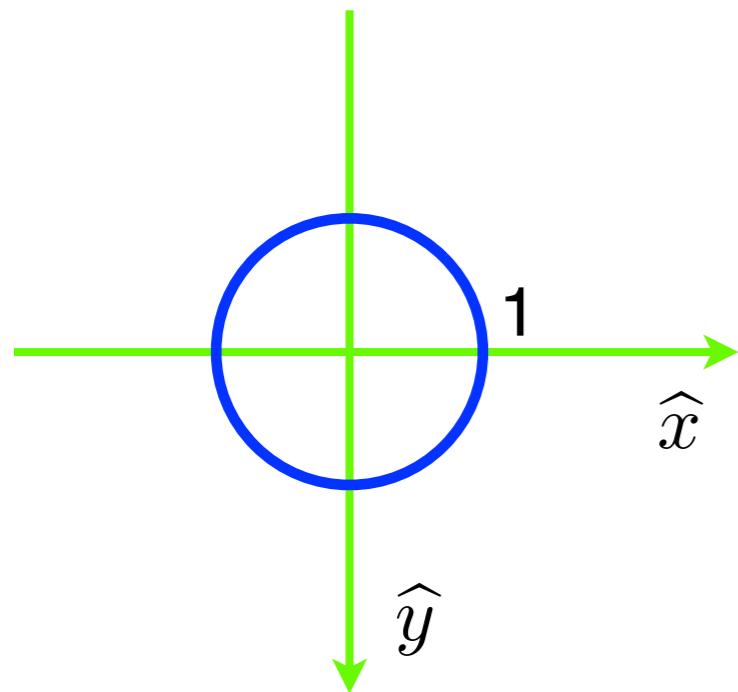


- **frames** has a feature frame for each column containing a stacked affine transformation (A, T)

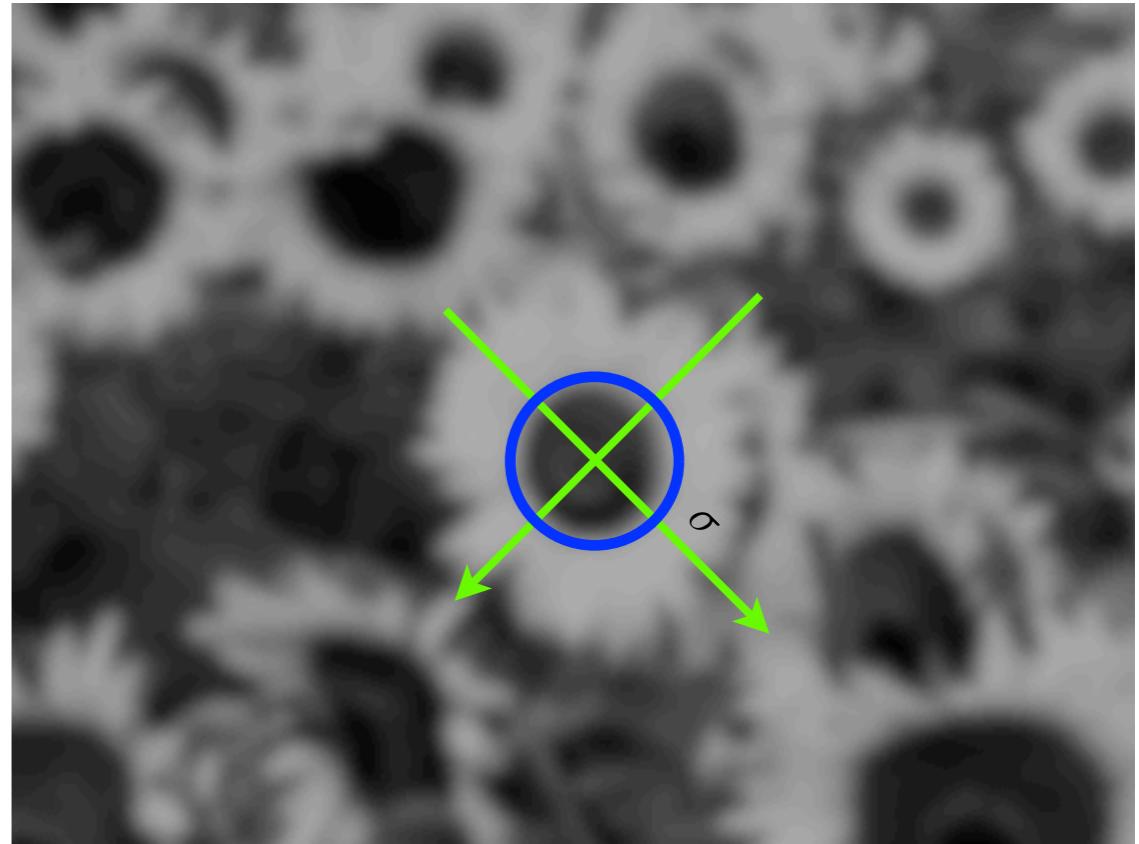
$$\text{frames}(:, i) = \begin{bmatrix} t_1 \\ t_2 \\ a_{11} \\ a_{21} \\ a_{12} \\ a_{22} \end{bmatrix} \quad \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} \hat{x} \\ \hat{y} \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

A frame as an oriented ellipse

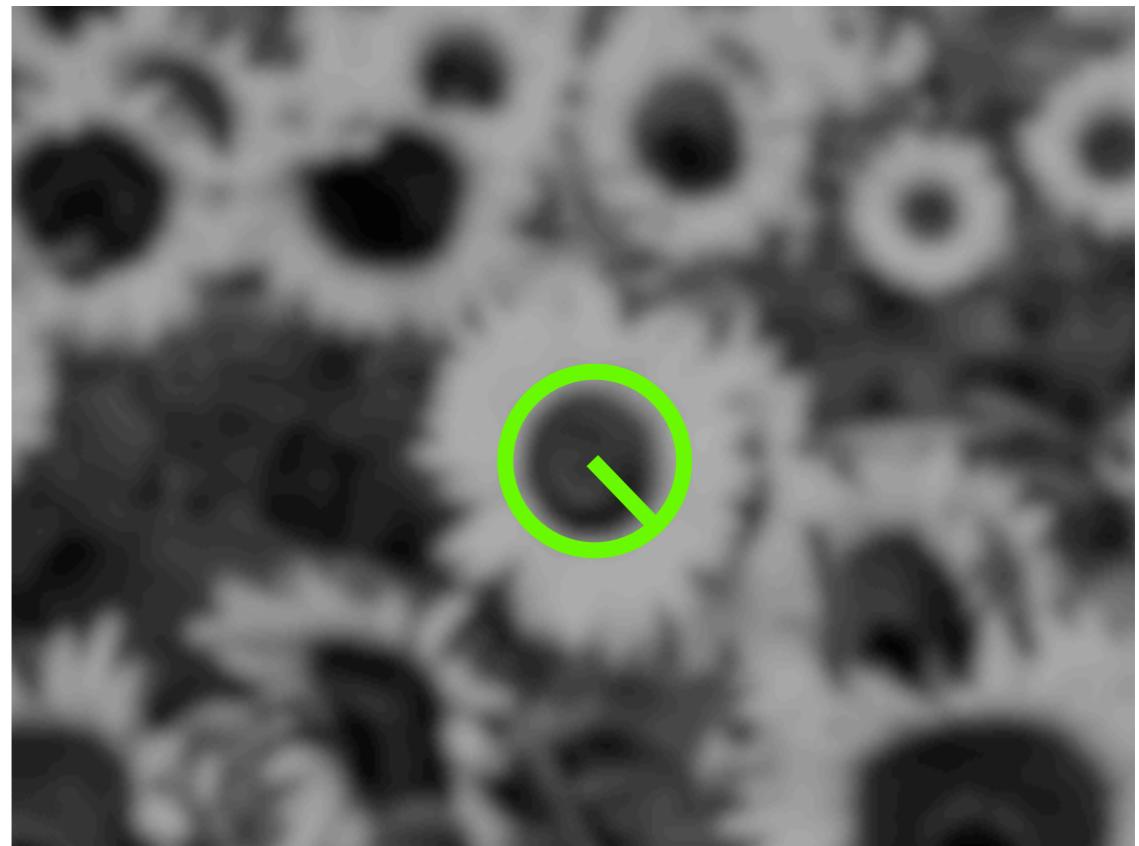
reference = two vectors



uniquely identifies
the affine
transformation



reference = oriented ellipse



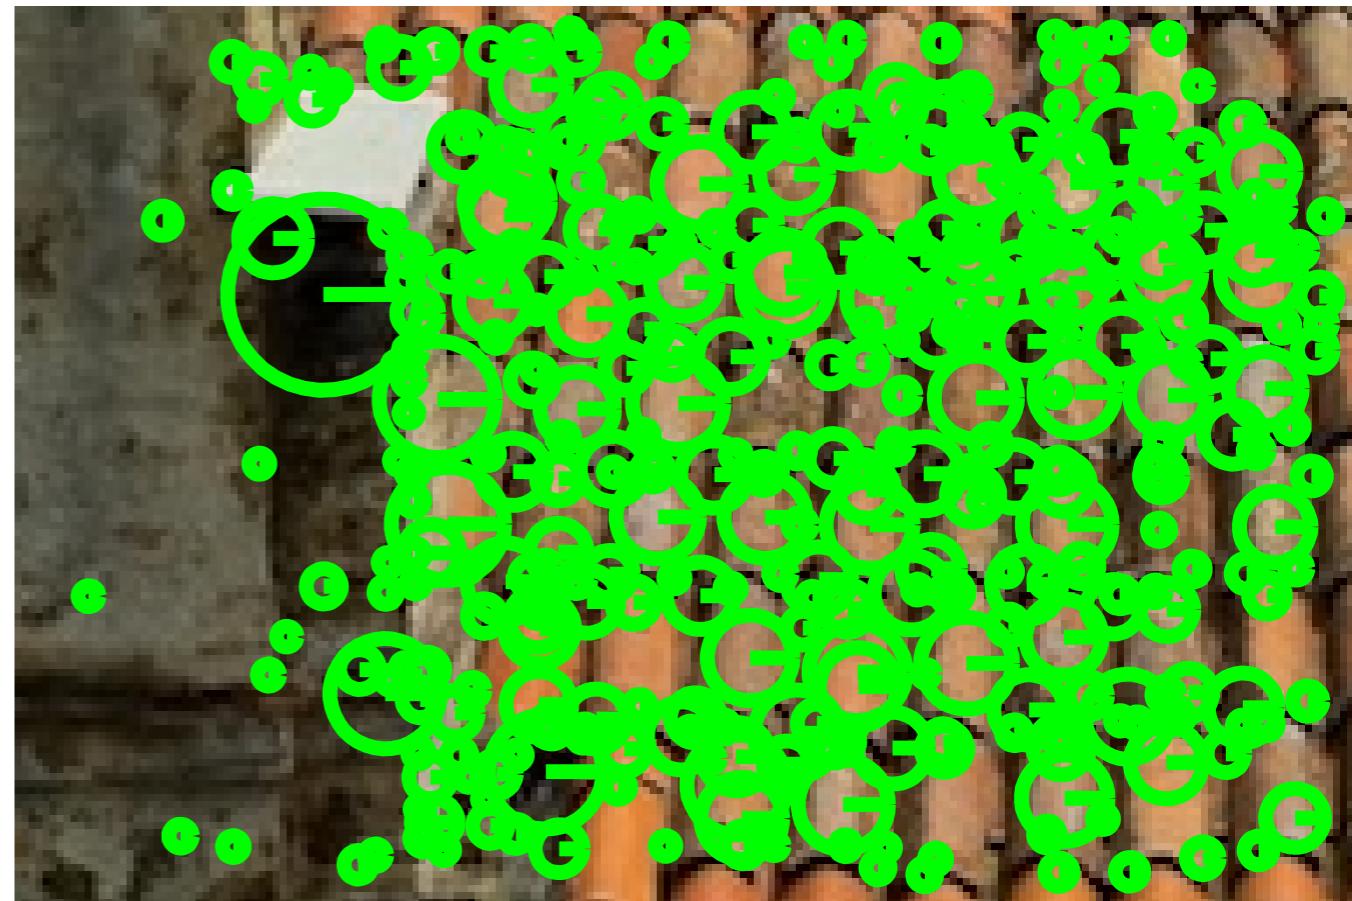
Different cornerness measures

```
frames = vl_covdet(imgs, 'method', 'Hessian') ;
```

DoG



Hessian



- Difference of Gaussian (Laplacian, trace of Hessian, or SIFT)
- Hessian (determinant of)

Different cornerness measures

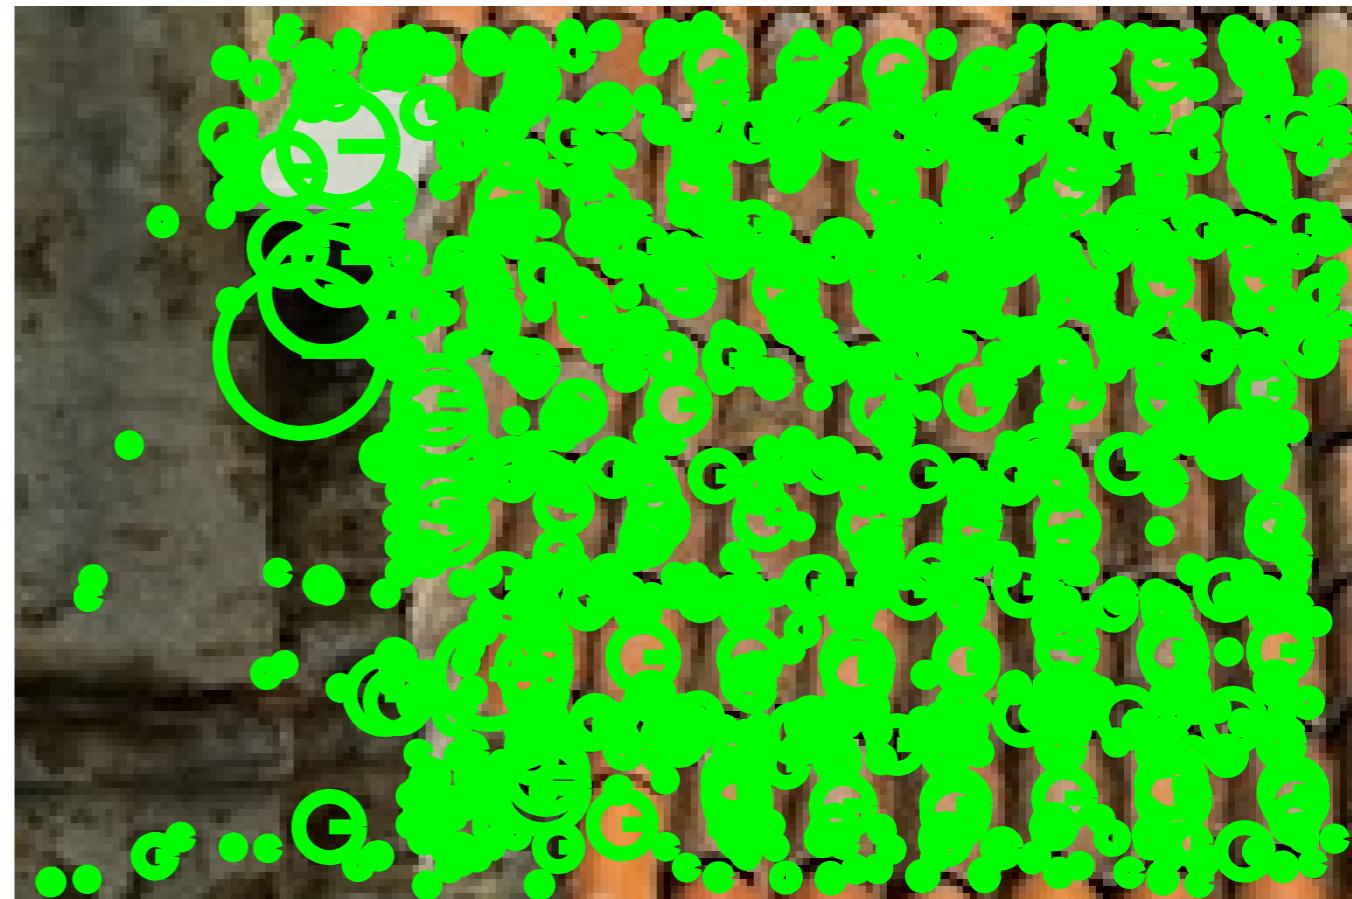
12

```
frames = vl_covdet(imgs, 'method', 'HarrisLaplace') ;
```

HarrisLaplace



HessianLaplace



- Harris Laplace (Harris cornerness, Laplacian for scale selection)
- Hessian Laplace

Different cornerness measures

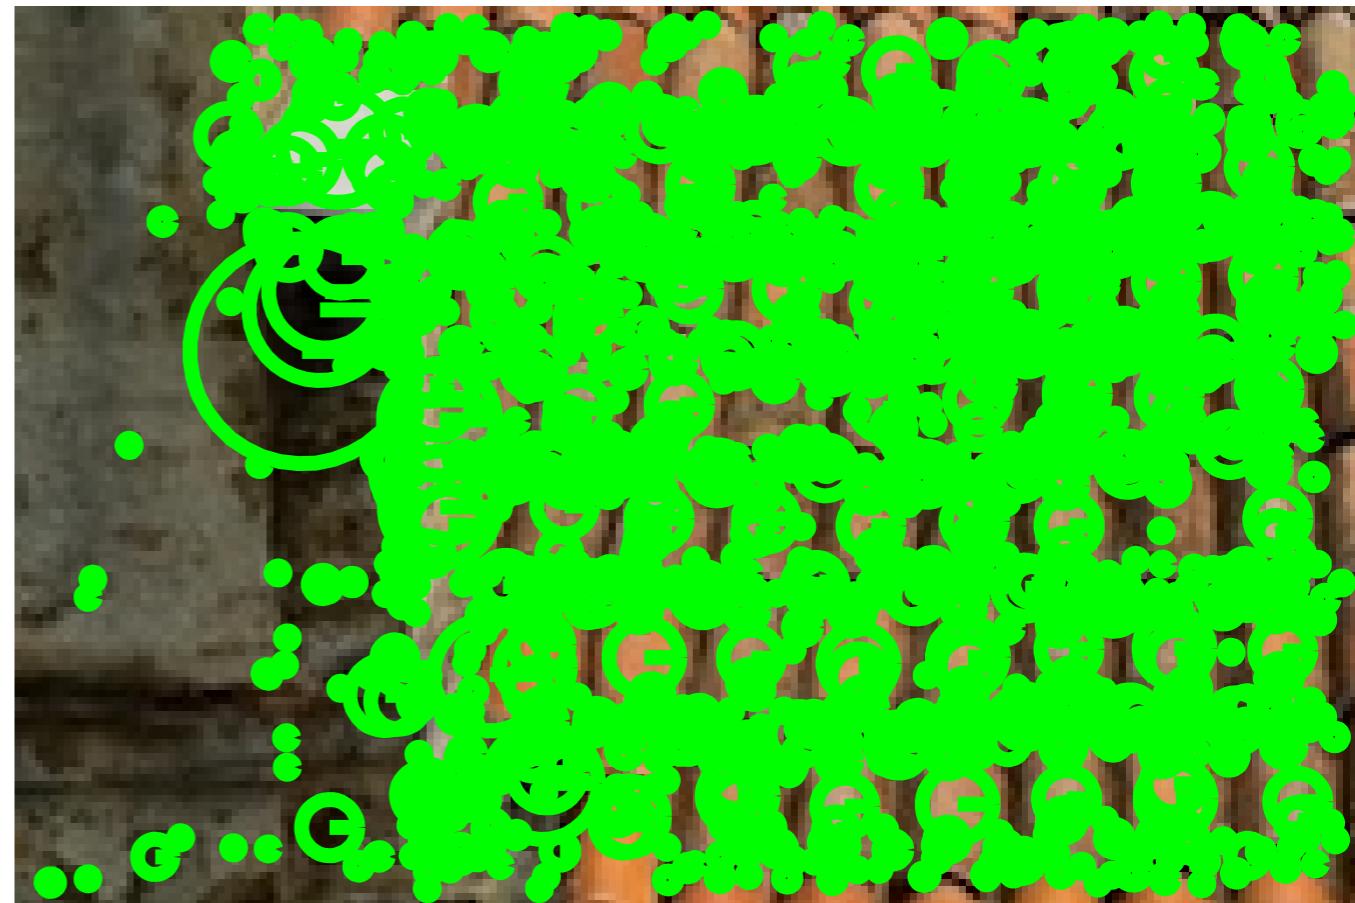
13

```
frames = vl_covdet(imgs, 'method', 'MultiscaleHarris') ;
```

MultiscaleHarris



MultiscaleHessian

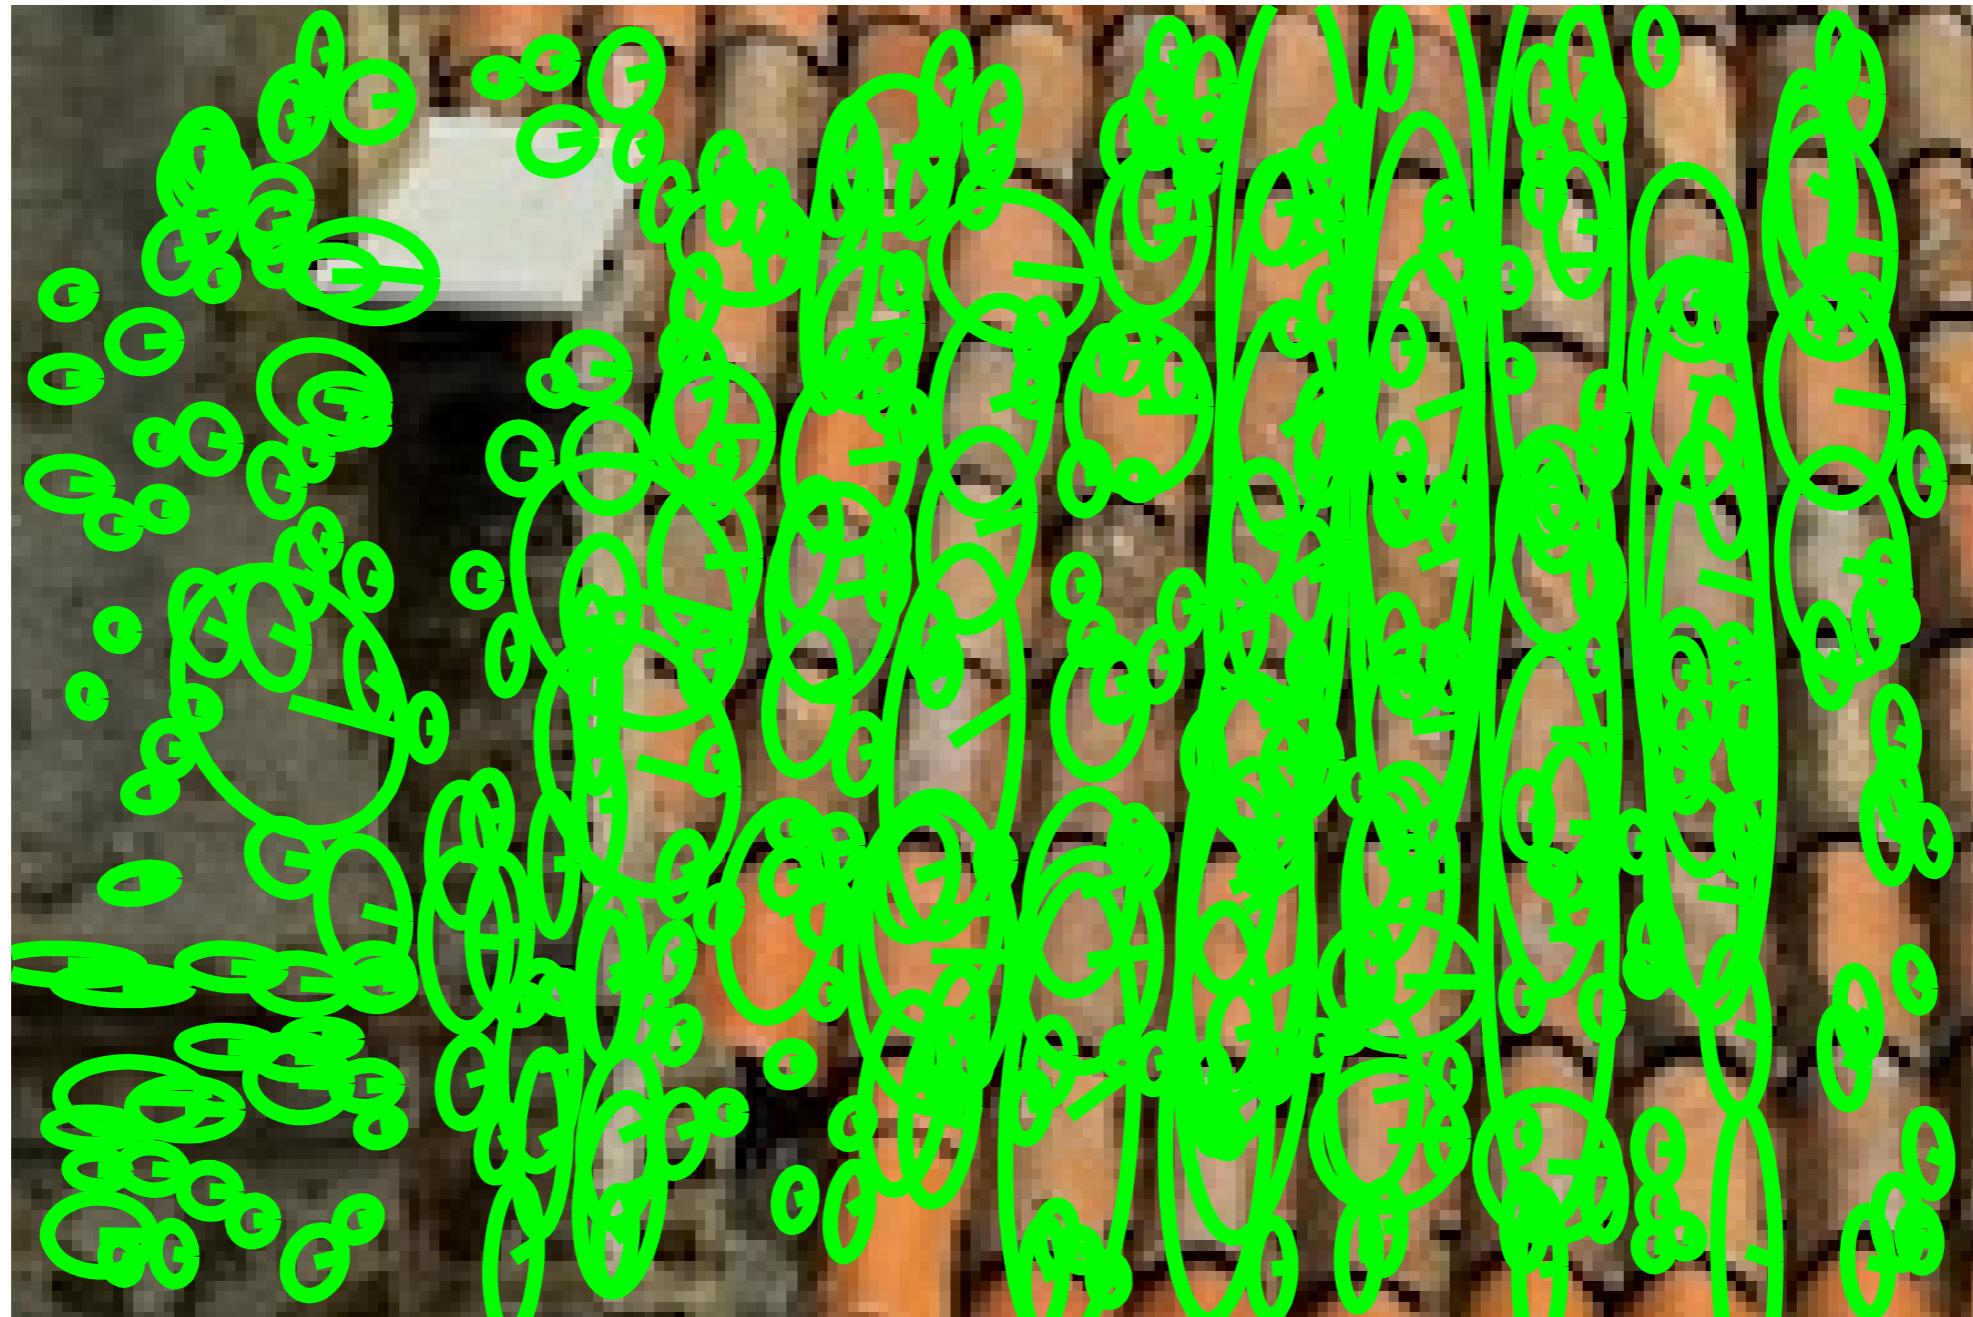


- Multiscale Harris (just a bunch of Harris)
- Multiscale Hessian

Affine adaptation

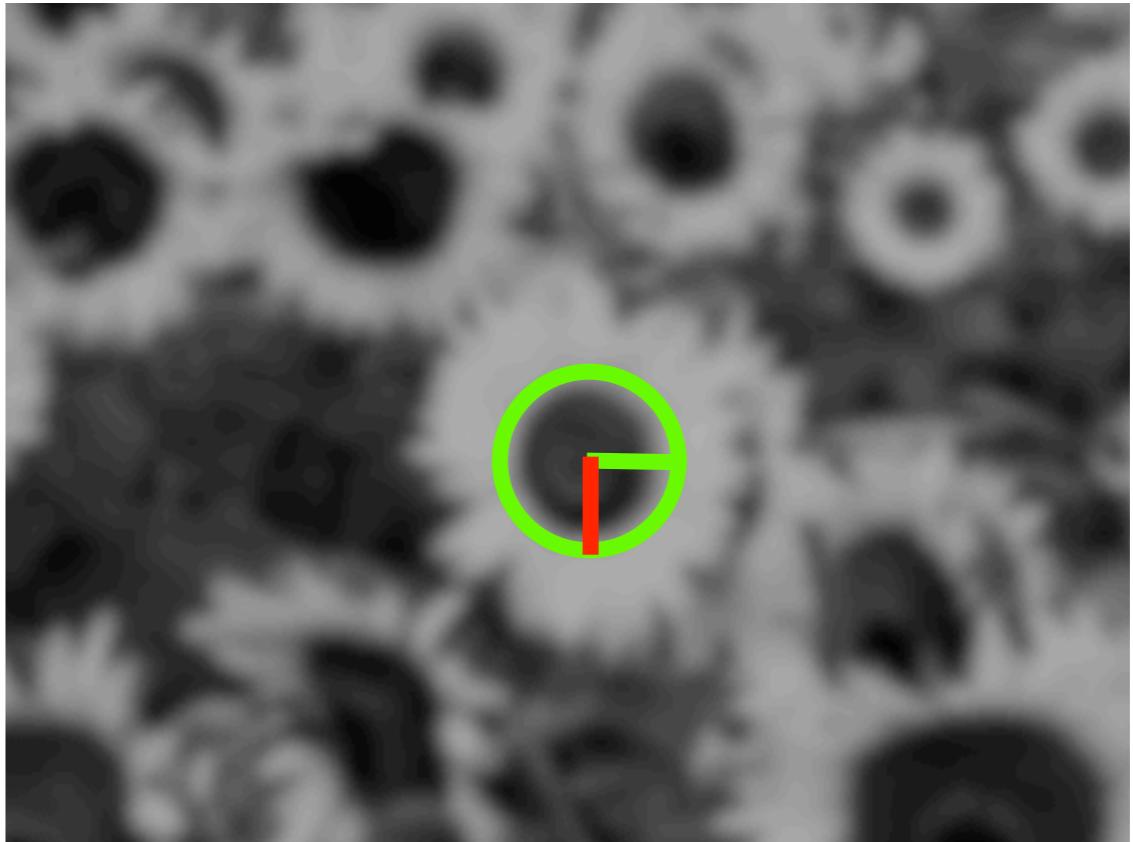
14

```
frames = vl_covdet(imgs, 'EstimateAffineShape', true);
```



Upright features

By default the rotation is selected so that the vertical axis stays vertical.



In formulas

$$\begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} a_{11} & 0 \\ a_{21} & a_{22} \end{bmatrix} \begin{bmatrix} \hat{x} \\ \hat{y} \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$

Remark: for general affine transformations the horizontal direction may appear rotated even for upright features.



Orientation detection

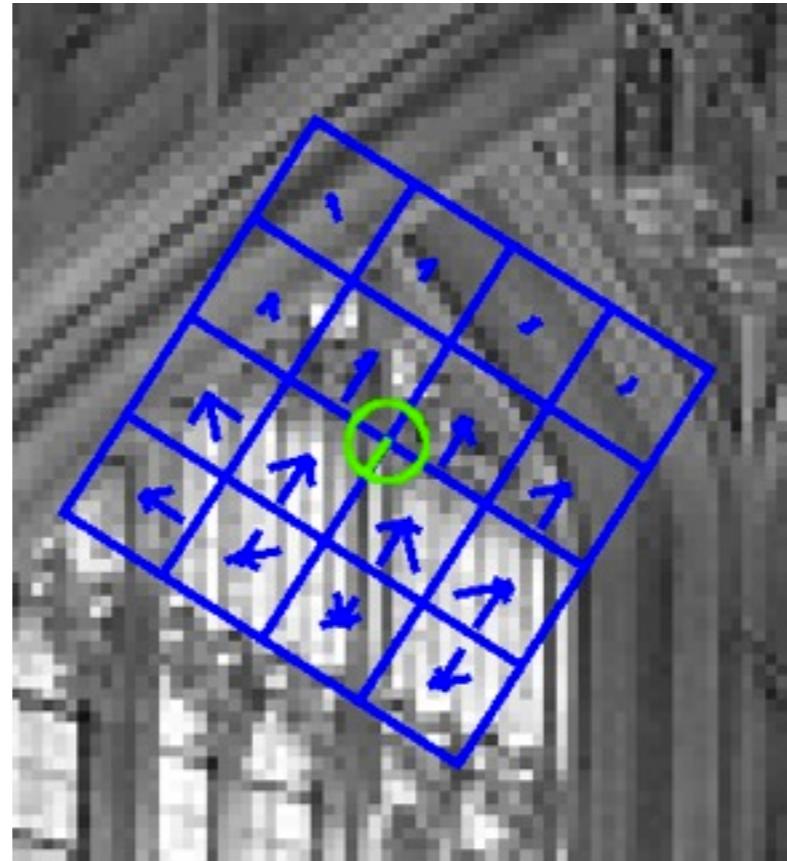
16

```
frames = vl_covdet(imgs, 'EstimateOrientations', true);
```



Computing descriptors

```
[frames, descrs] = vl_covdet(imgs) ;
```

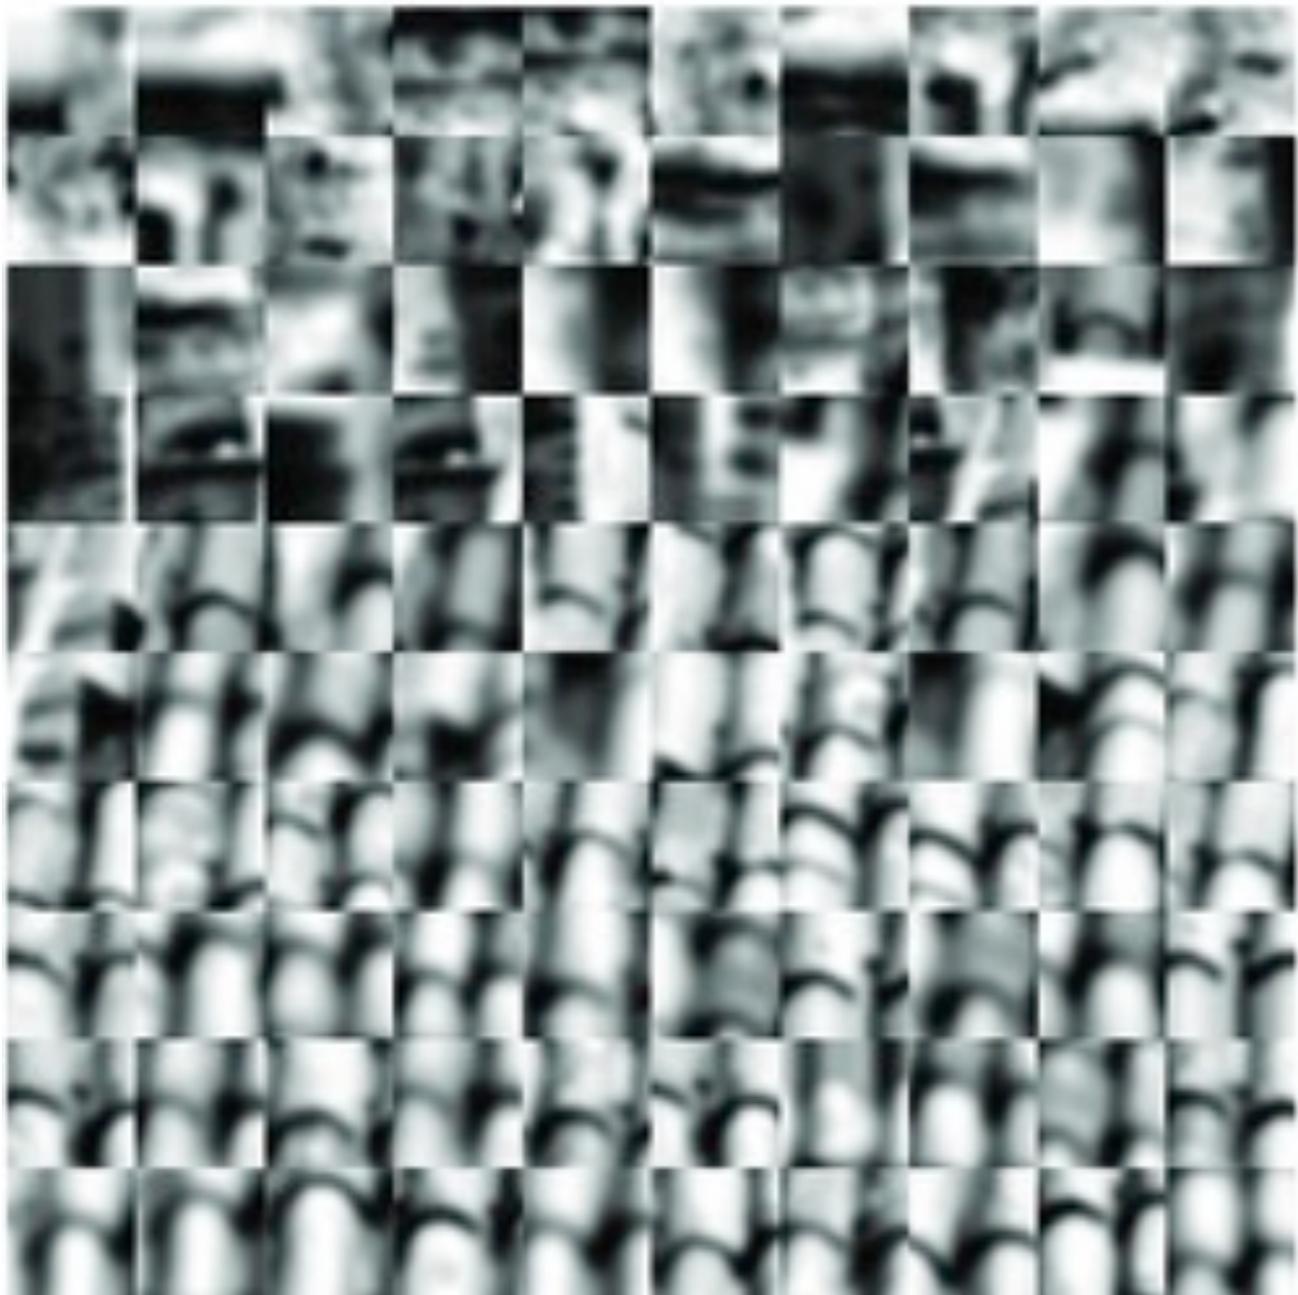


- Each column of **descrs** is a stacked SIFT descriptor
- How is a SIFT descriptor associated to a feature?
- What about other descriptors?

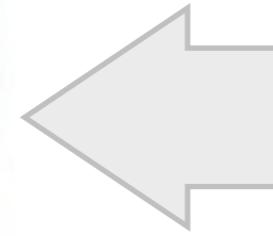
Computing descriptors

18

```
[frames, descrs] = vl_covdet(imgs, 'Descriptor', 'patch');
```

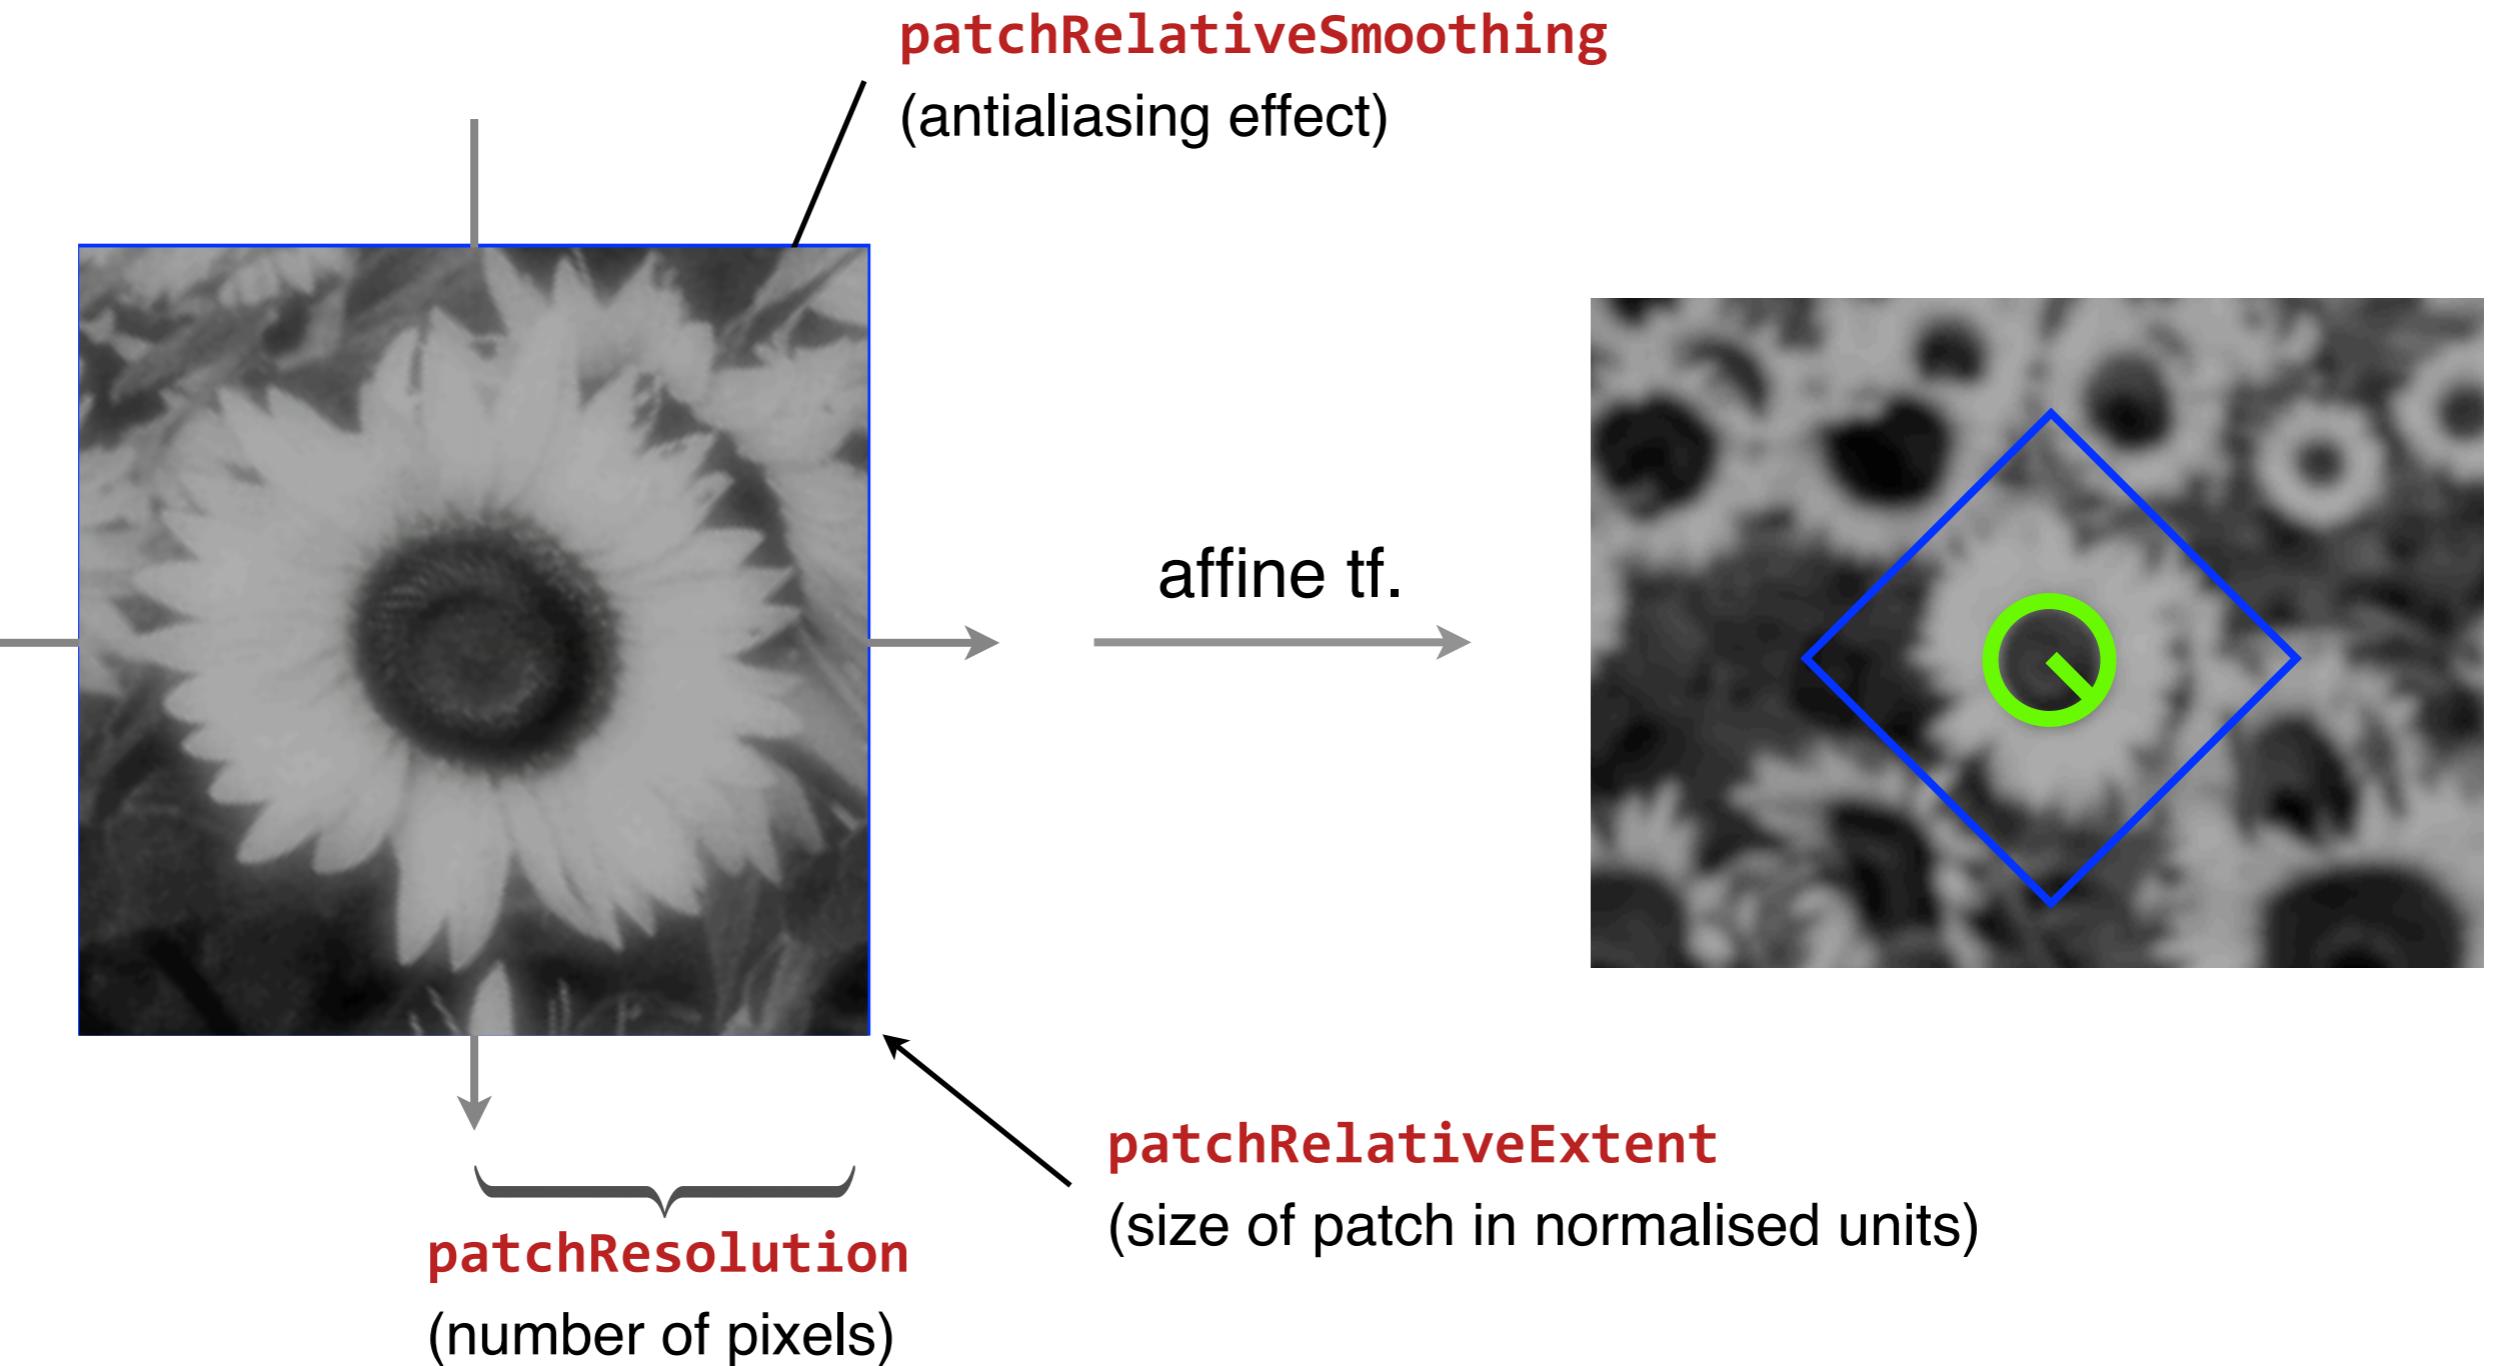


Each column of **descrs** is a stacked normalised image patch.



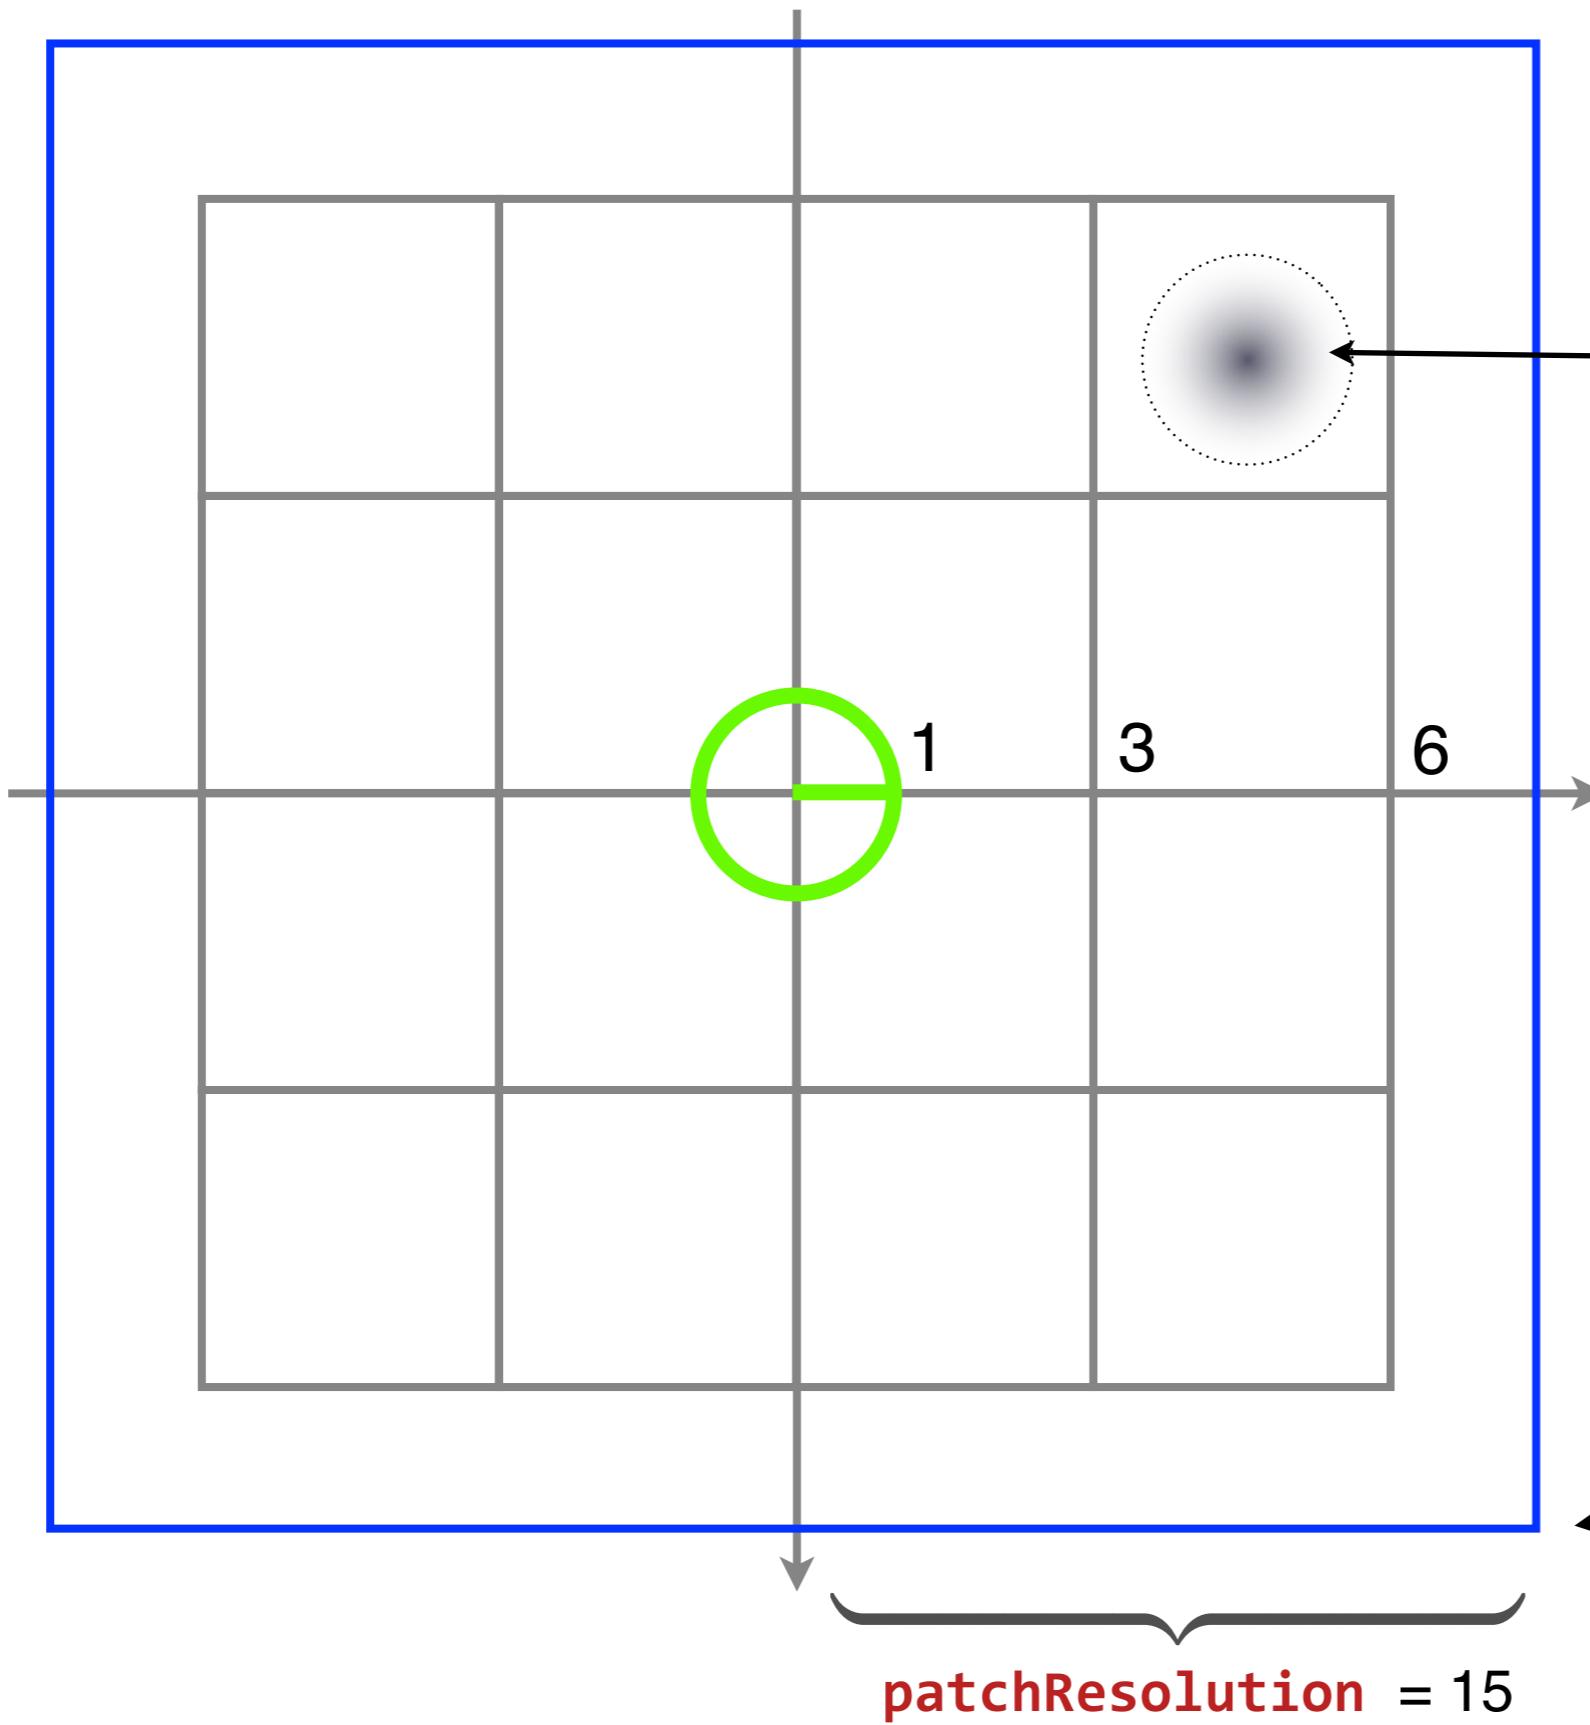
The geometry of an normalised patch

- Normalisation and patch resampling is controlled by three parameters:

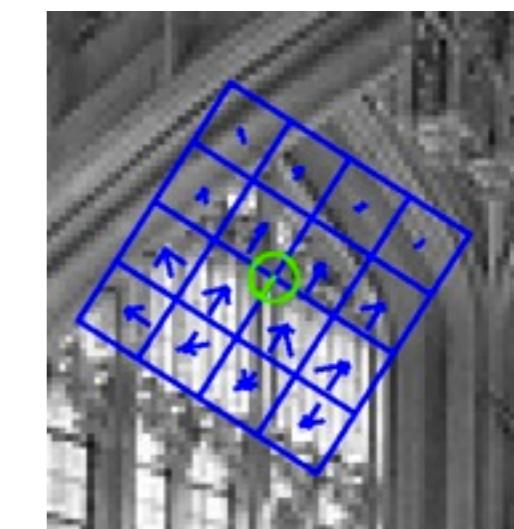


Example: normalised patch for SIFT

20



patchRelativeSmoothing = 1



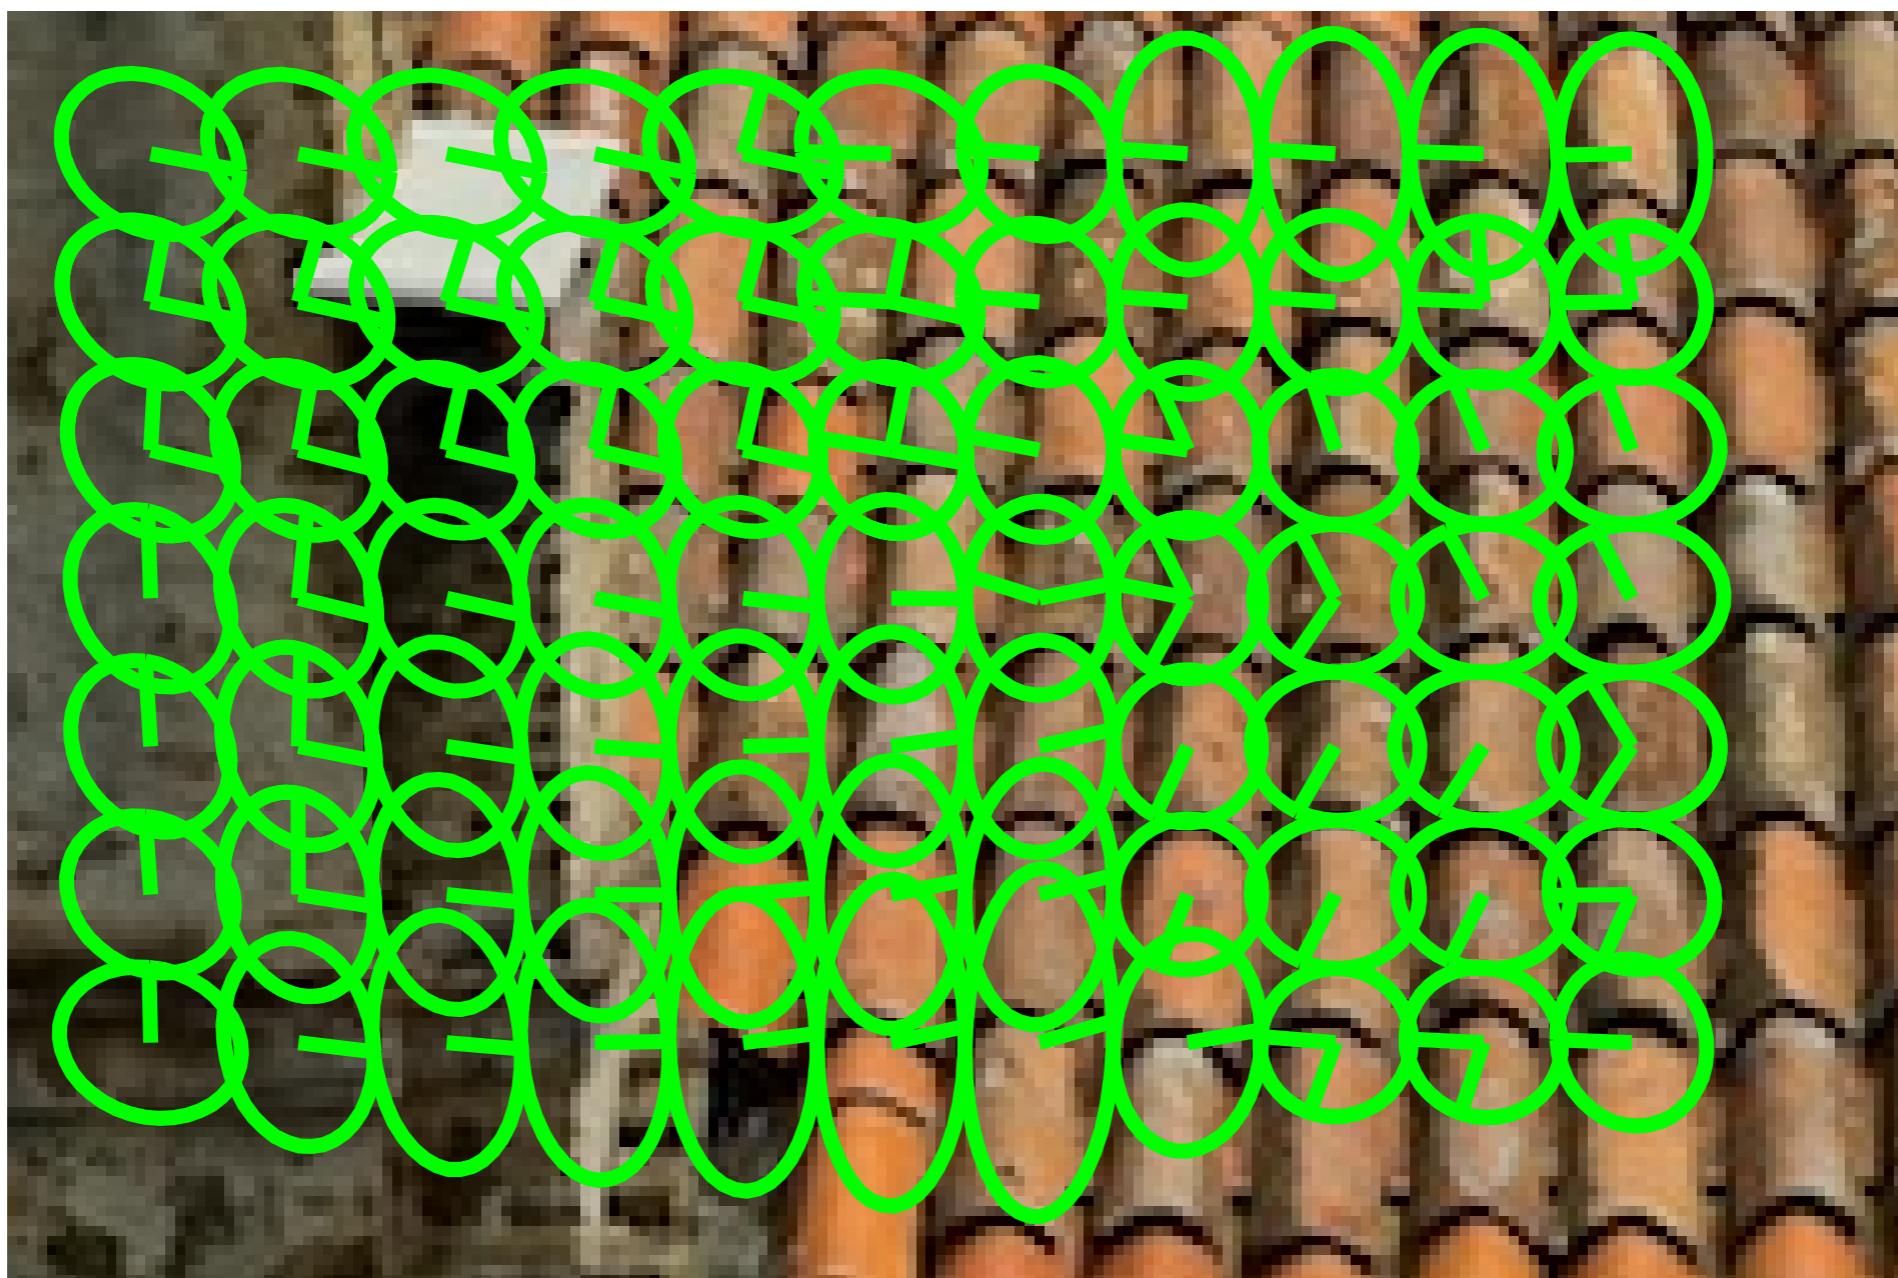
patchRelativeExtent = 7.5

patchResolution = 15

Custom frames

Example: running affine adaptation on a grid

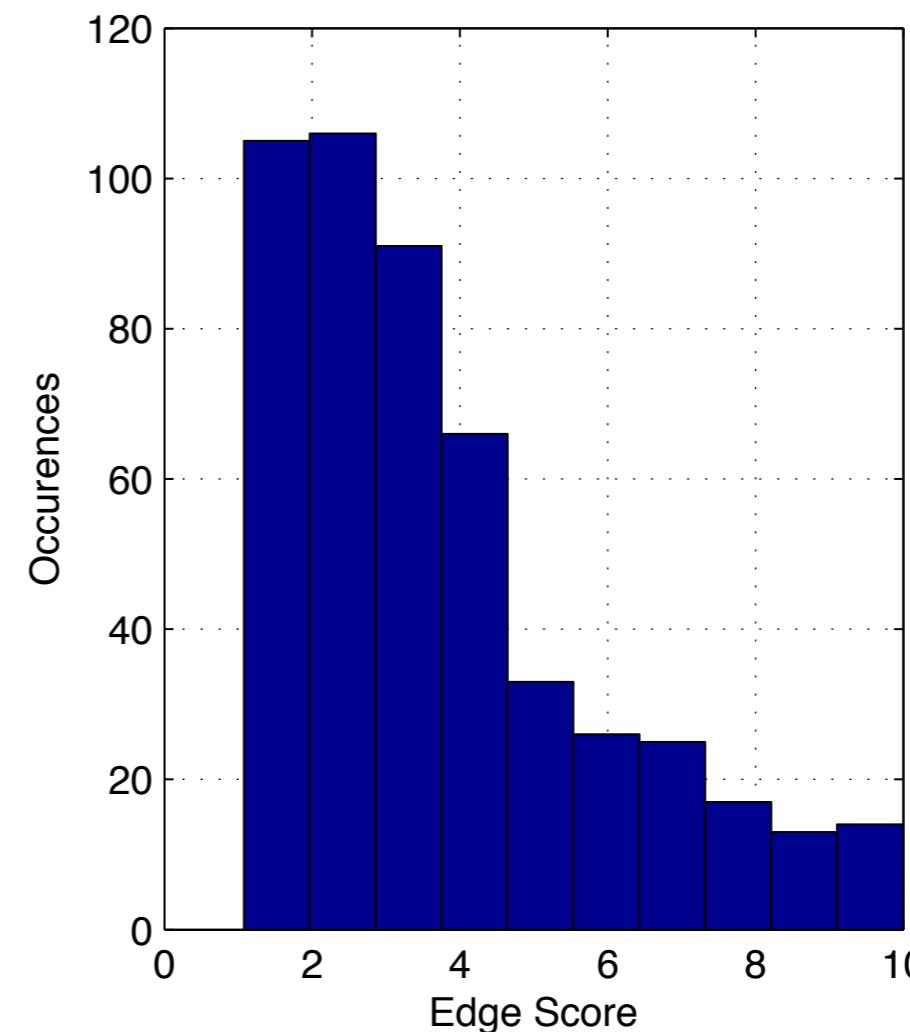
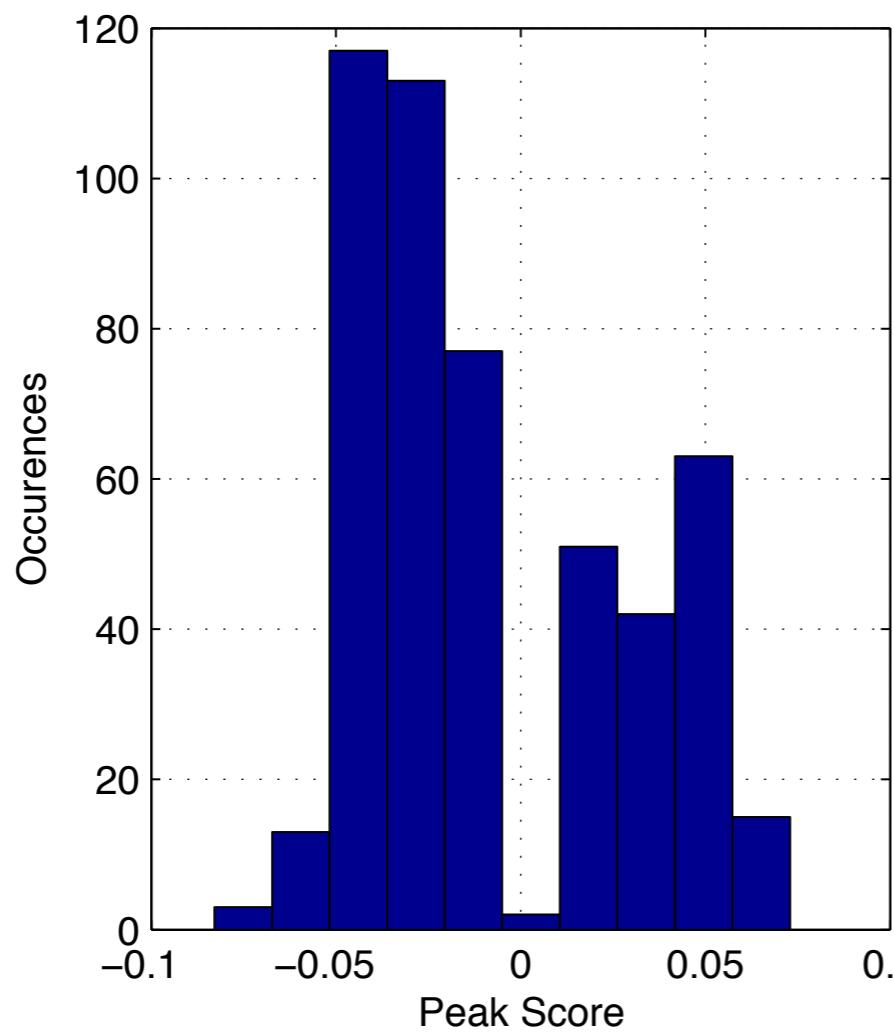
```
frames = vl_covdet(imgs, ...
    'Frames', frames, ...
    'EstimateAffineShape', true, ...
    'EstiamteOrientation', true) ;
```



Obtaining additional information on the features

22

```
[frames, descrs, info] = vl_covdet(imgs);  
  
info =  
    gss: [1x1 struct]  
    css: [1x1 struct]  
    peakScores: [1x351 single]  
    edgeScores: [1x351 single]  
    orientationScores: [1x351 single]  
    laplacianScaleScores: [1x351 single]
```



Obtaining the scale space

```
[frames, descrs, info] = vl_covdet(imgs);
```

```
info =
```

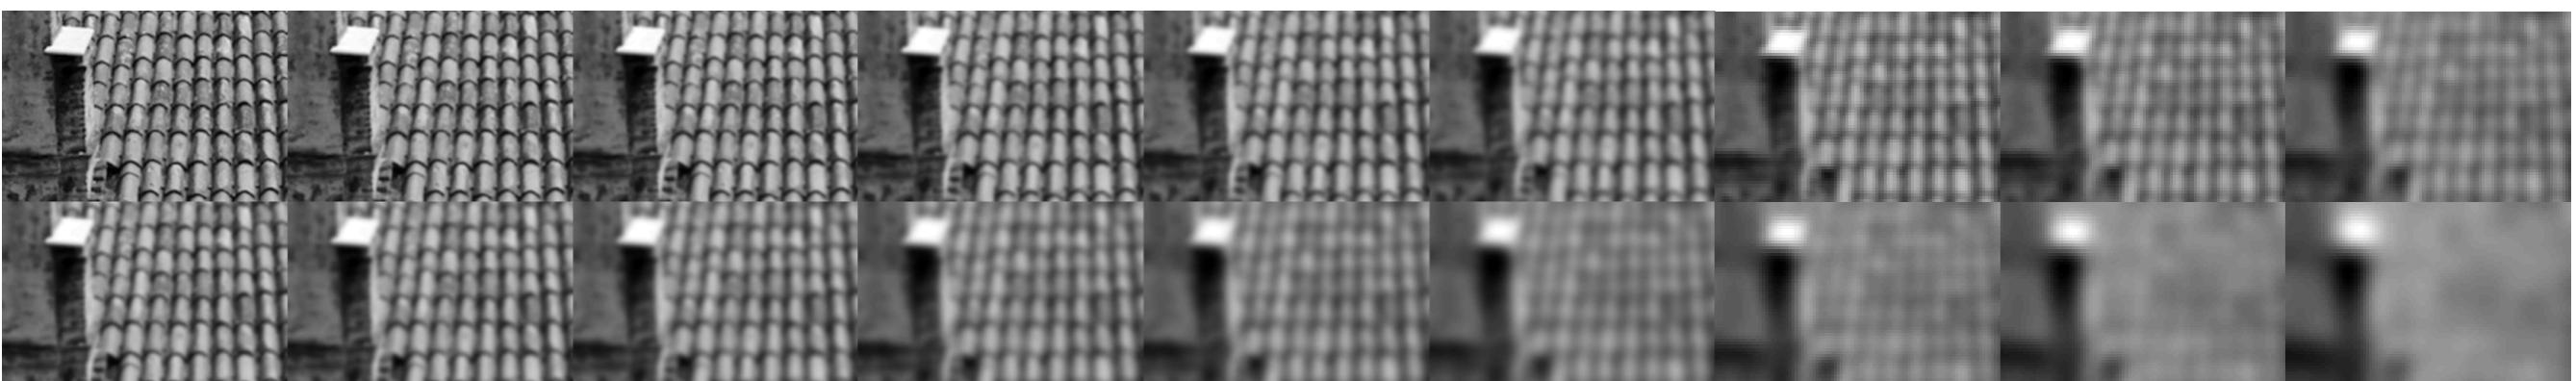
```
    gss: [1x1 struct]
    css: [1x1 struct]
    peakScores: [1x351 single]
    edgeScores: [1x351 single]
    orientationScores: [1x351 single]
    laplacianScaleScores: [1x351 single]
```

Gaussian
scale space

octave -1

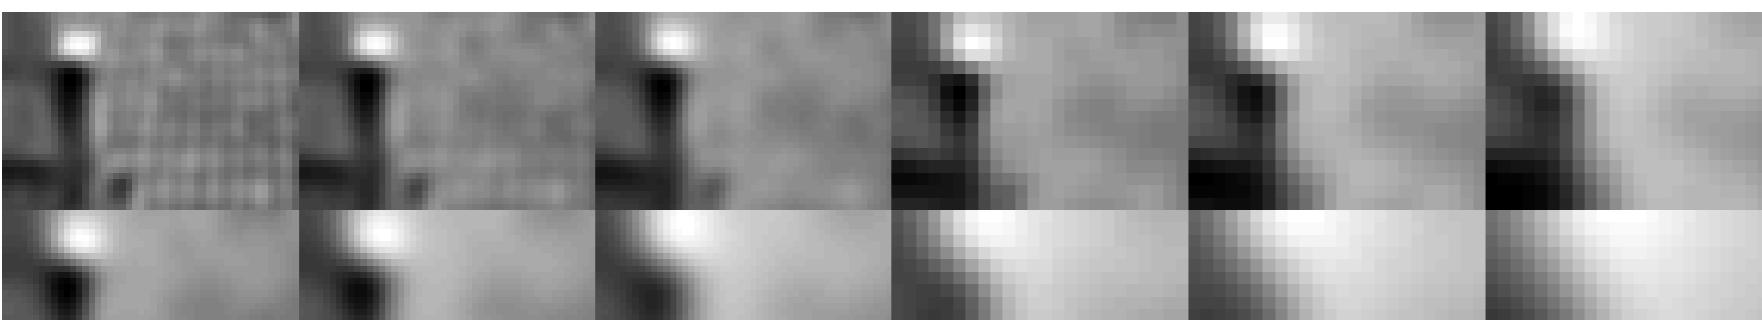
octave 0

octave 1



octave 2

octave 3



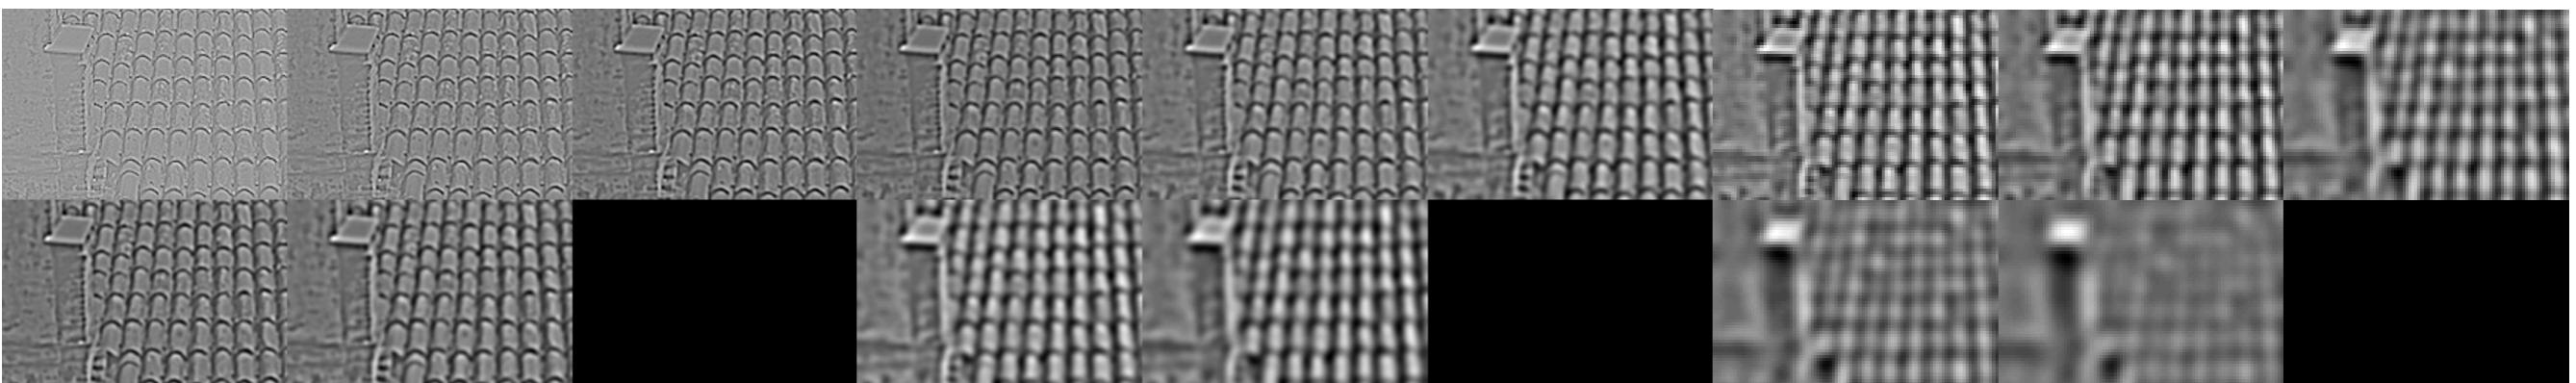
Obtaining the scale space

```
[frames, descrs, info] = vl_covdet(imgs);  
  
info =  
    gss: [1x1 struct]  
    css: [1x1 struct] ← cornerness measure  
    peakScores: [1x351 single]  
    edgeScores: [1x351 single]  
    orientationScores: [1x351 single]  
    laplacianScaleScores: [1x351 single]
```

octave -1

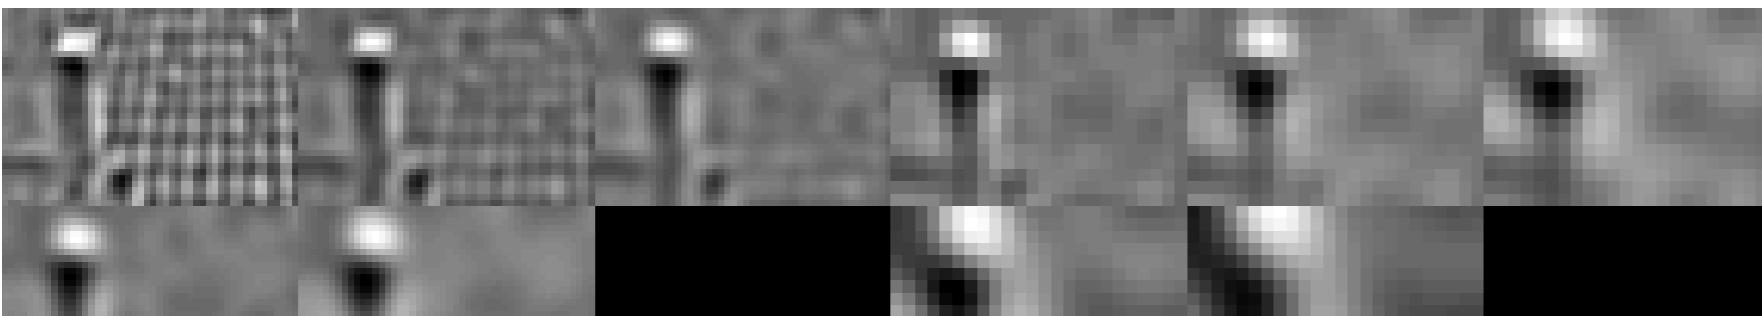
octave 0

octave 1



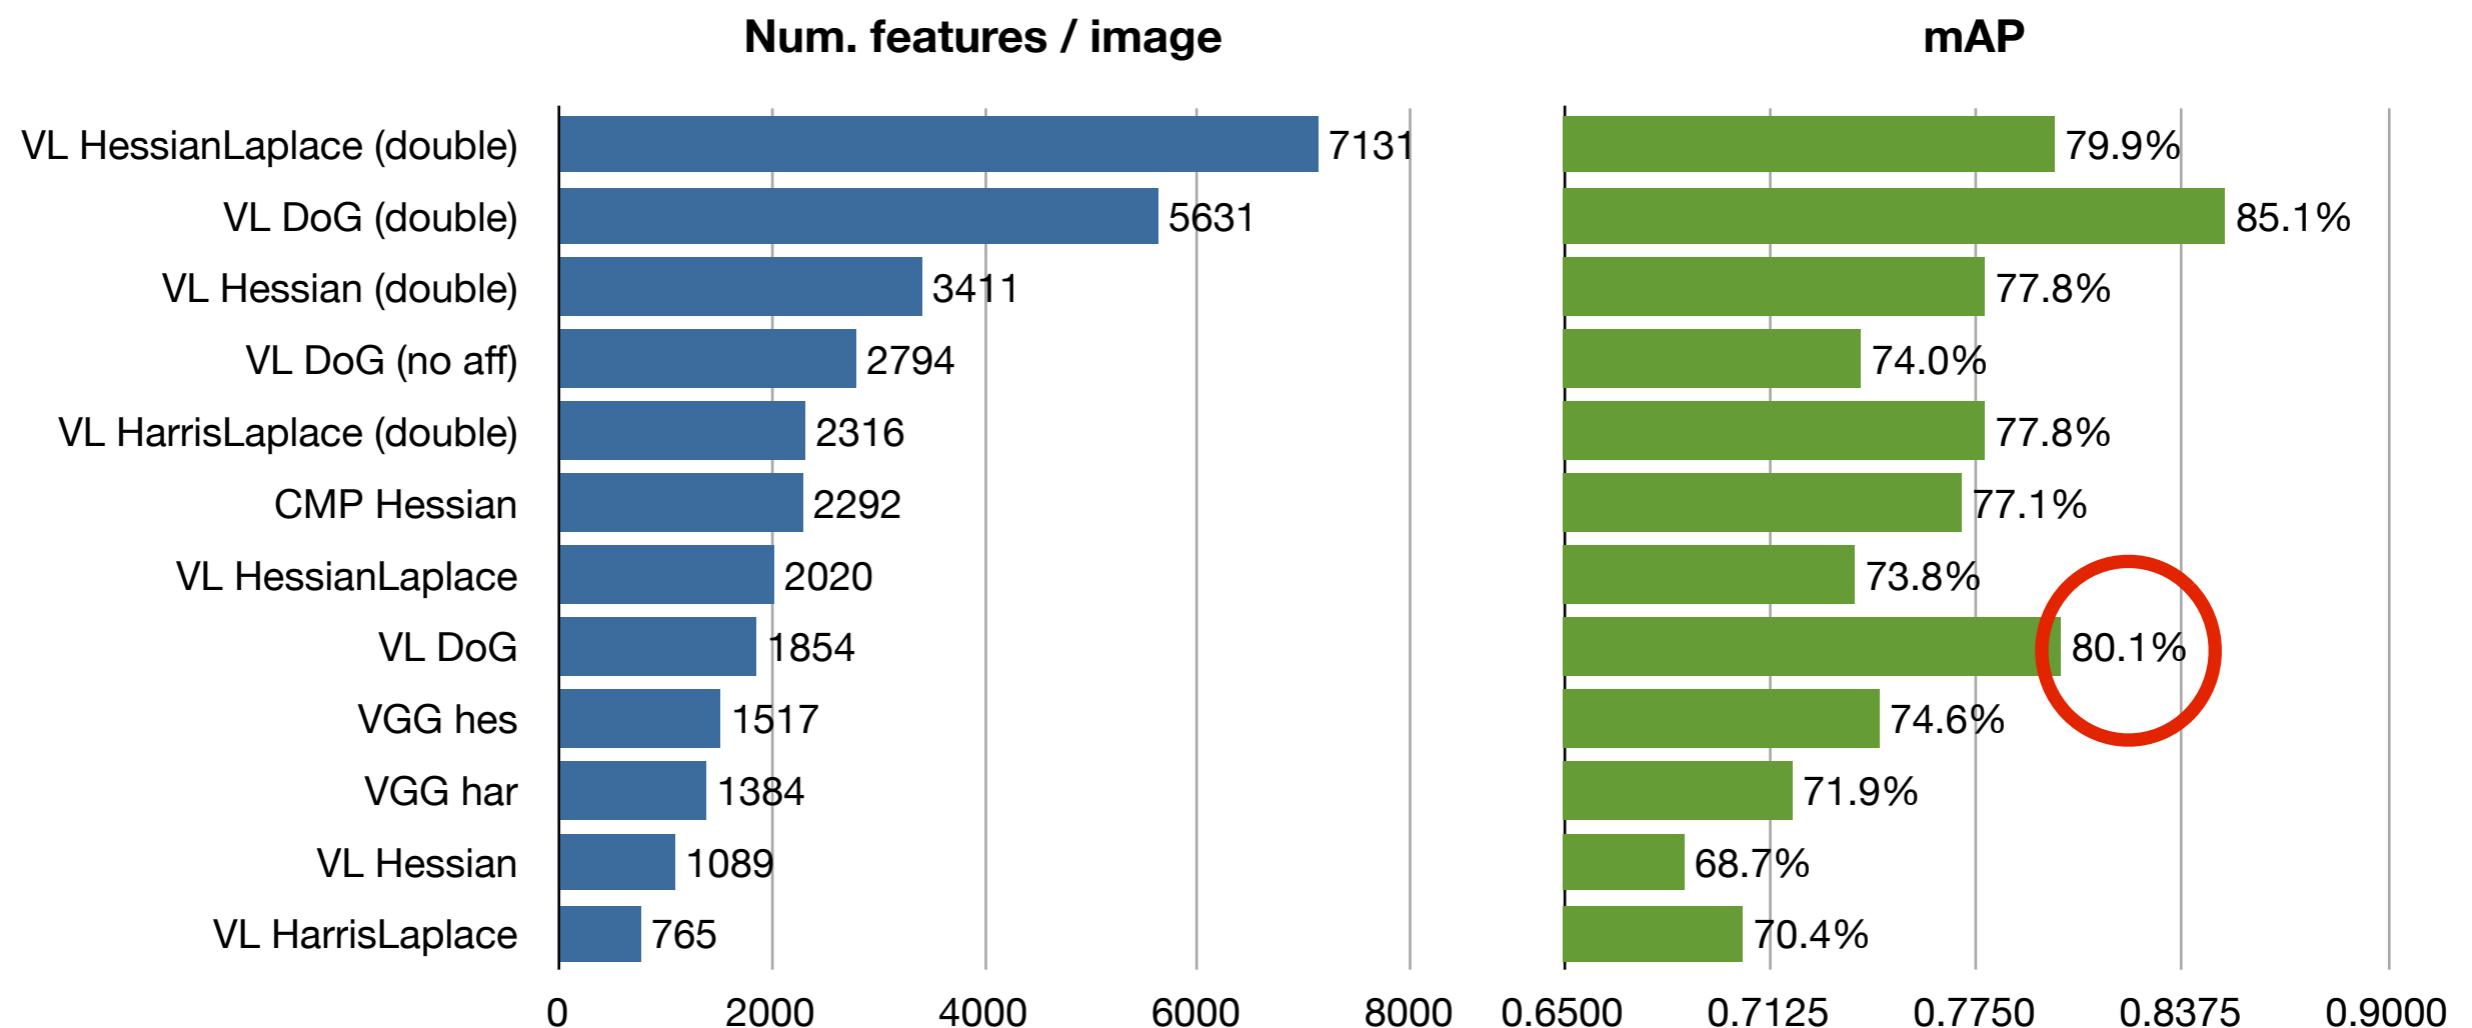
octave 2

octave 3

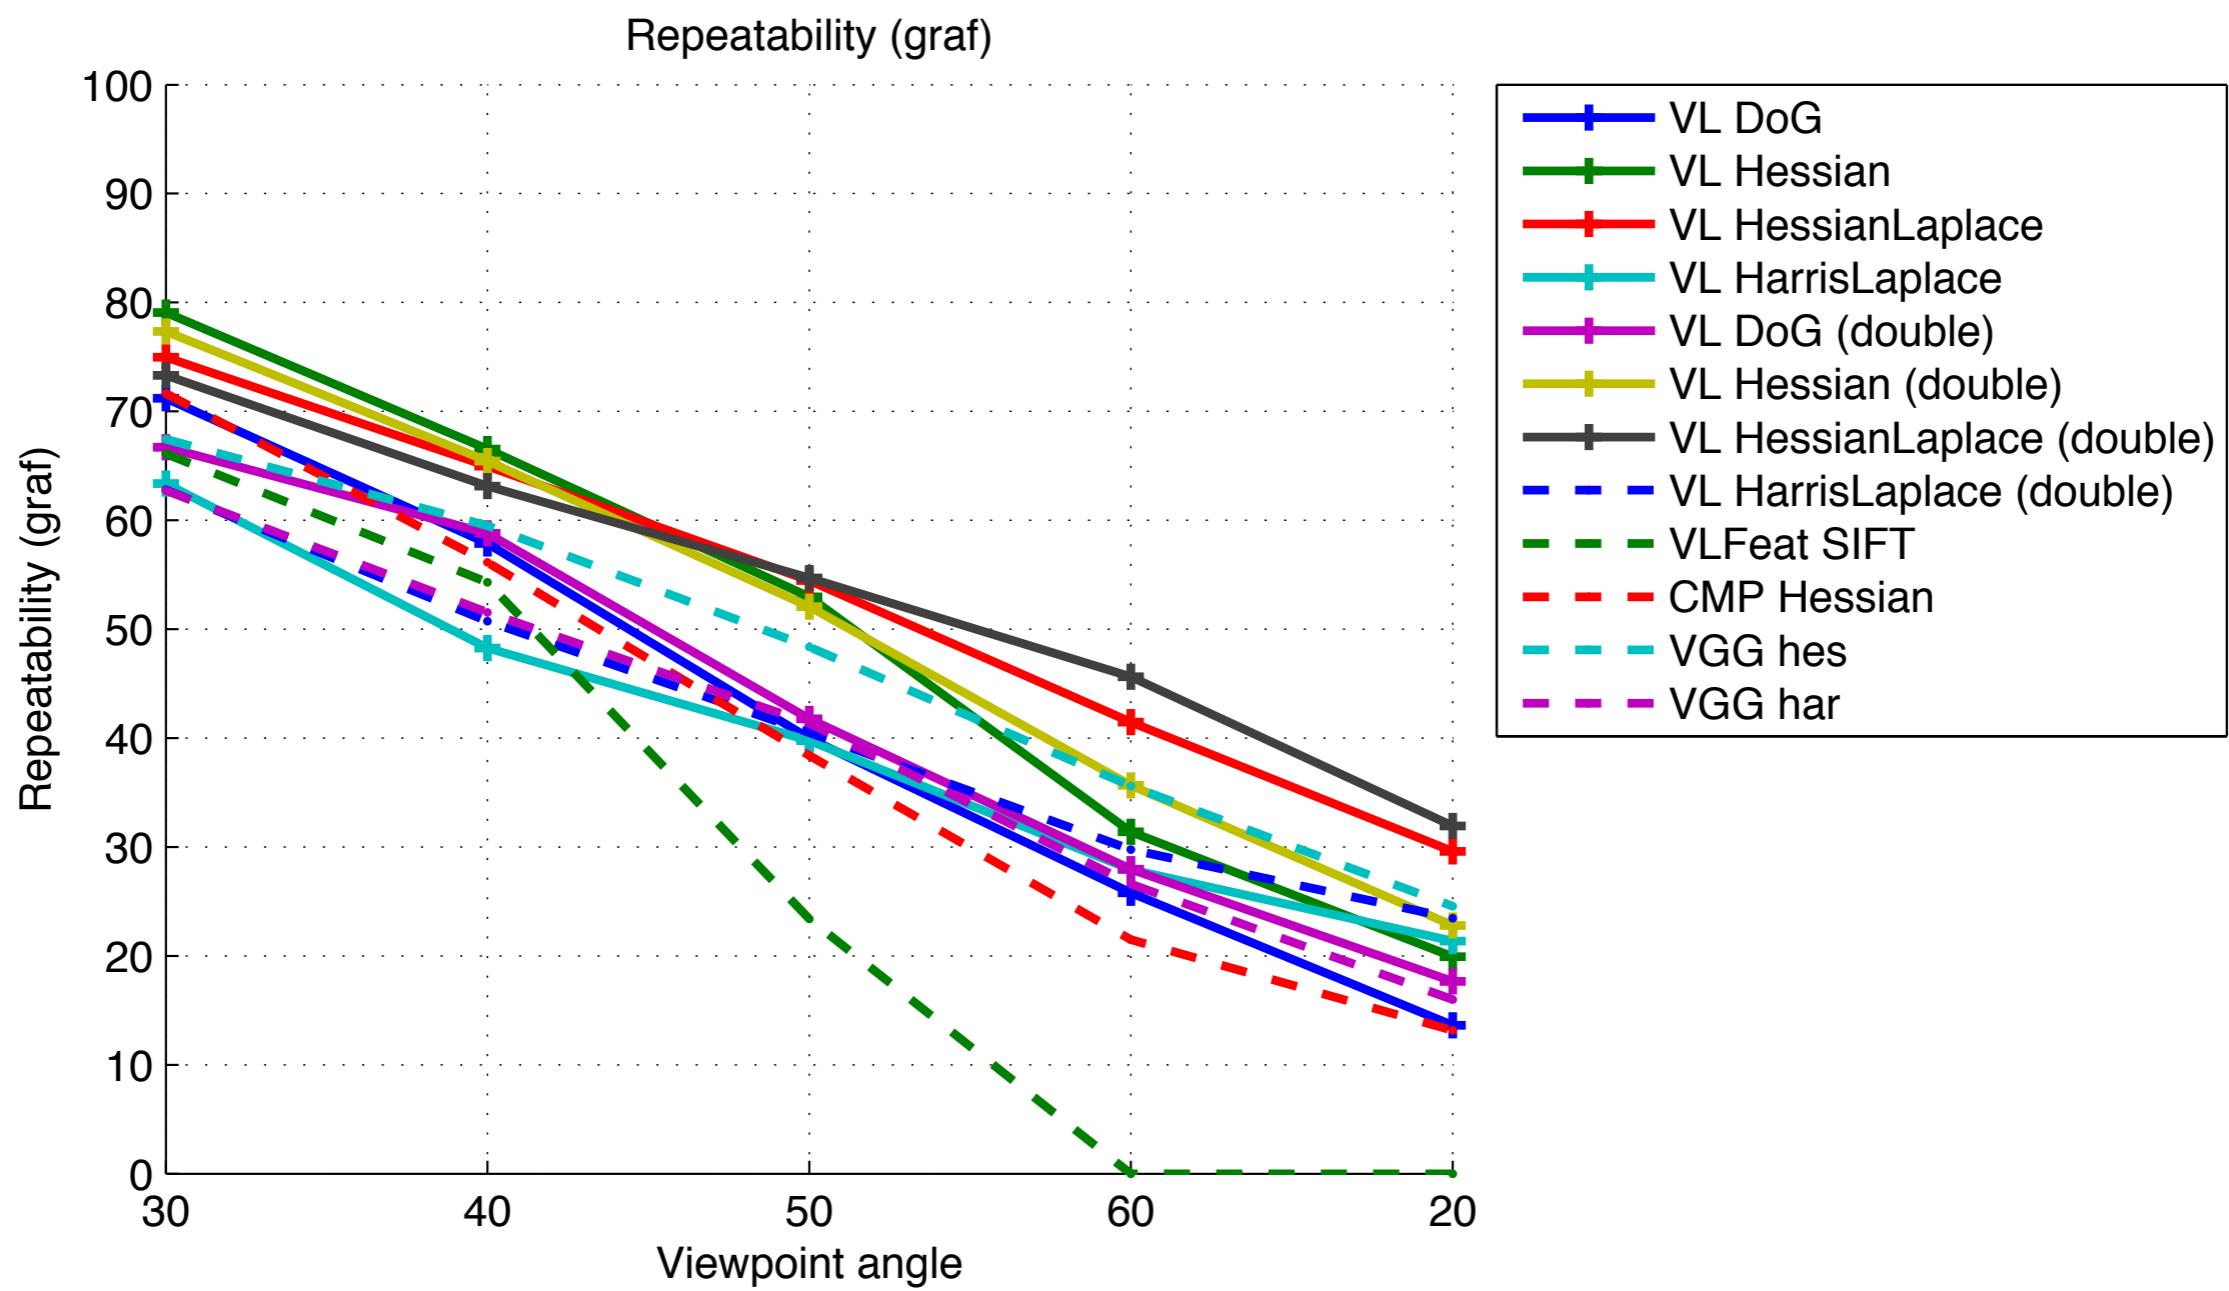
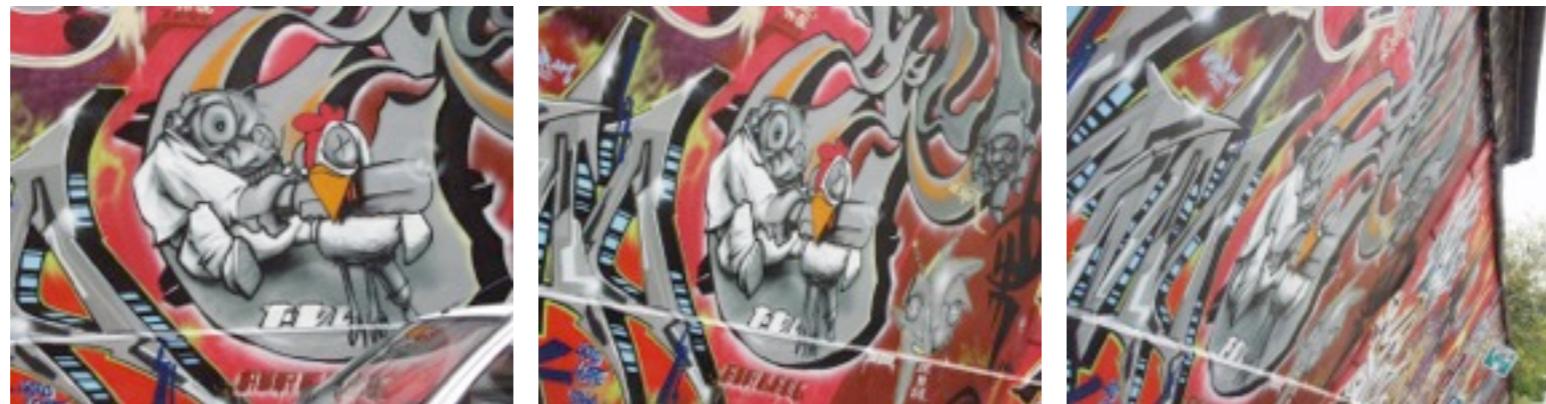


Performance comparison

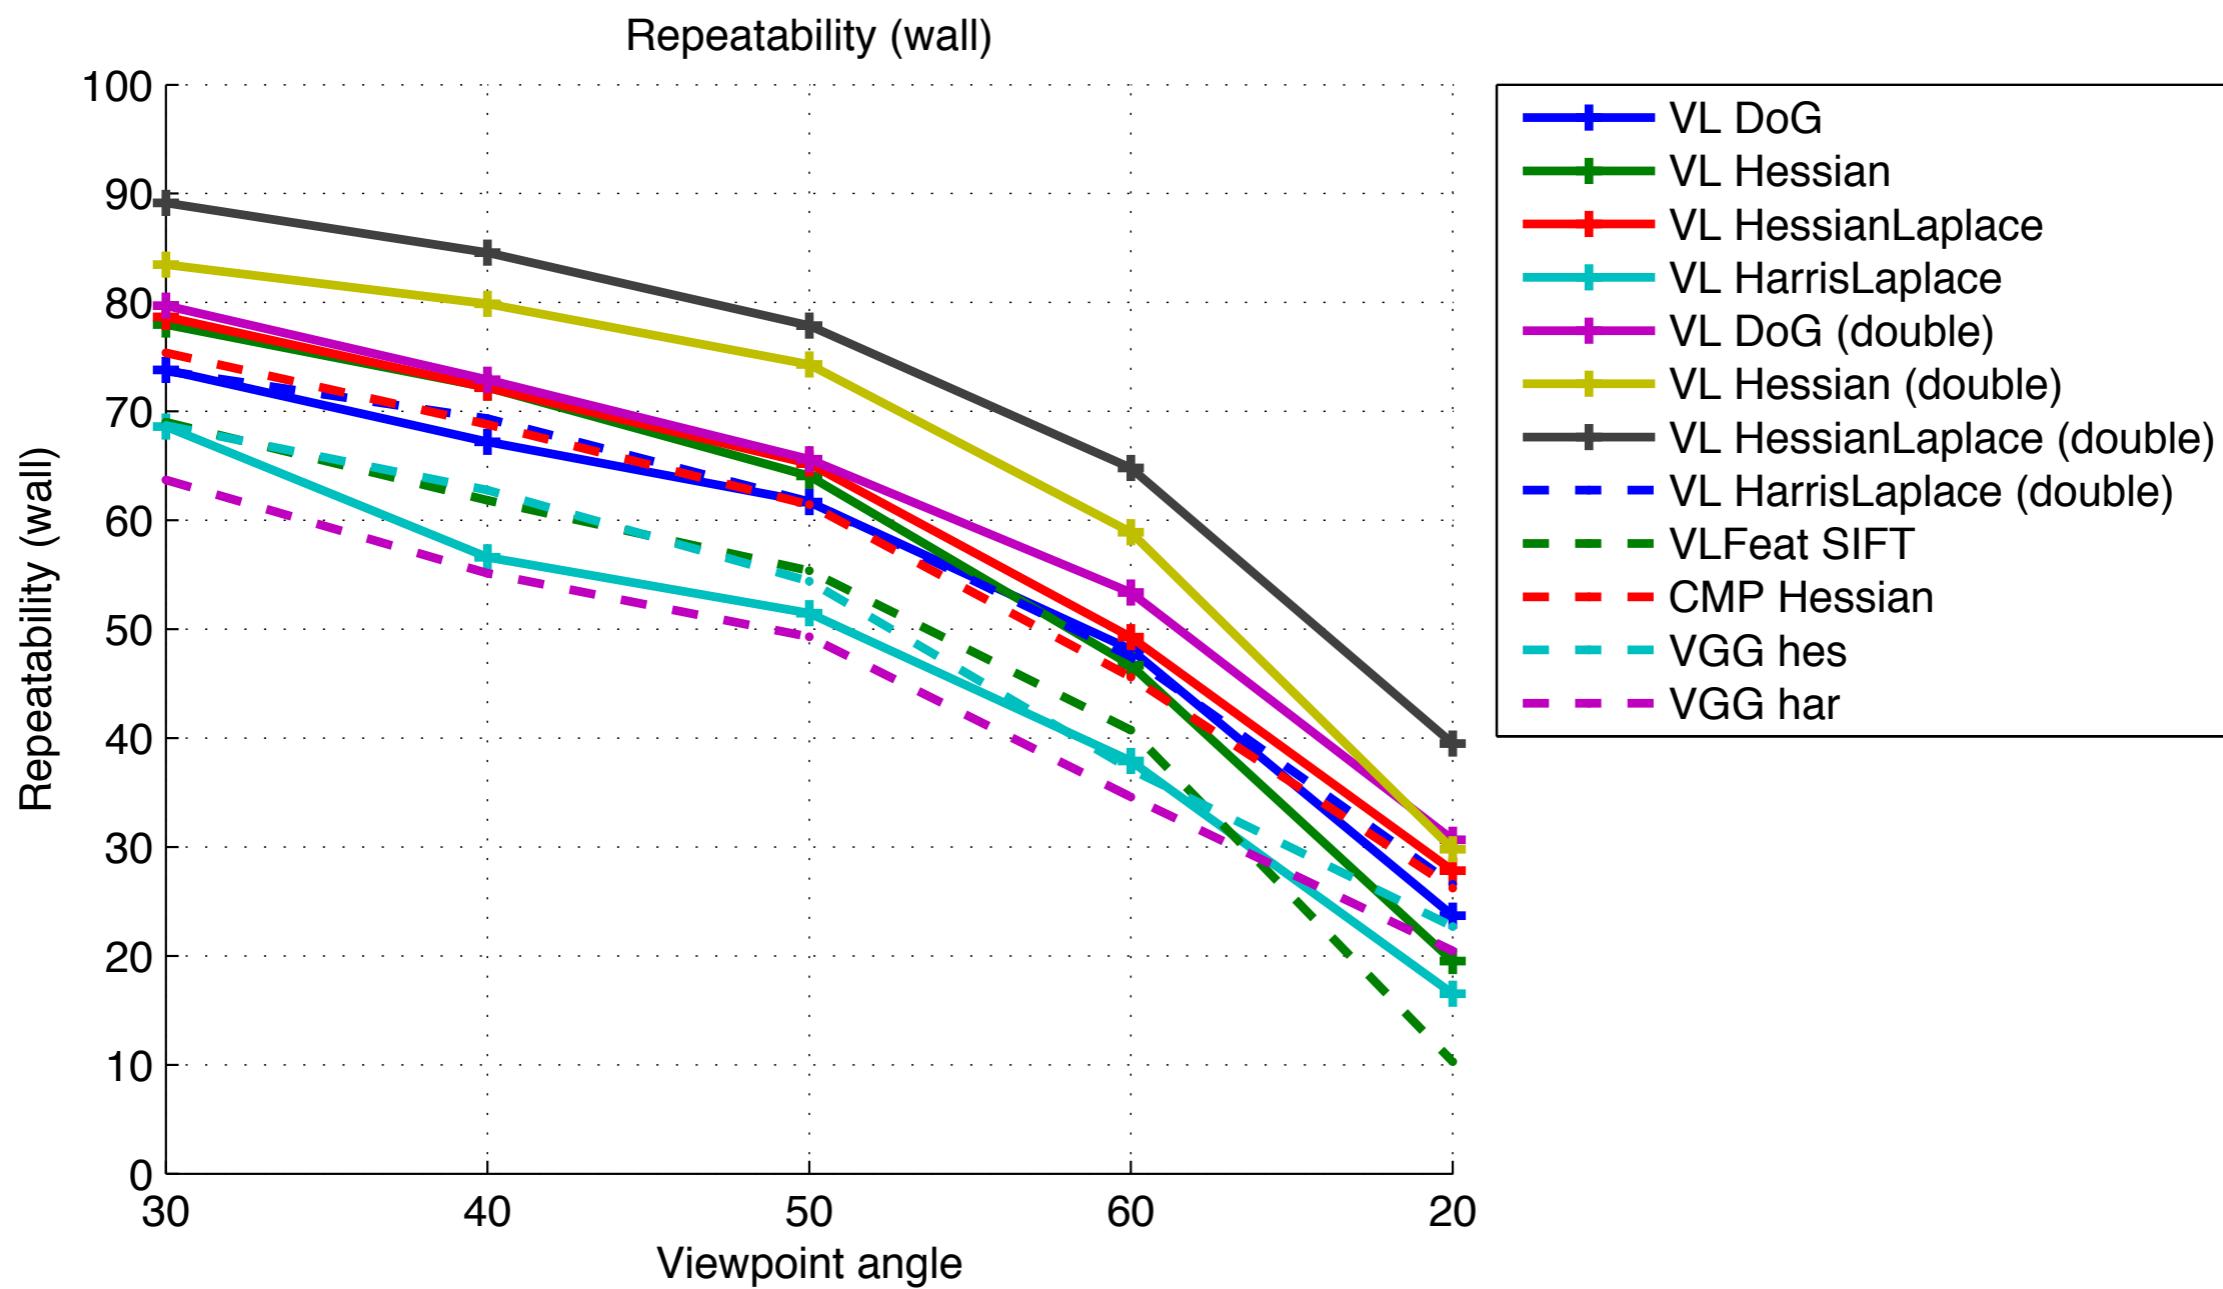
Oxford 5K lite retrieval benchmark



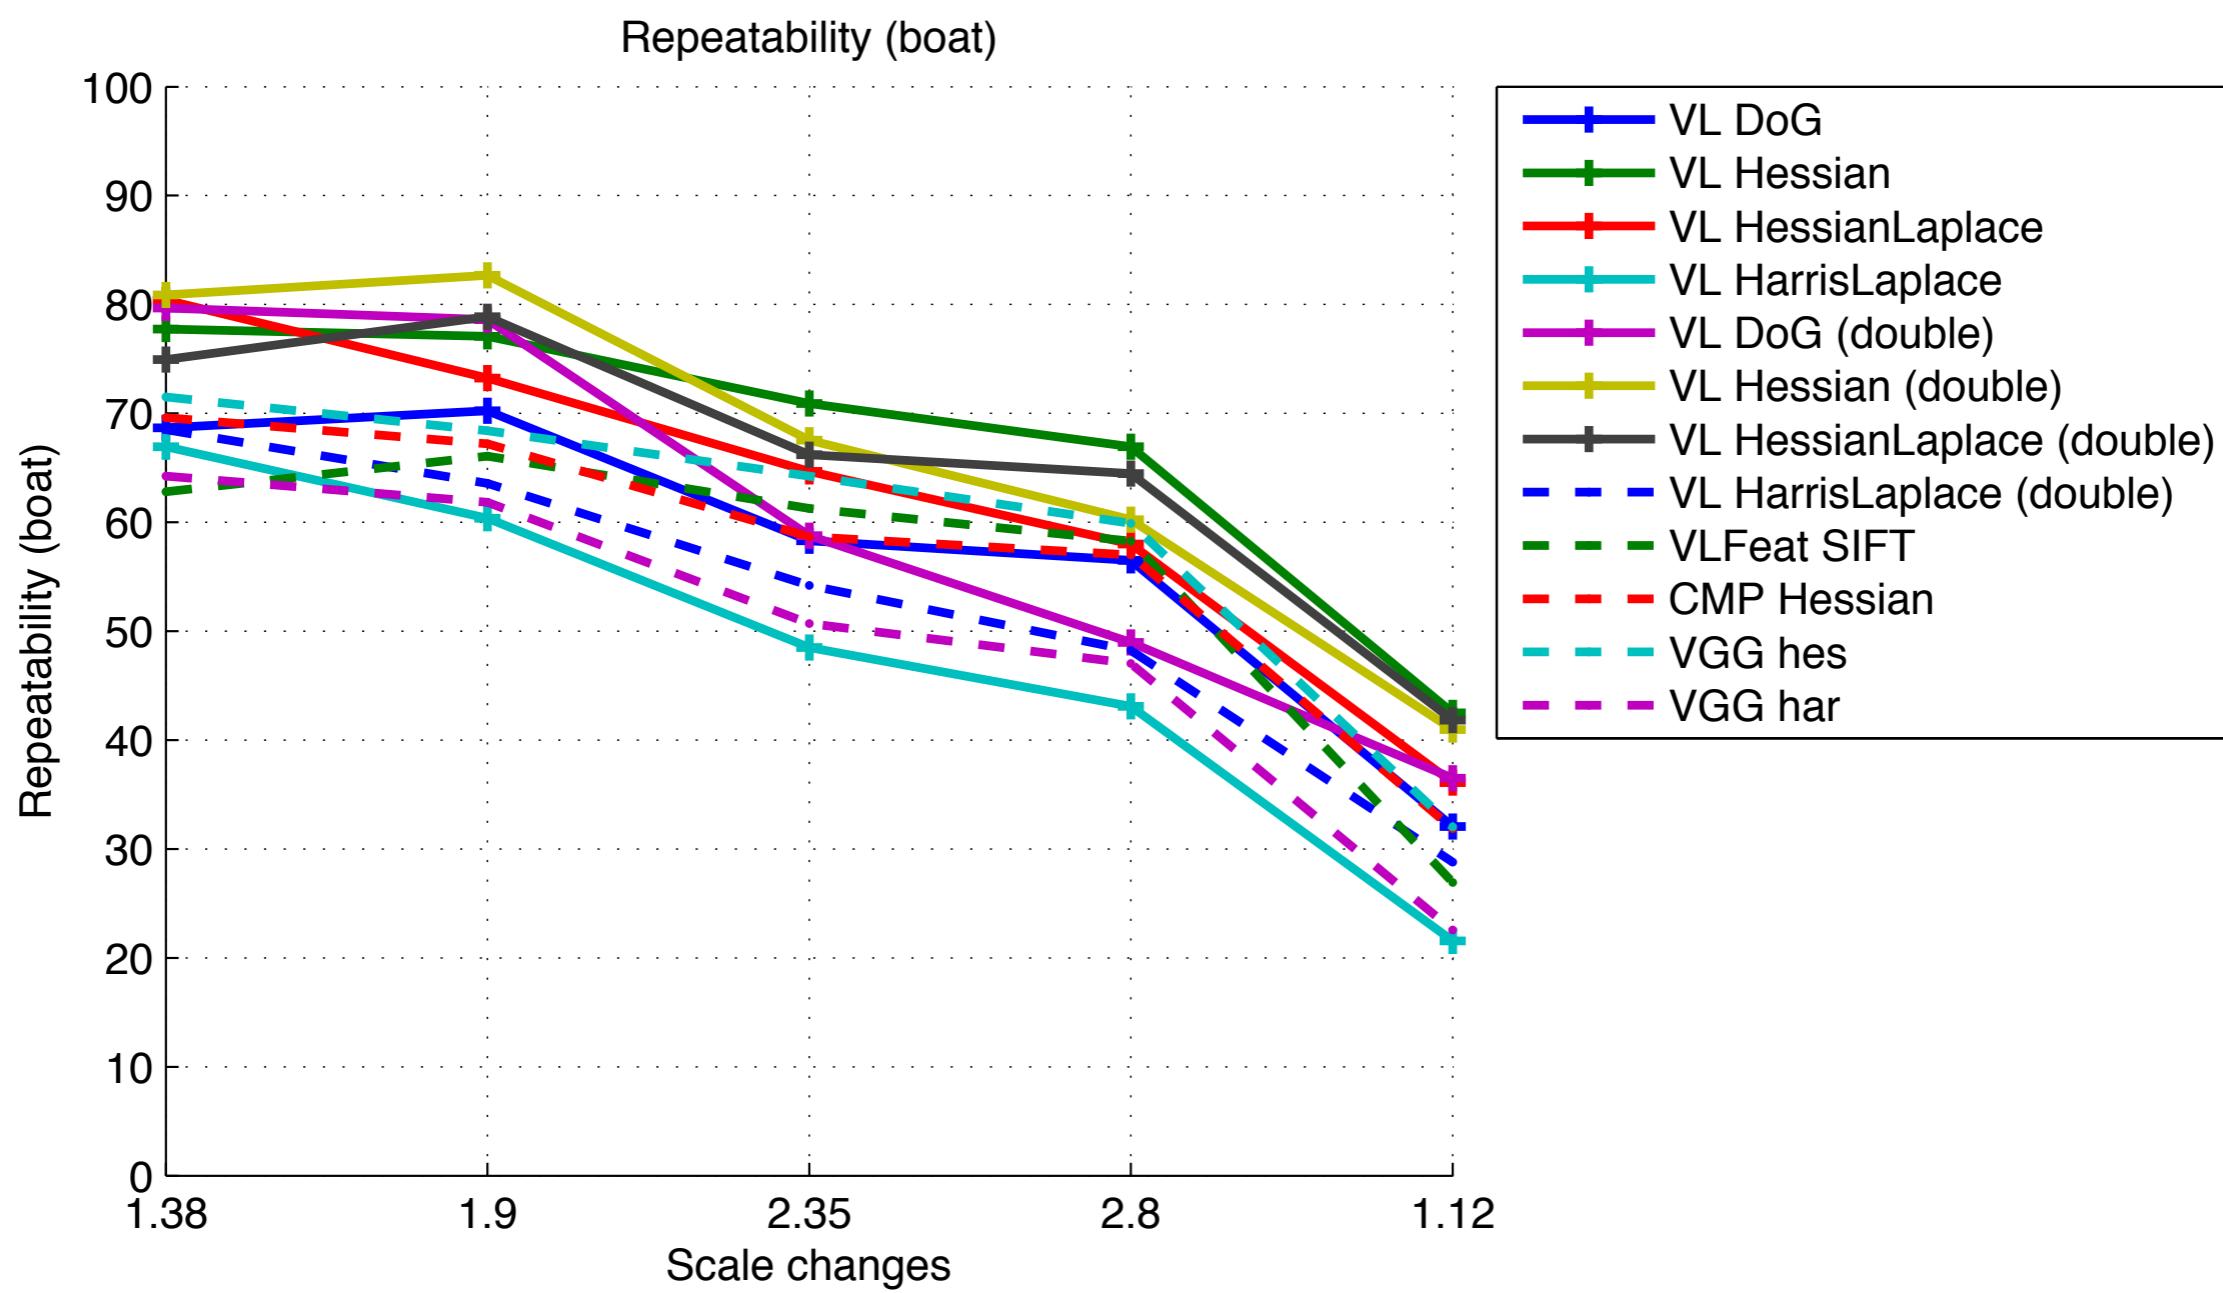
Repeatability: viewpoint



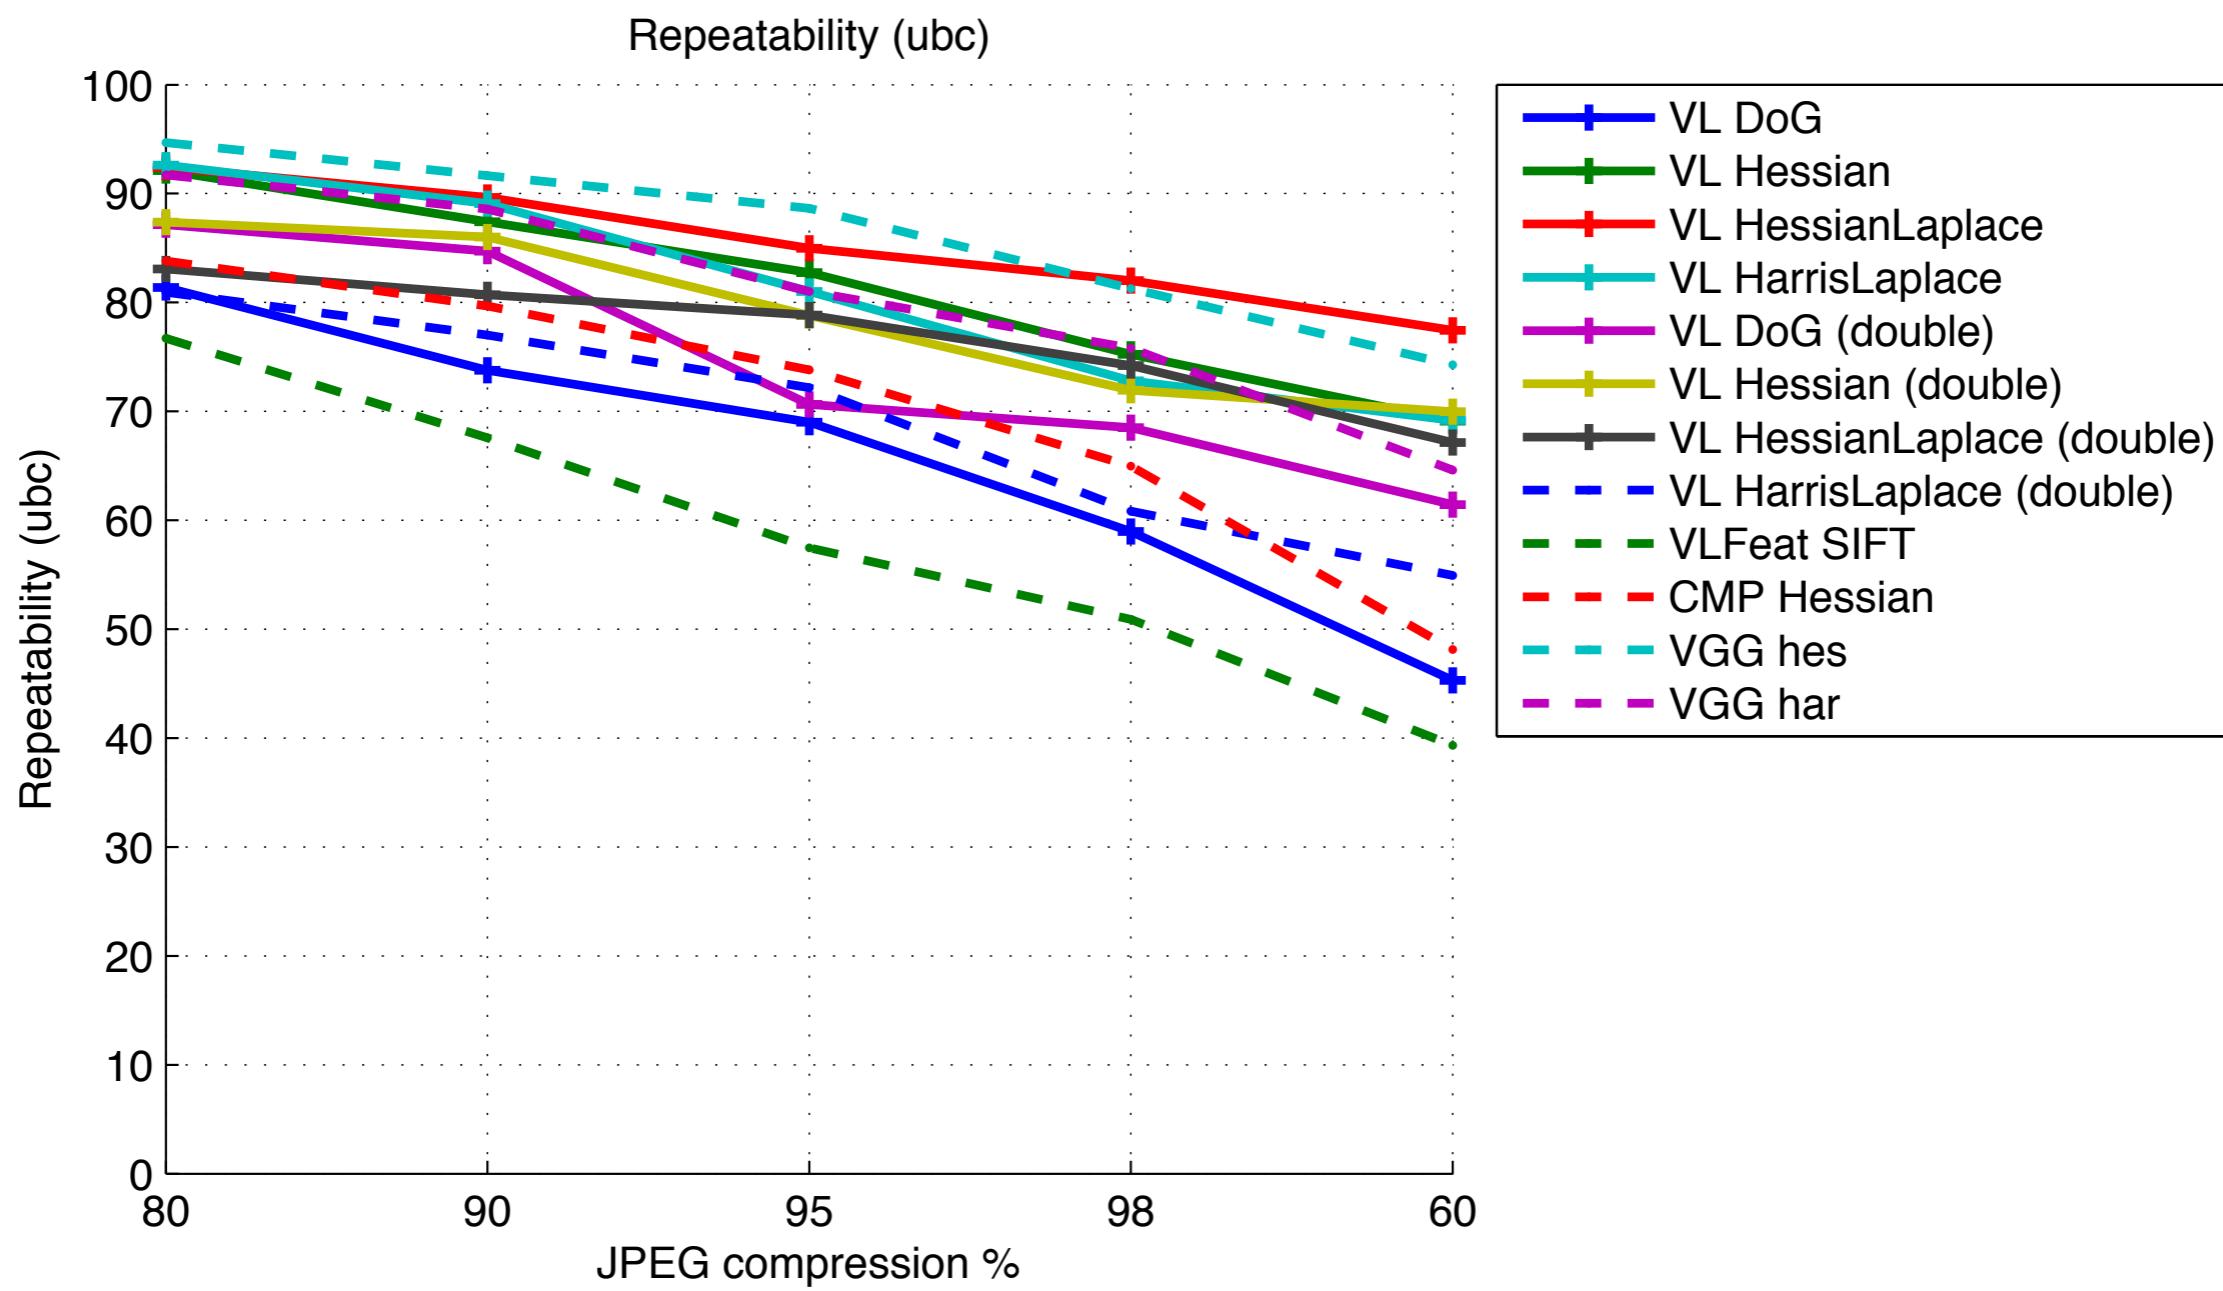
Repeatability: viewpoint



Repeatability: scale



Repeatability: JPEG compression



Maximally Stable Extremal Regions (MSER)

30

[Matas *et al.* 2002]



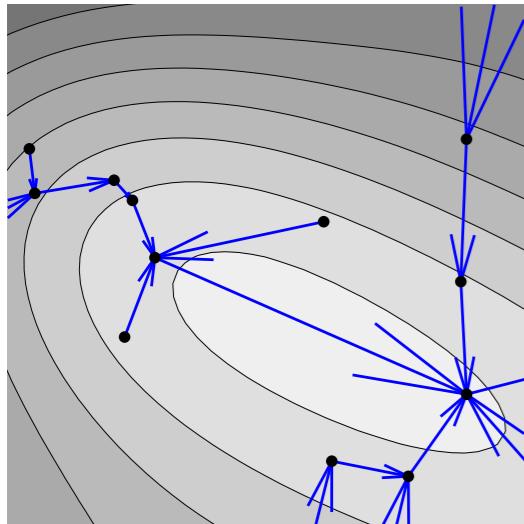
- Extracts MSER regions **regions** and fitted ellipses **frames**:

```
imgs = uint8(rgb2gray(I)) ;  
[regions, frames] = vl_mser(imgs) ;
```

- Control region stability:

```
[regions, frames] = vl_mser(imgs, ...  
    'MinDiversity', 0.7, ...  
    'MaxVariation', 0.2, ...  
    'Delta', 10) ;
```

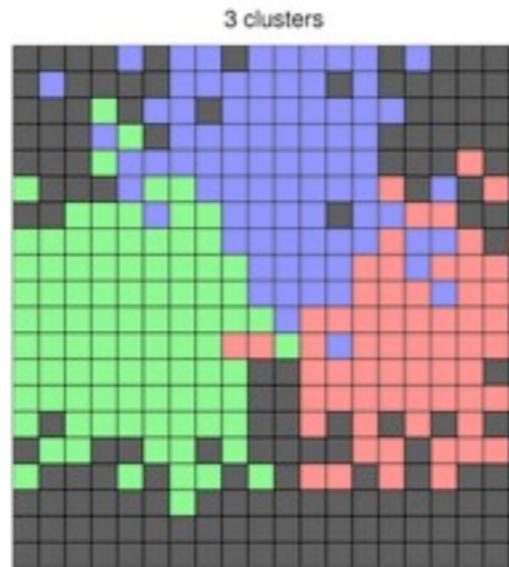
Many other features and tools



superpixels



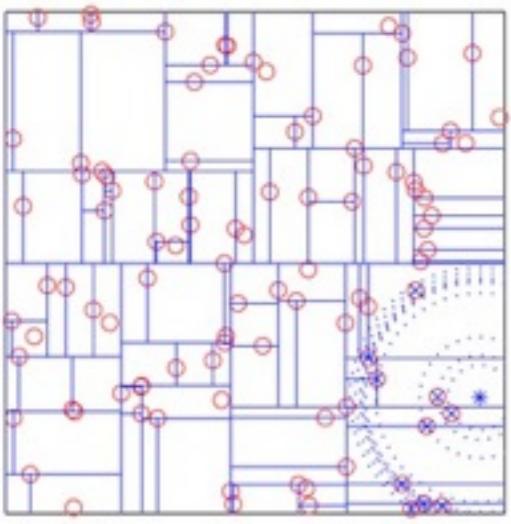
MSER



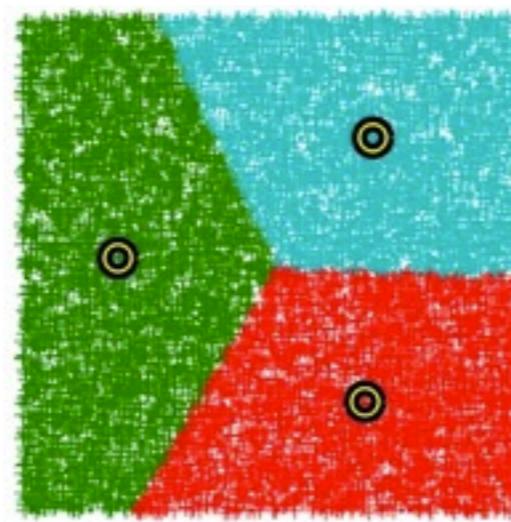
agglomerative information bottleneck



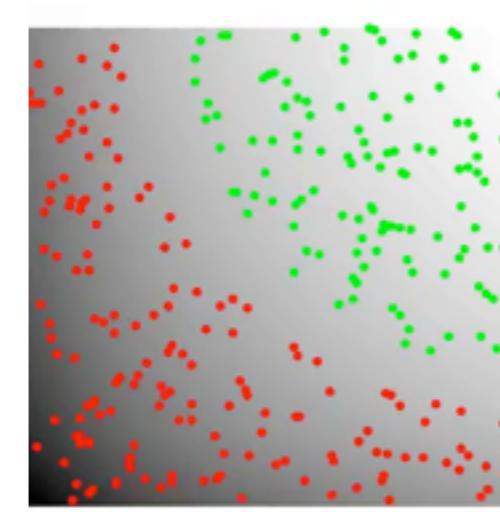
distance transform



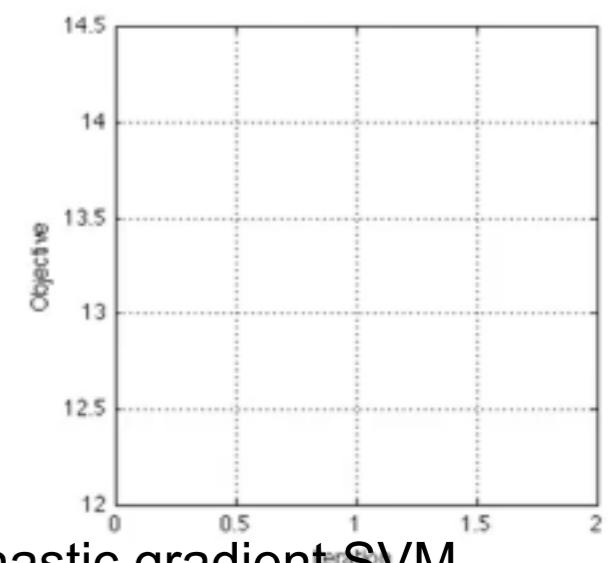
randomized kd-trees



fast k-means



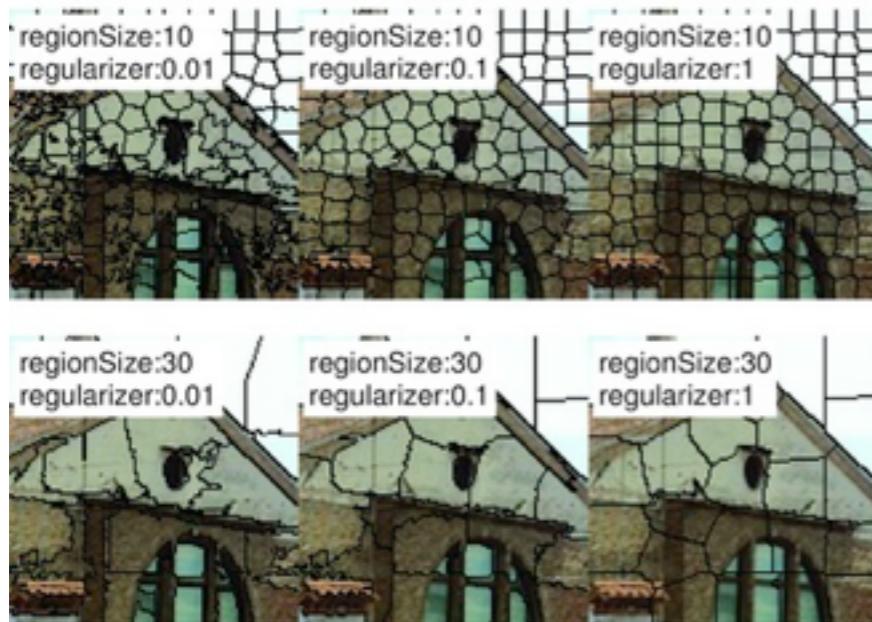
feature maps, stochastic gradient SVM



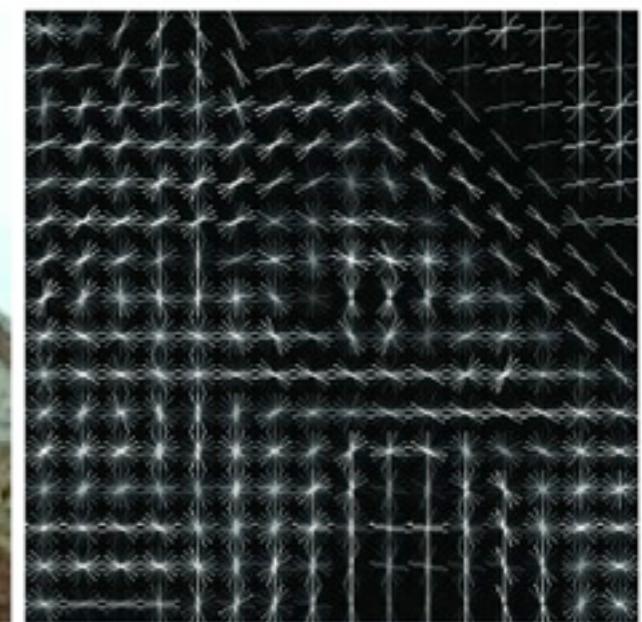
Fast Dense SIFT

Many other features and tools

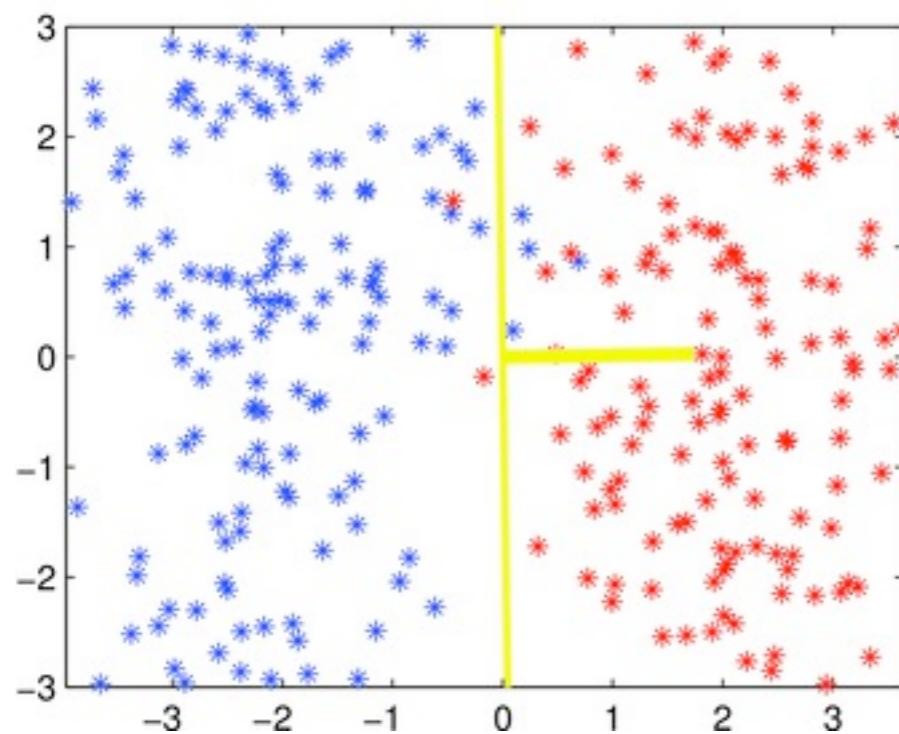
32



SLIC superpixels



Histogram of Oriented Gradients (HOG)



New SGD Learning of SVMs



OpenCV

- The largest, best known vision library
- **Features**
 - *OpenCV 2.1:*
Harris, Tomasi&Shi, MSER, DoG, SURF, STAR and FAST,
DoG + SIFT, SURF, HARRIS, GoodFeaturesToTrack
 - *OpenCV 2.4:*
ORB, BRIEF, and FREAK

Example: computing ORB features.

```
#include <opencv2/core/core.hpp>
#include <opencv2/features2d/features2d.hpp>
#include <vector>

Mat image = importImage8UC1(in[IN_I]);
Mat descriptors;
std::vector<KeyPoint> frames;

ORB orb(numFeatures, scaleFactor, numLevels, edgeThreshold, 0,
        wtak, scoreTypeToORBPar(scoreType), patchSize);

orb(image, Mat(), frames);
```

This is *not* all folks

- Abundance of open source / free software for CV
 - Always check out authors home pages
- Examples

Rob Hess SIFT

<http://robwhess.github.com/opensift/>

SURF

<http://www.vision.ee.ethz.ch/~surf/>

Open SURF

<http://www.chrisevansdev.com/computer-vision-opensurf.html>

FAST corners

<http://www.edwardrosten.com/work/fast.html>

LIOP descriptors

<http://vision.ia.ac.cn/Students/bfan/index.htm>



Credits



Karel Lenc



Daniele Perrone



Michal Perdoch



Krystian Mikolajczyk

Tinne Tuytelaars

Jiri Matas

Cordelia Schmid

Andrew Zisserman

Feature evaluation

Old and new benchmarks, and new software

Andrea Vedaldi, University of Oxford

Benchmarking: why and how

- Dozens of feature detectors and descriptors have been proposed
- **Benchmarks**
 - compare methods empirically
 - select the best method for a task
- **Public benchmarks**
 - reproducible research
 - simplify your life!
- **Ingredients** of a benchmark:

Theory

Data

Software

Indirect evaluation

Repeatability and matching score

Data: affine covariant testbed

Direct evaluation

Image retrieval

Data: oxford 5k

Software

VLBenchmarks

Indirect evaluation

Repeatability and matching score

Data: affine covariant testbed

Direct evaluation

Image retrieval

Data: oxford 5k

Software

VLBenchmarks

Indirect feature evaluation

- **Intuition**

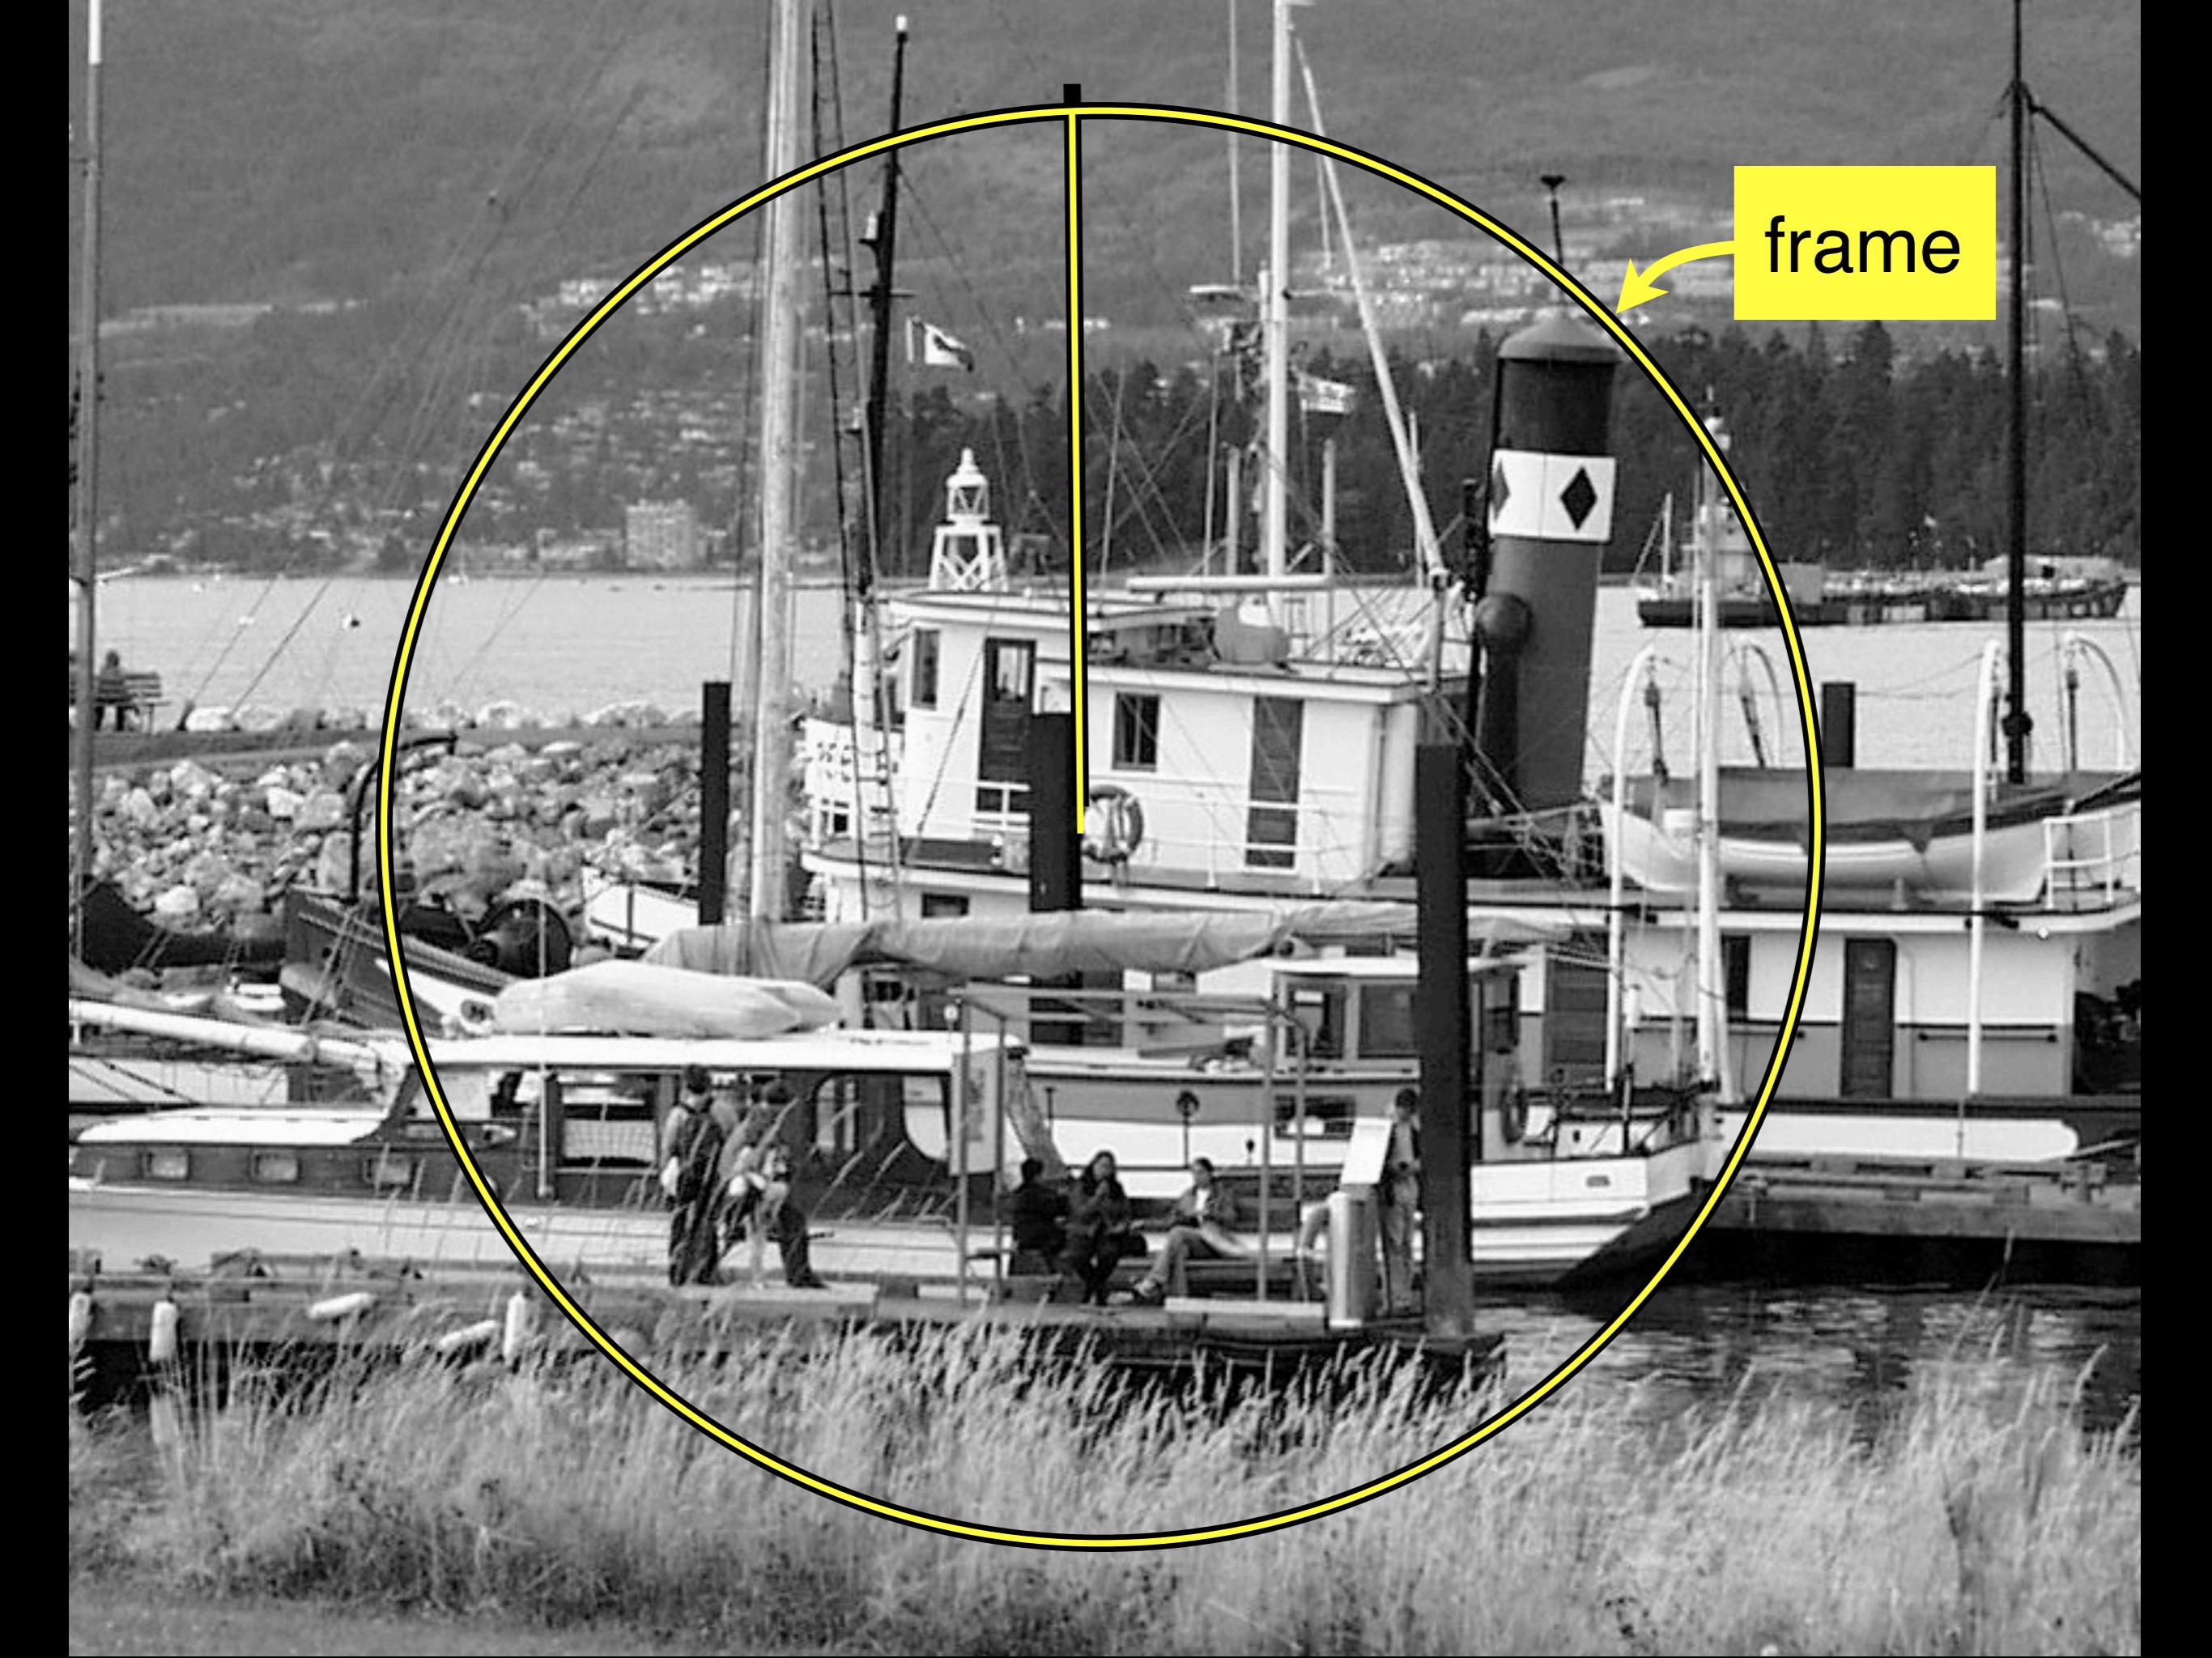
Test how well features **persist** and **can be matched** across image transformations.

- **Data**

- *must be representative* of the transformations
(viewpoint, illumination, noise, etc.)

- **Performance measures**

- **repeatability**
persistence of features
 - **matching score**
matchability of features



frame









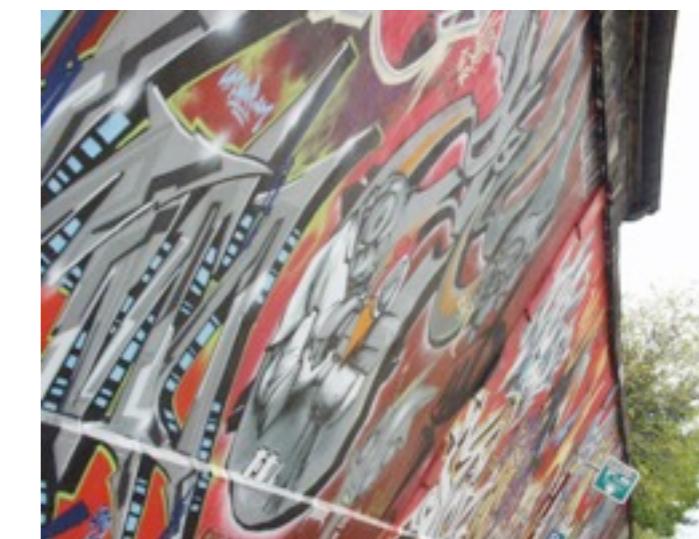






Affine Testbed

Viewpoint, scale, rotation



Affine Testbed

Lighting, compression, blur

50



Detector repeatability

Intuition



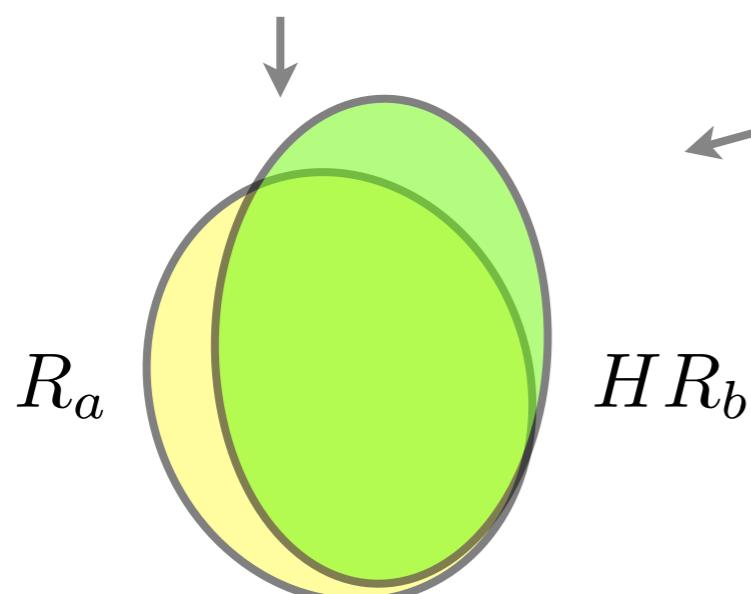
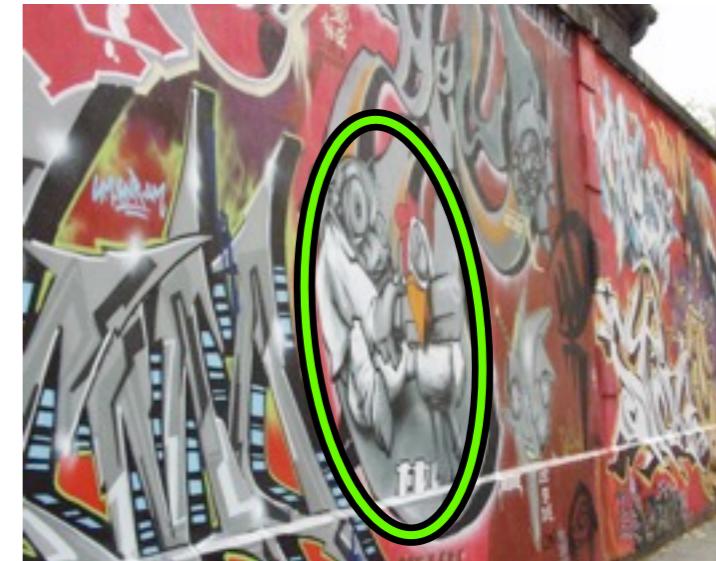
- two pairs of features correspond
- another pair of features does not
- repeatability = 2/3

Region overlap

Formal definition



H
homography

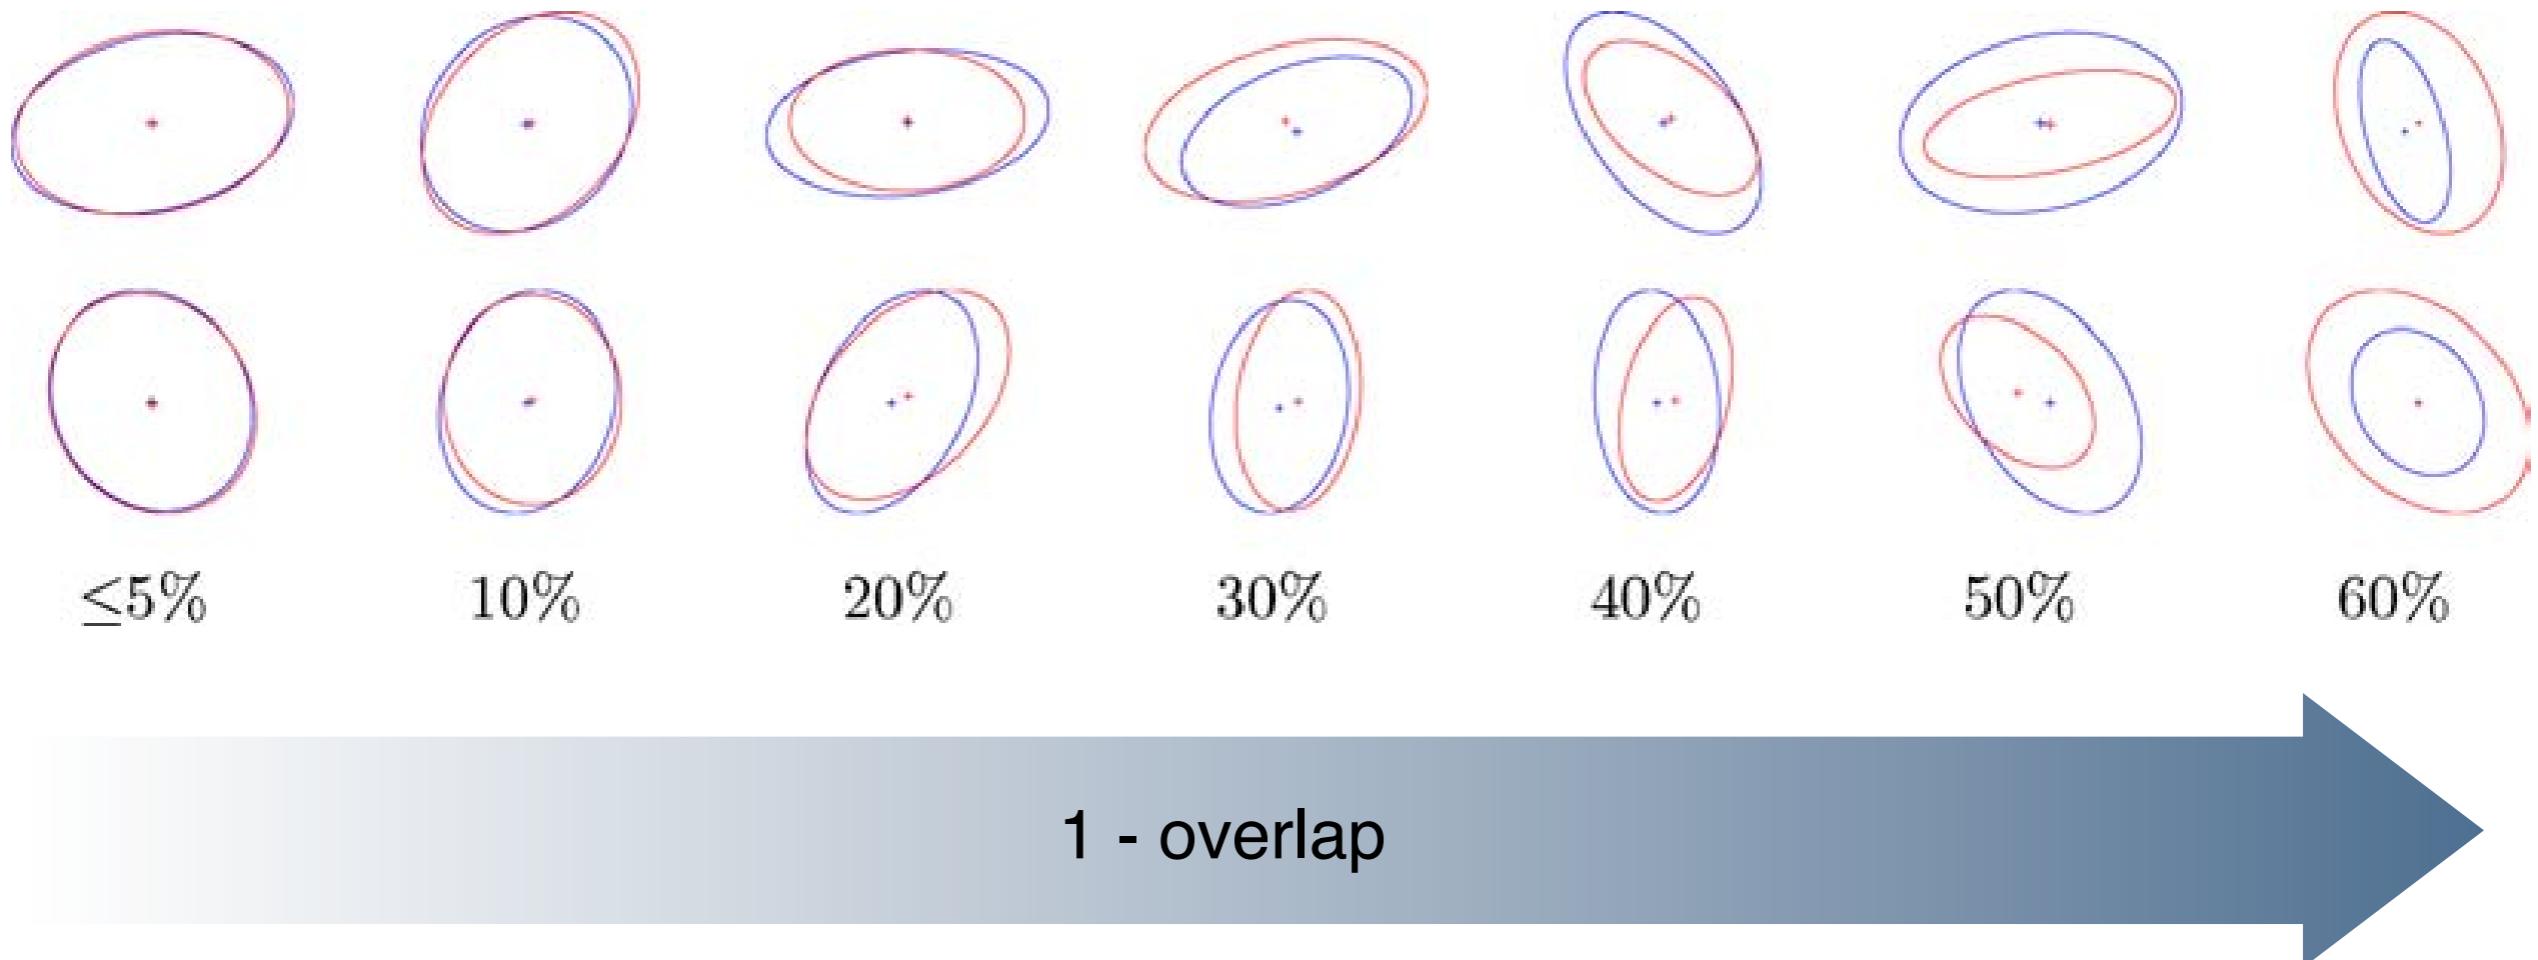


$$\text{overlap}(a, b) = \frac{|R_a \cap HR_b|}{|R_a \cup HR_b|}$$

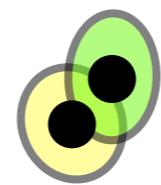
$$= \frac{\text{area intersection}}{\text{area union}}$$

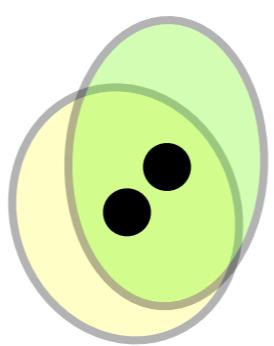
Region overlap

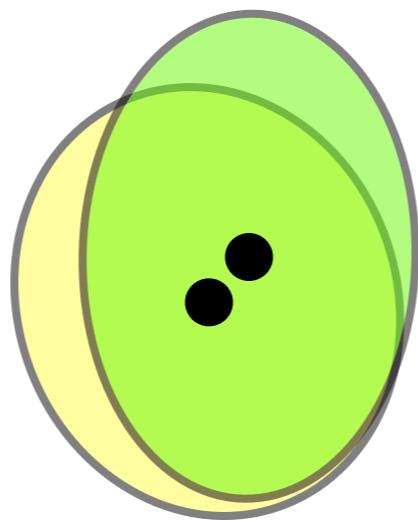
Intuition

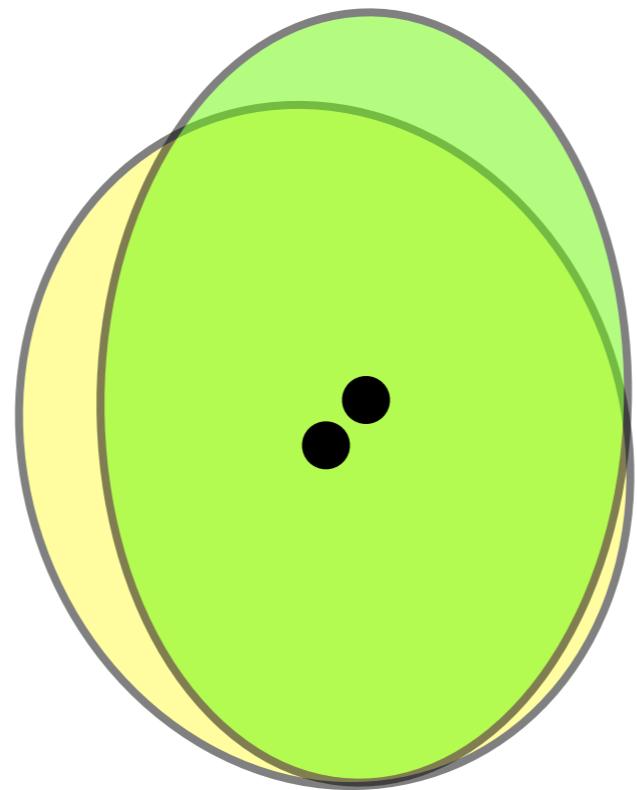


- Examples of ellipses overlapping by different amounts
- Usually, features are tested at 40% overlap *error* (= 60% overlap)



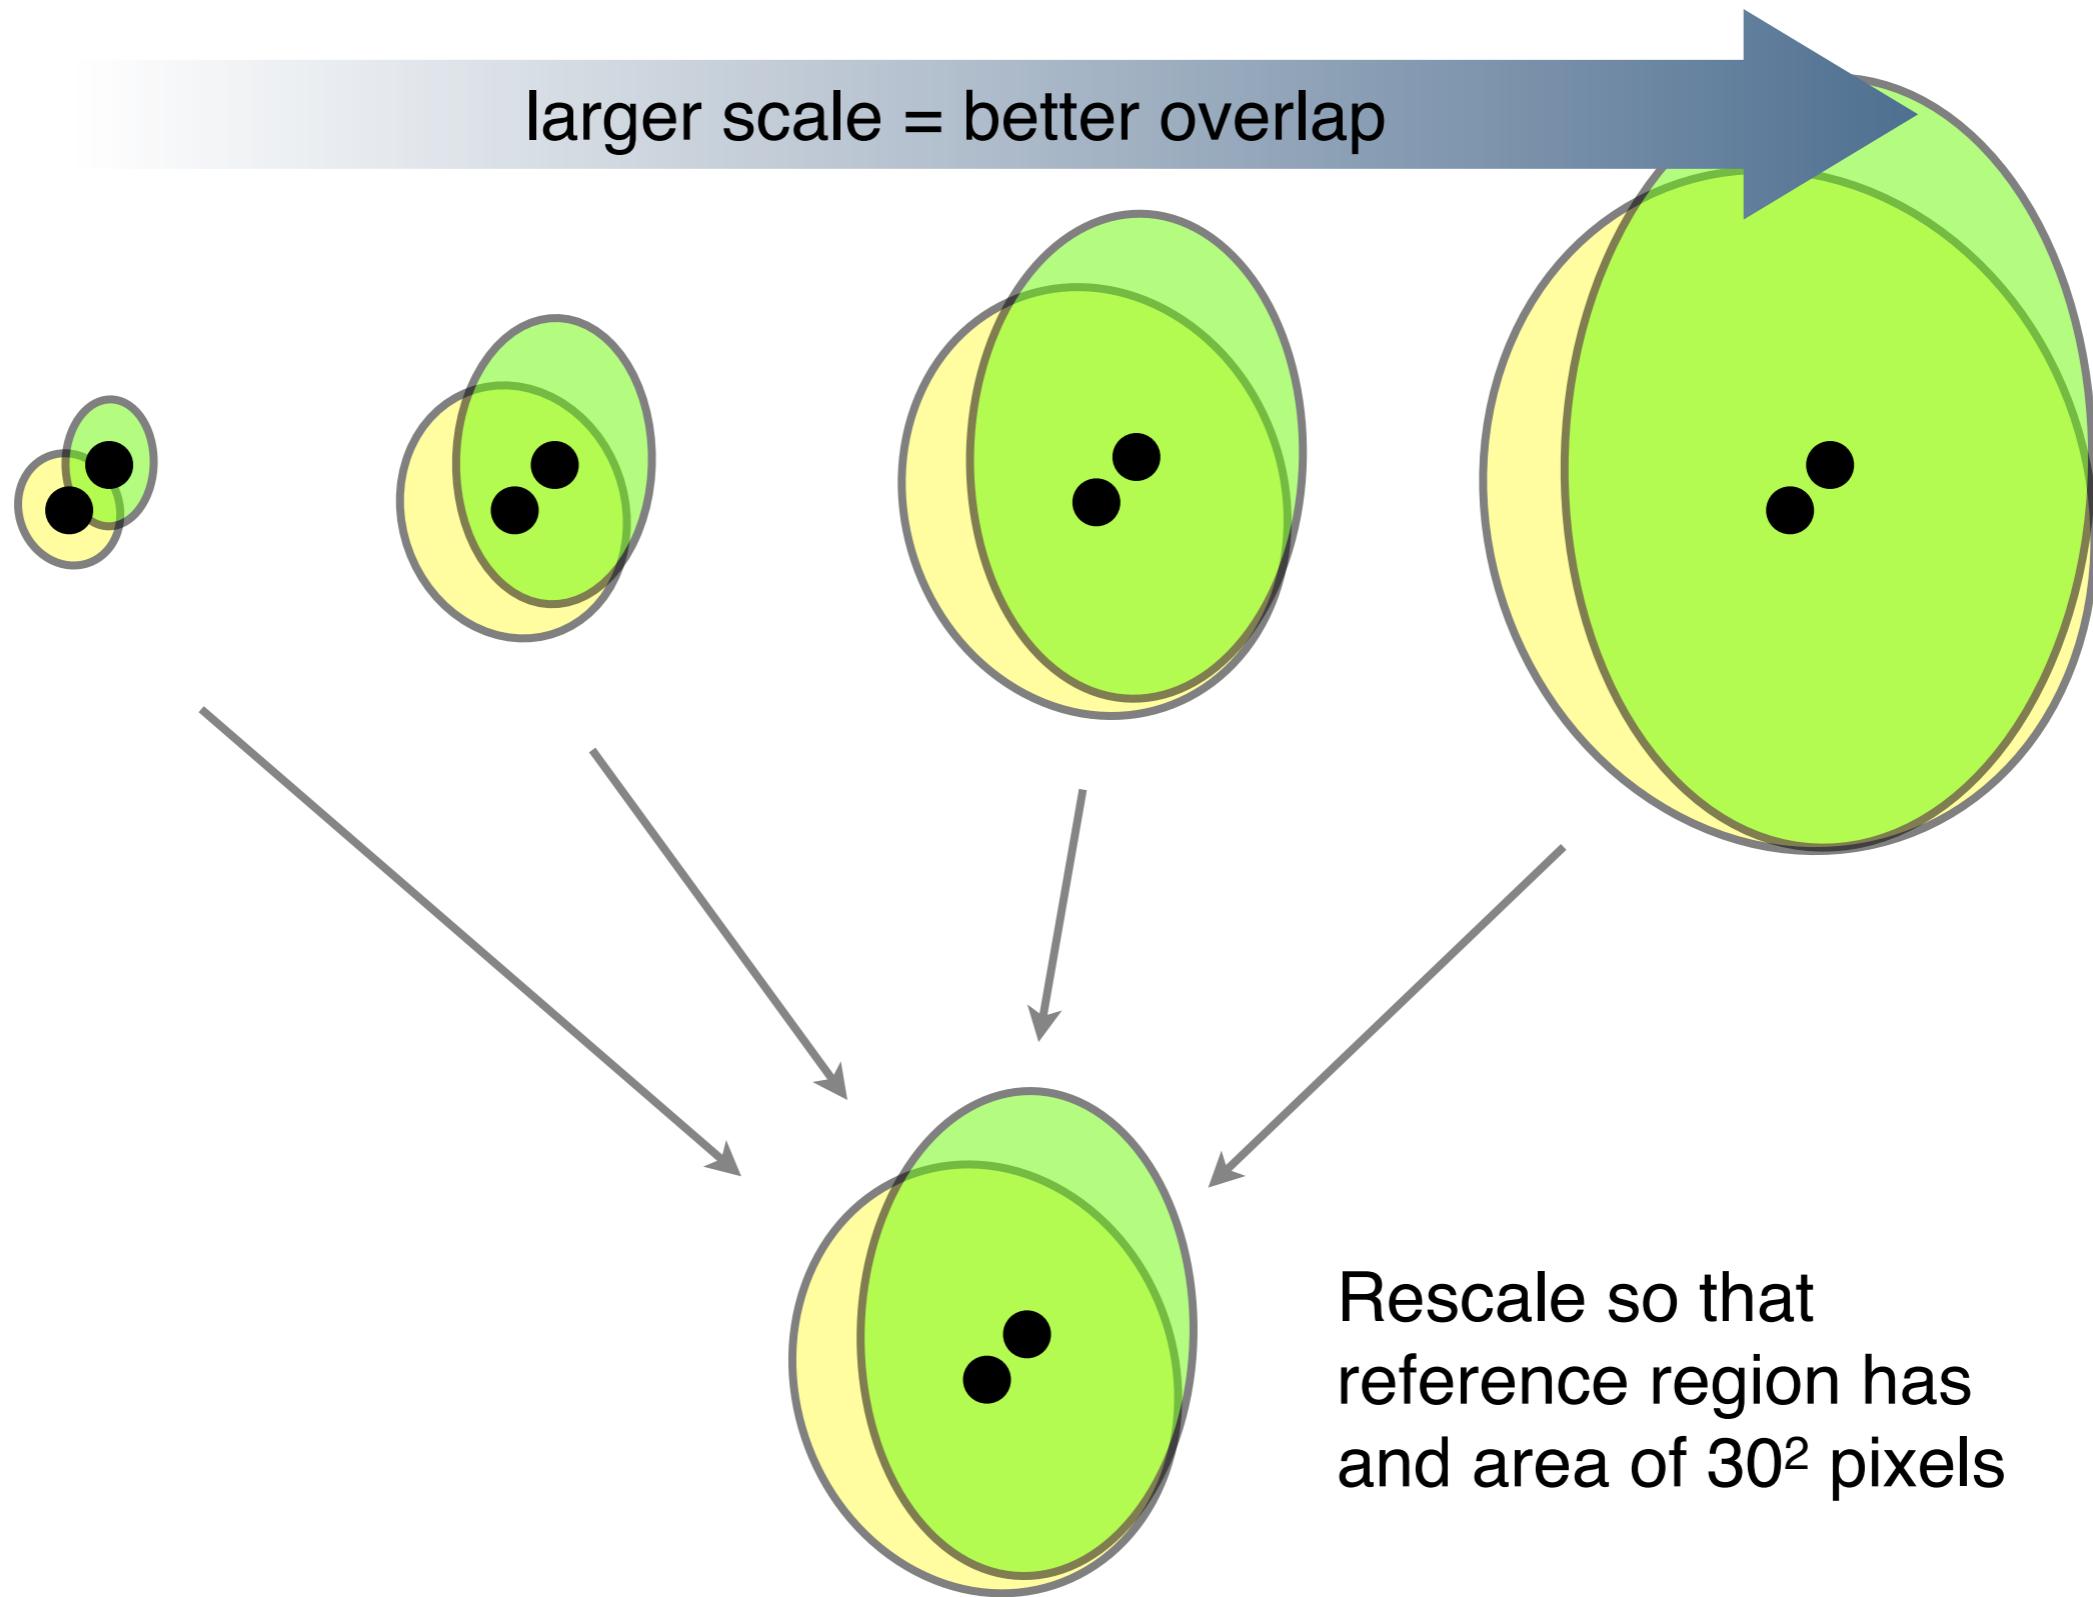






Normalised region overlap

Intuition



Normalised region overlap

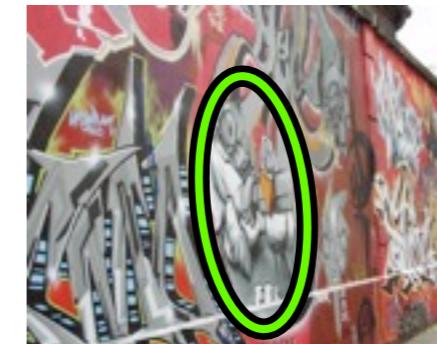
Formal definition

1. Detect



$$R_a$$

$$\xleftarrow[H]{\text{homography}}$$



$$R_b$$

2. Warp

$$R_a$$

$$HR_b$$

3. Normalise

$$sR_a$$

$$sHR_b$$

$$s = |R_a|/30^2$$

4. Intersection
over union

$$\text{overlap}(a, b) = \frac{|sR_a \cap sHR_b|}{|sR_a \cup sHR_b|}$$

Detector repeatability

Formal definition

1. Find features in common area



$$\{R_a : a \in A\}$$

$\xleftarrow[H]{\text{homography}}$



$$\{R_b : b \in B\}$$

2. Threshold the overlap score

$$s_{ab} = \begin{cases} \text{overlap}(a, b), & \text{overlap}(a, b) \geq 1 - \epsilon_o \\ -\infty, & \text{otherwise.} \end{cases}$$

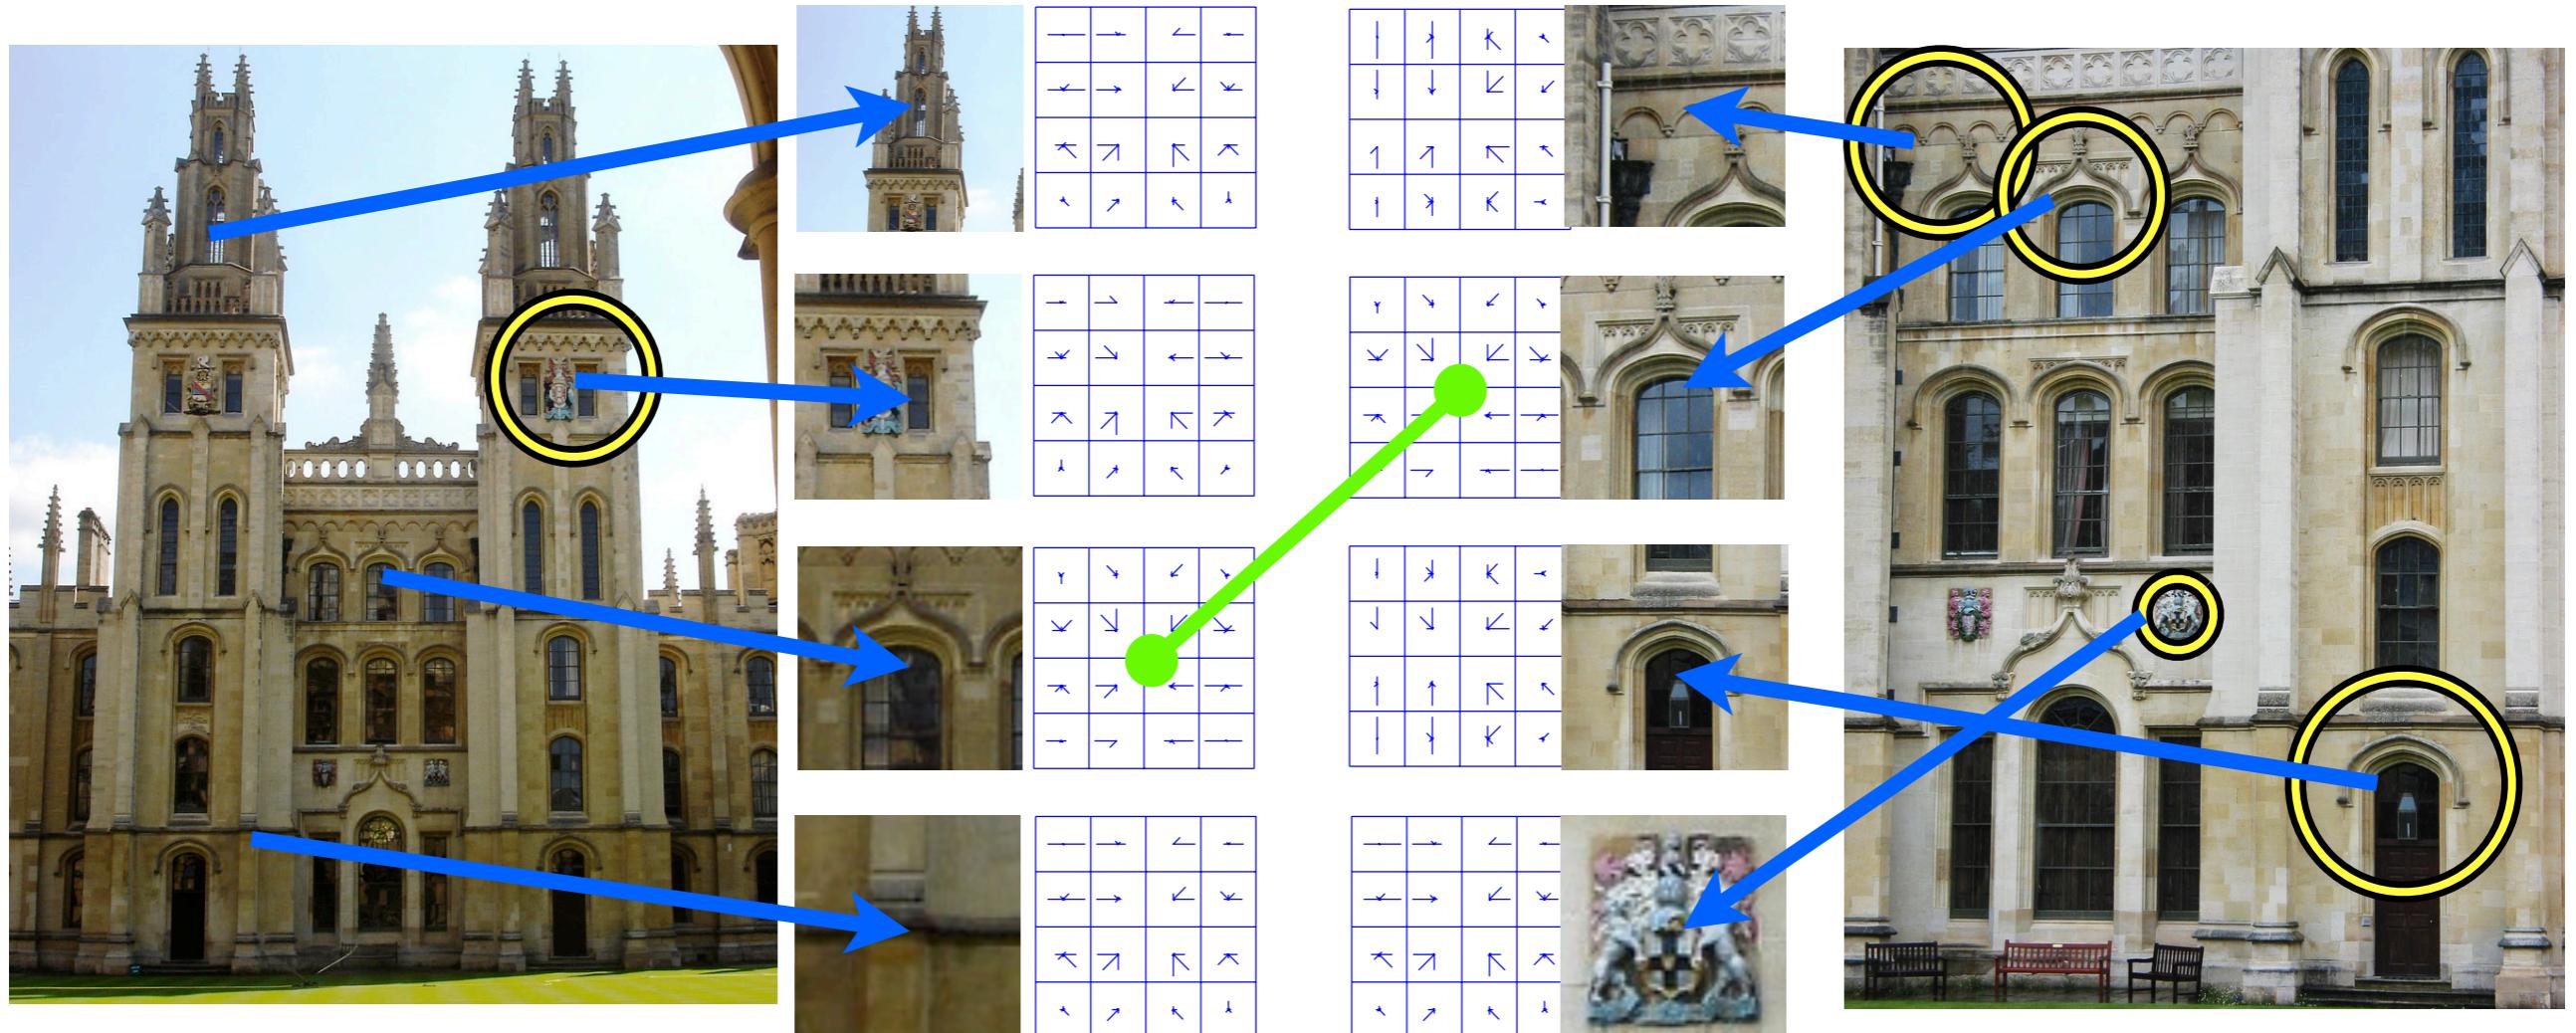
3. Find geometric matches

$$\mathcal{M}^* = \max_{\mathcal{M} \text{ bipartite}} \sum_{(a,b) \in \mathcal{M}} s_{ab} \quad \xleftarrow{\text{greedy approximation}}$$

$$\text{repeatability}(A, B) = \frac{|\mathcal{M}^*|}{\min\{|A|, |B|\}}$$

Descriptor matching score

Intuition



- In addition to being stable, features must be **visually distinctive**
- **Descriptor matching score**
 - similar to repeatability
 - but matches are constructed by comparing descriptors

Descriptor matching score

Formal definition

1. Find features in common area



$$\{R_a : a \in A\}$$



$$\{R_b : b \in B\}$$

2. Descriptor distances

$$\{\mathbf{d}_a : a \in A\}$$

$$\{\mathbf{d}_b : b \in B\} \quad d_{ab} = \|\mathbf{d}_a - \mathbf{d}_b\|_2$$

3. Descriptor matches

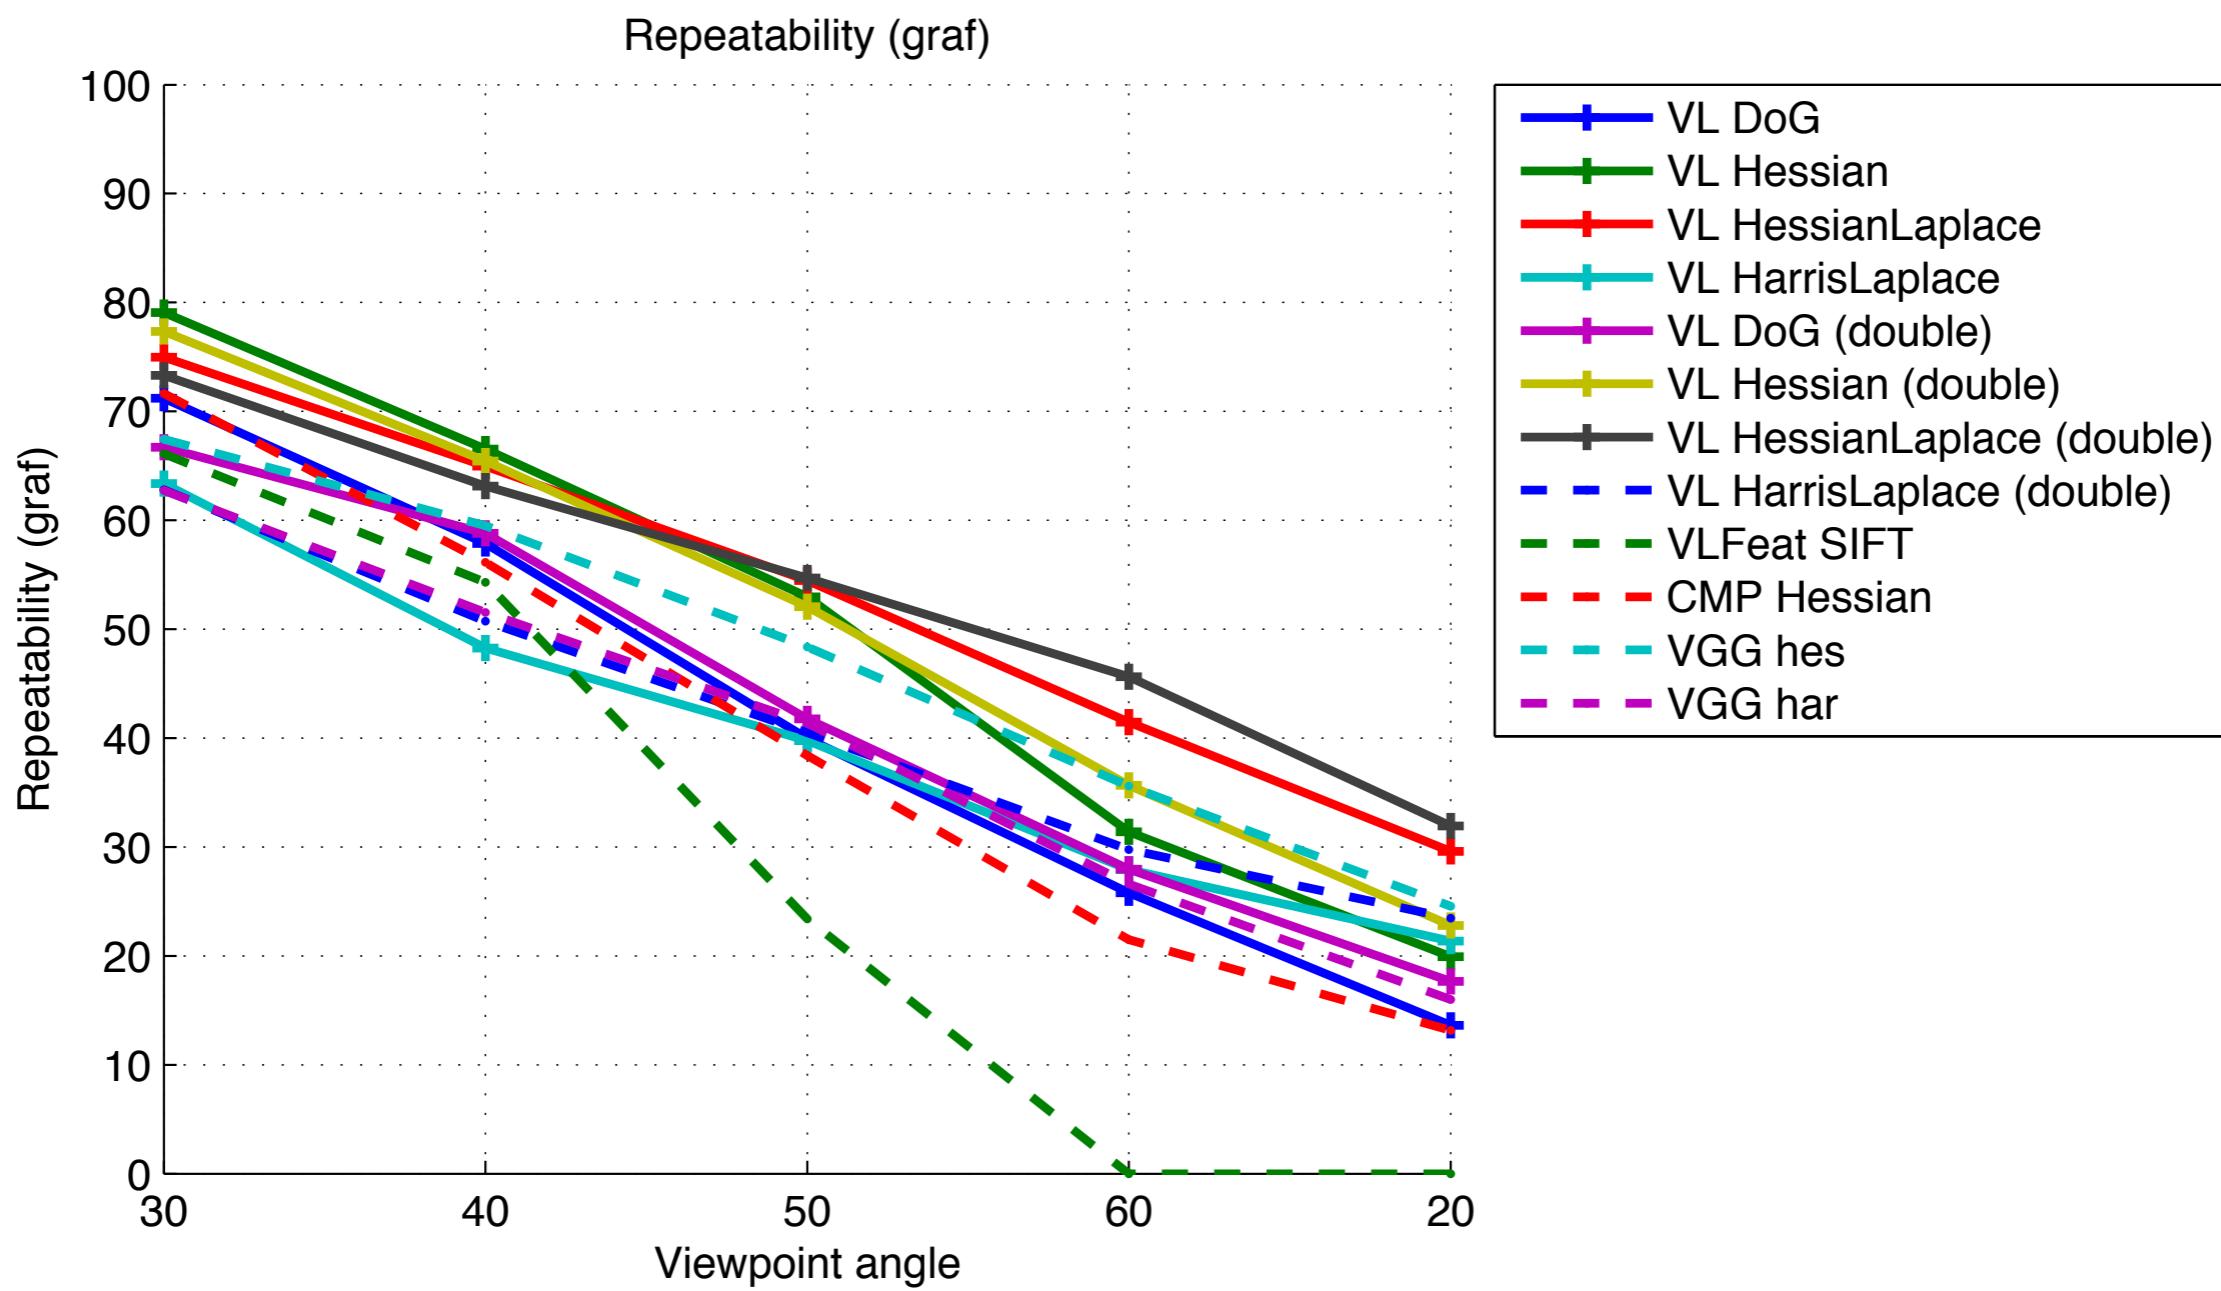
$$\mathcal{M}_d^* = \min_{\mathcal{M} \text{ bipartite}} \sum_{(a,b) \in \mathcal{M}} d_{ab}$$

4. Geometric matches (as before)

$$\mathcal{M}^* = \max_{\mathcal{M} \text{ bipartite}} \sum_{(a,b) \in \mathcal{M}} s_{ab}$$

$$\text{match-score}(A, B) = \frac{|\mathcal{M}^* \cap \mathcal{M}_d^*|}{\min\{|A|, |B|\}}$$

Example of a repeatability graph



Indirect evaluation

Repeatability and matching score

Data: affine covariant testbed

Direct evaluation

Image retrieval

Data: oxford 5k

Software

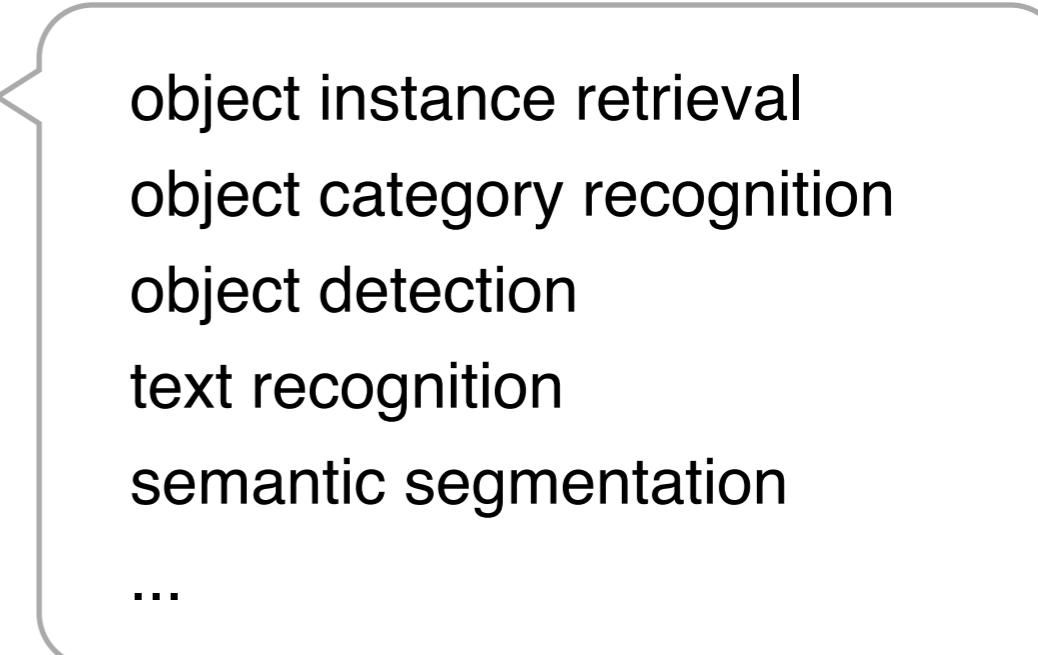
VLBenchmarks

Indirect evaluation

- **Indirect evaluation**
 - a “synthetic” performance measure in a “synthetic” setting
- **The good**
 - independent of a specific application / implementation
 - allow to evaluate single components, e.g.
 - repeatability of detector
 - matching score of descriptor
- **The bad**
 - difficult to design well
 - unclear correlation to the performance in applications

Direct evaluation

- **Direct evaluation**
 - performance of a real system using a feature
- **The good**
 - tied to the “real” performance of the feature
- **The bad**
 - tied to one *application*
 - worse, tied to one *implementation*
 - difficult to evaluate single aspects of a feature
- In the follow up we will focus on **object instance retrieval**



- object instance retrieval
- object category recognition
- object detection
- text recognition
- semantic segmentation

...

Image retrieval

Used to evaluate features

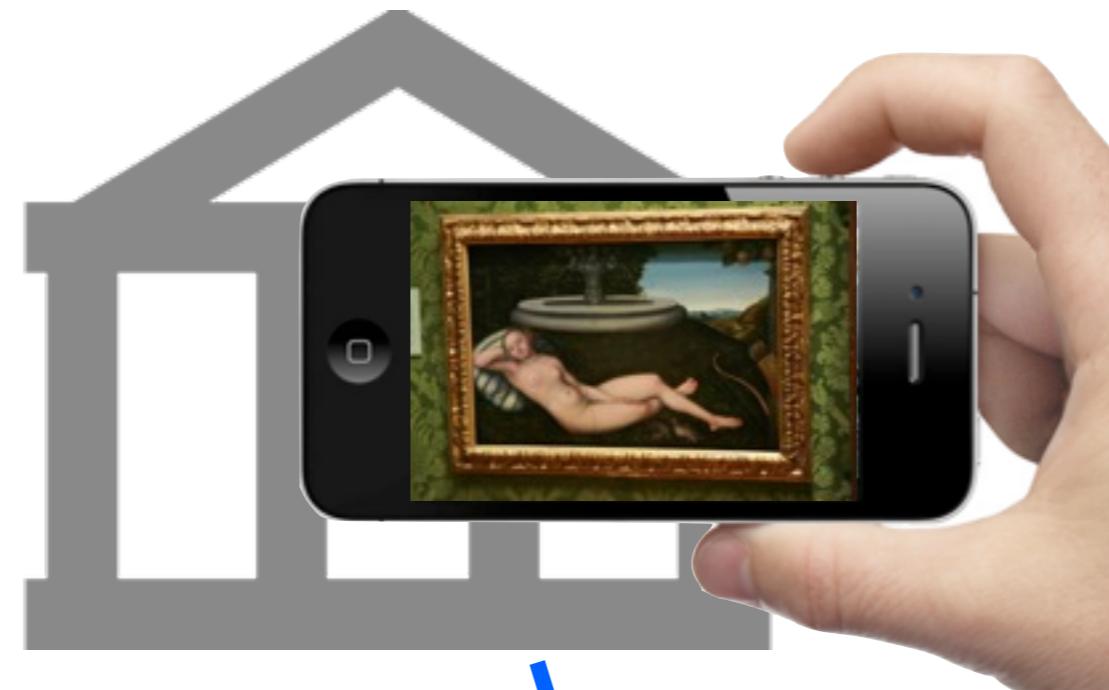


Image retrieval pipeline

Represent images as bags of features

input image



detector



descriptor



$$\{\mathbf{f}_1, \dots, \mathbf{f}_n\}$$

$$\{\mathbf{d}_1, \dots, \mathbf{d}_n\}$$

Harris (Laplace)

SIFT

Hessian (Laplace)

LIOP

DoG

BRIEF

MSER

Jets

Harris Affine

...

Hessian Affine

....

Image retrieval pipeline

Step 1: find neighbours of each query descriptor

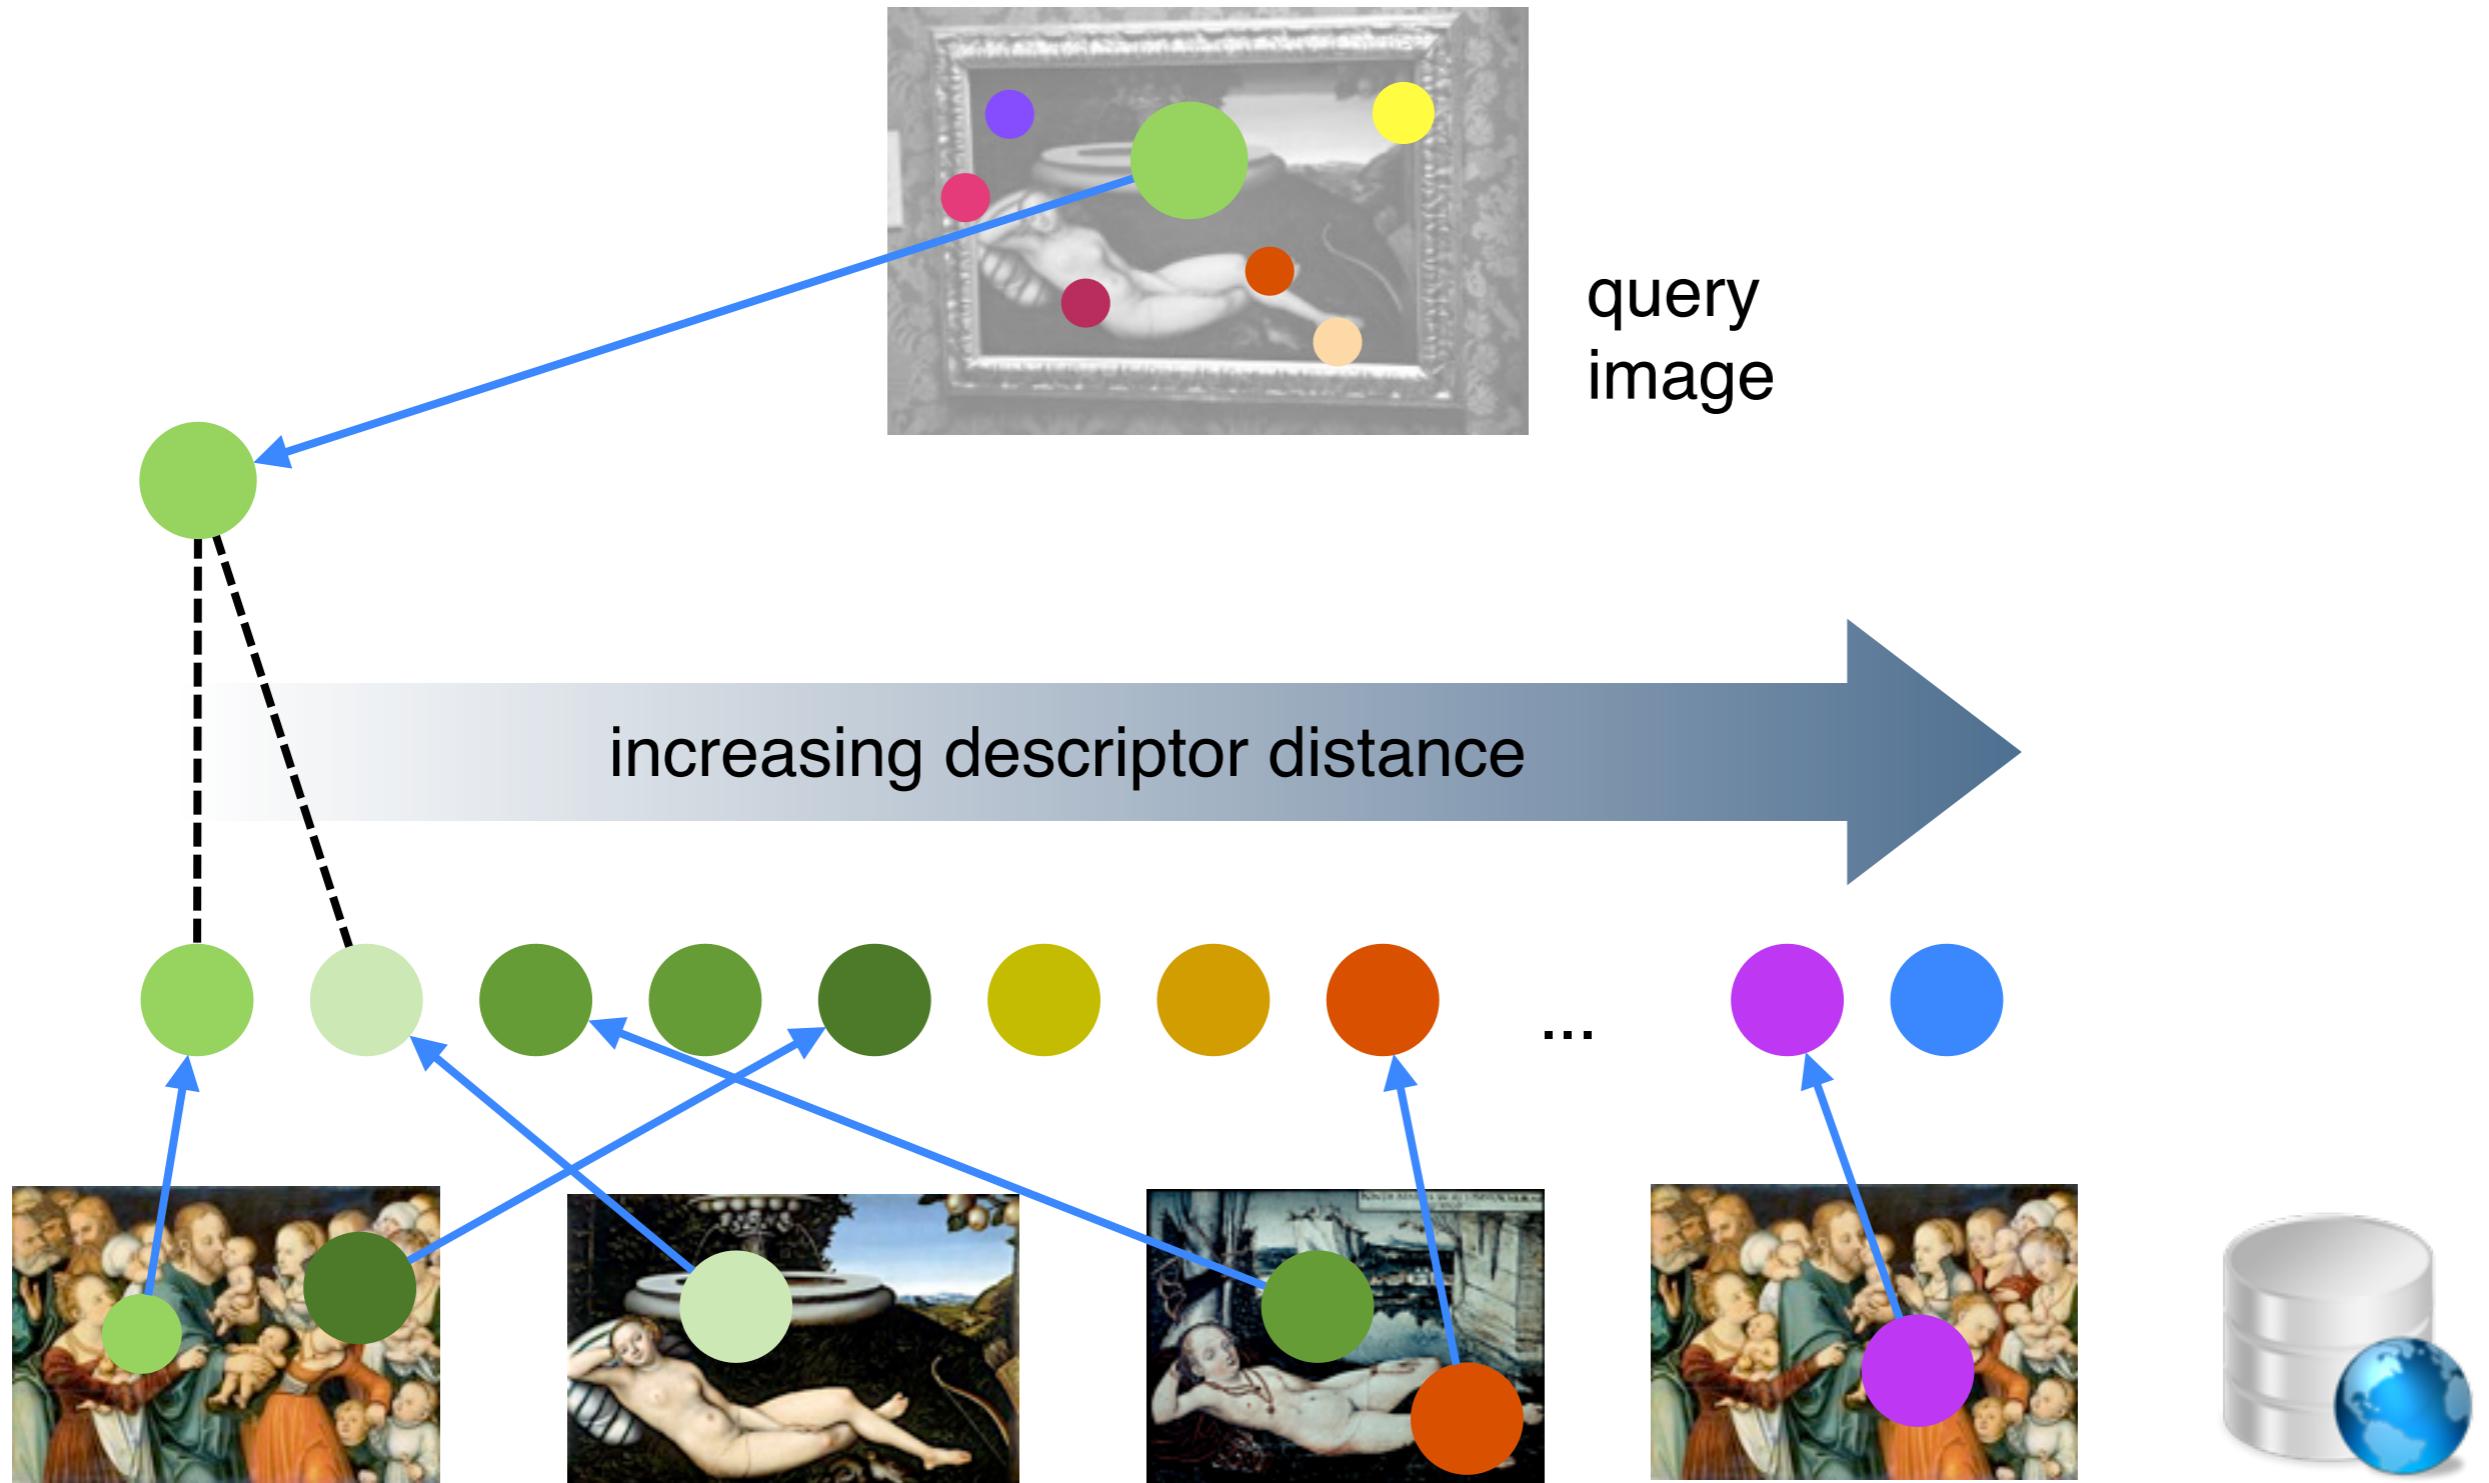


Image retrieval pipeline

Step 2: each query descriptor casts a vote for each DB image

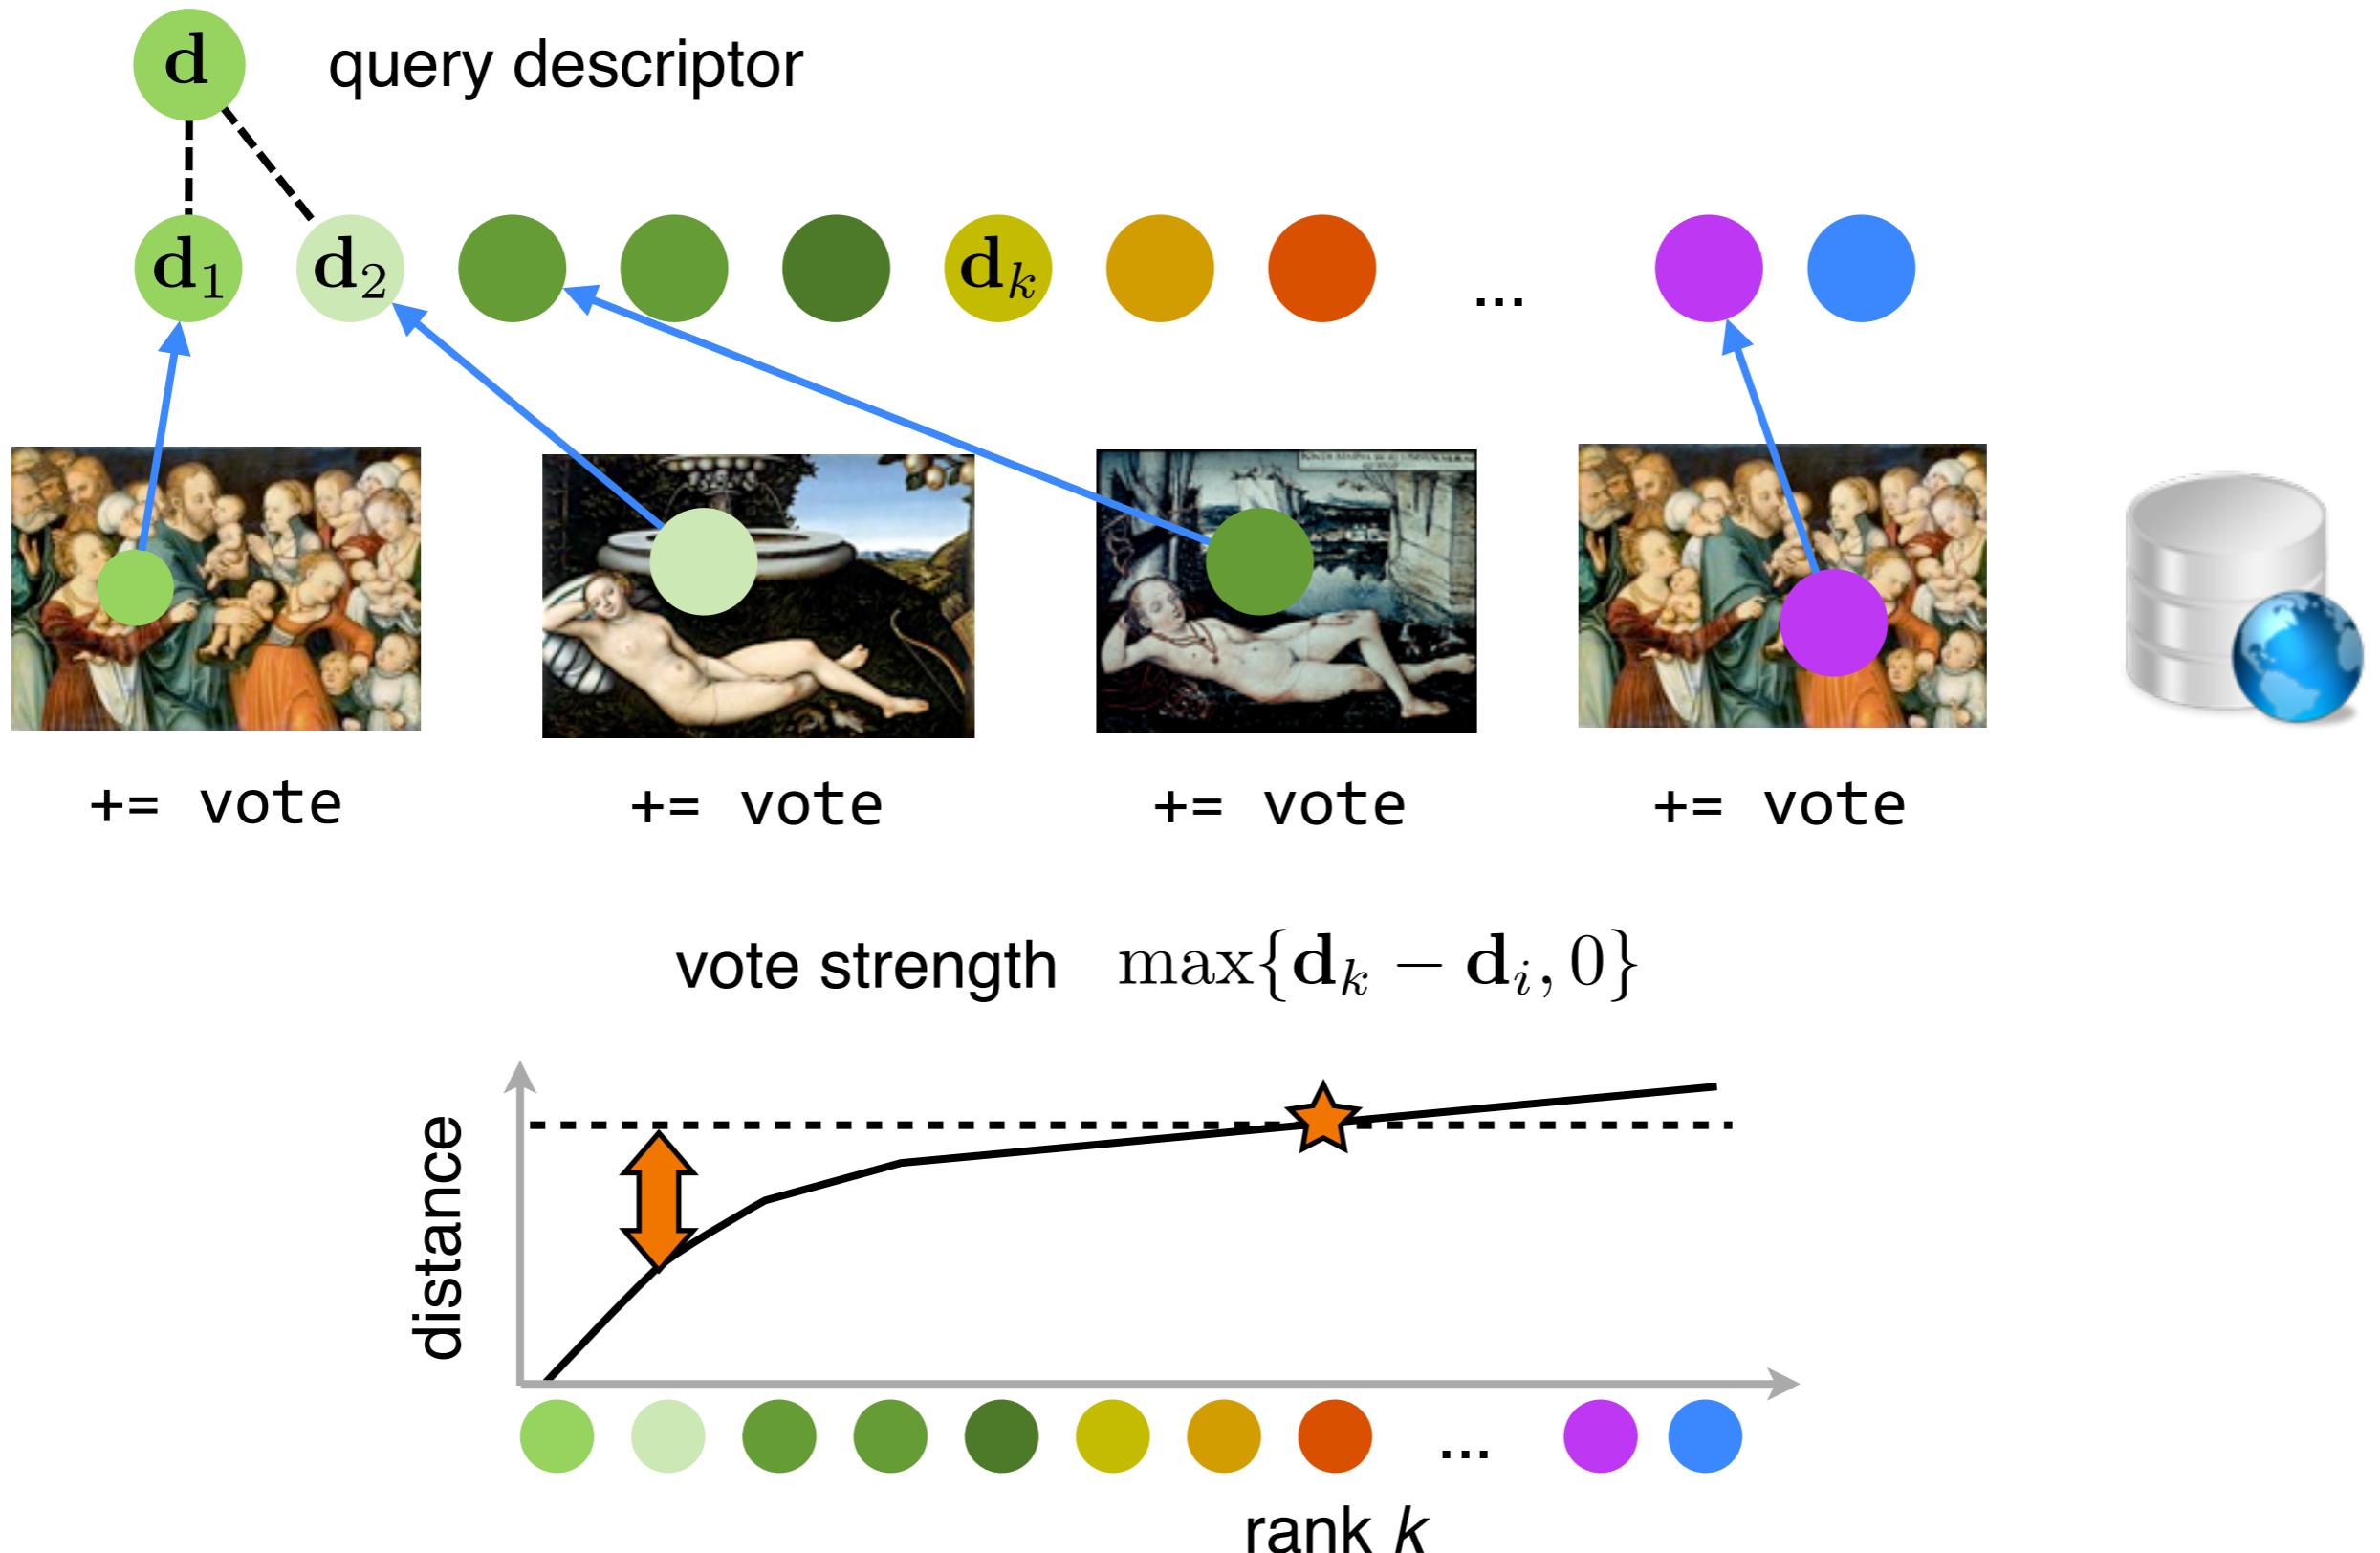
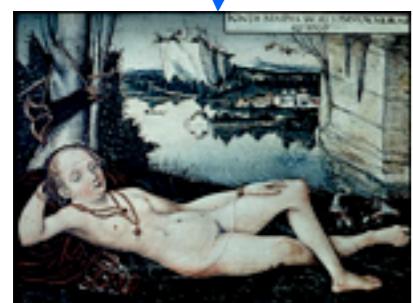


Image retrieval pipeline

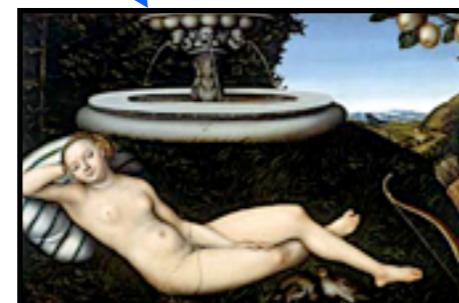
Step 3: sort DB images by decreasing total votes



query image



✗

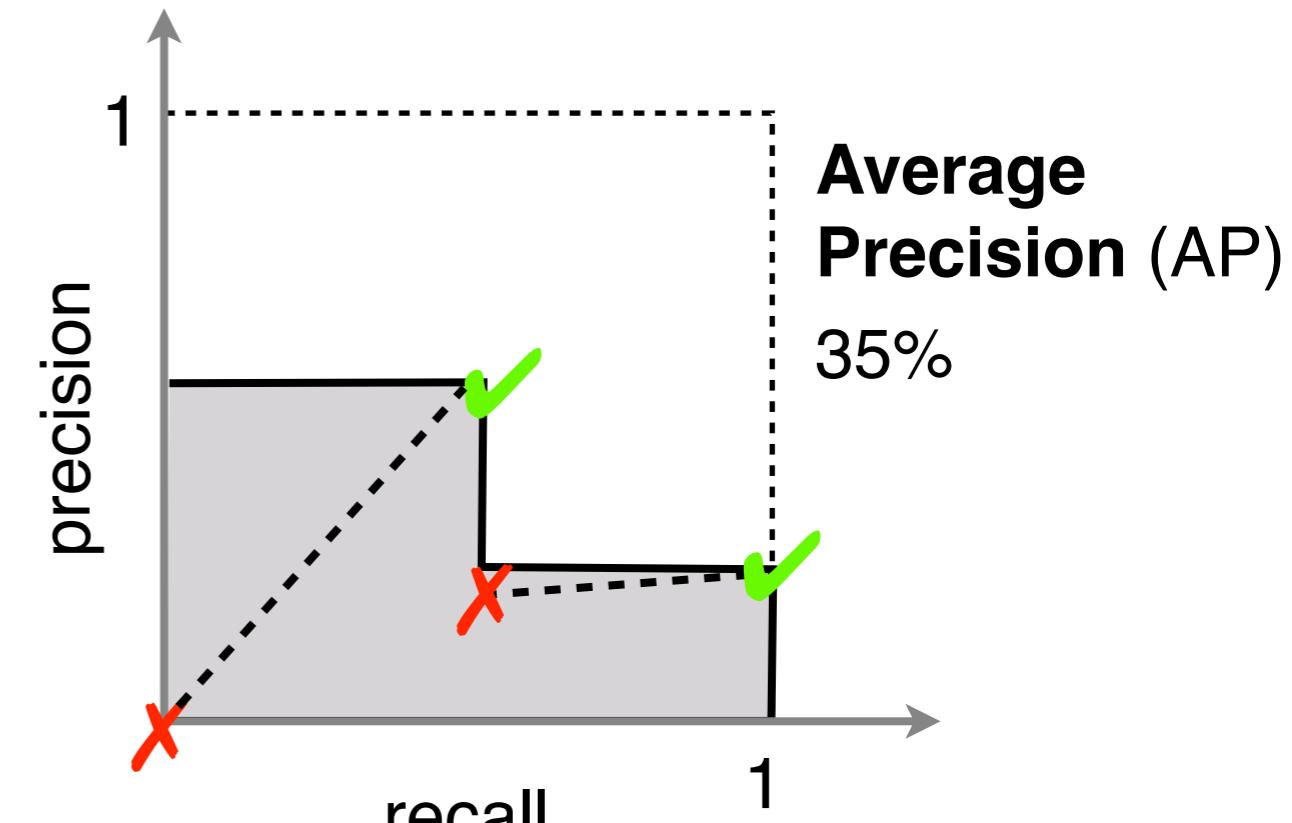


✓



✗

decreasing total votes



...



Image retrieval pipeline

Step 4: Overall performance score

query



retrieval results



✗



✓



✗

AP

35%



✓

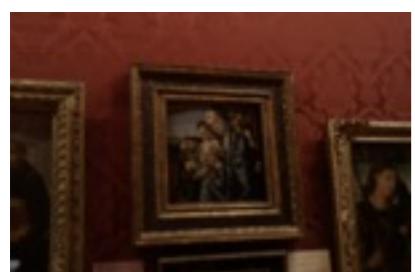


✗



✗

100%



✓



✗



✓

75%

...

...

...

...

Mean Average Precision (mAP) 53%

Oxford 5K data

A retrieval benchmark dataset

Query



Retrieved Images



...

- ~ 5K images of Oxford
 - For each of 58 queries
 - about XX matching images
 - about XX confounders images
- Larger datasets are possible, but slow for extensive evaluation
- Relative ranking of features seems to be representative

Indirect evaluation

Repeatability and matching score

Data: affine covariant testbed

Direct evaluation

Image retrieval

Data: oxford 5k

Software

VLBenchmarks

VLBenchmarks

A new easy-to-use benchmarking suite

<http://www.vlfeat.org/benchmarks/index.html>

- A novel MATLAB framework for feature evaluation
 - **Repeatability and matching scores**
 - VGG affine testbed
 - **Image retrieval**
 - Oxford 5K
- **Goodies**
 - Simple to use MATLAB code
 - Automatically download datasets & run evaluations
 - Backward compatible with published results

VLBenchmarks

Obtaining and installing the code

- **Installation**

- Download the latest version
- Unpack the archive
- Launch MATLAB and type

`>> install`

- **Requirements**

- MATLAB R2008a (7.6)
- A C compiler (e.g. Visual Studio, GCC, or Xcode)
- Do not forget to setup MATLAB to use your C compiler

`mex -setup`

Example usage

```
import datasets.*;  
import localFeatures.*;  
import benchmarks.*;
```

choose a detector
(Harris Affine, VGG version)

```
detector = VlFeatSift() ;
```

choose a dataset
(graffiti sequence)

```
dataset = VggAffineDataset('category','graf') ;
```

```
benchmark = RepeatabilityBenchmark() ;
```

choose test
(detector repeatability)

```
repeatability = ...
```

```
benchmark.testFeatureExtractor(detector, ...  
                                dataset.getTransformation(2), ...  
                                dataset.getImagePath(1), ...  
                                dataset.getImagePath(2)) ;
```

run the evaluation
(repeatability = 0.66)

Testing on a sequence of images

78

```
import datasets.*; import localFeatures.*; import benchmarks.*;

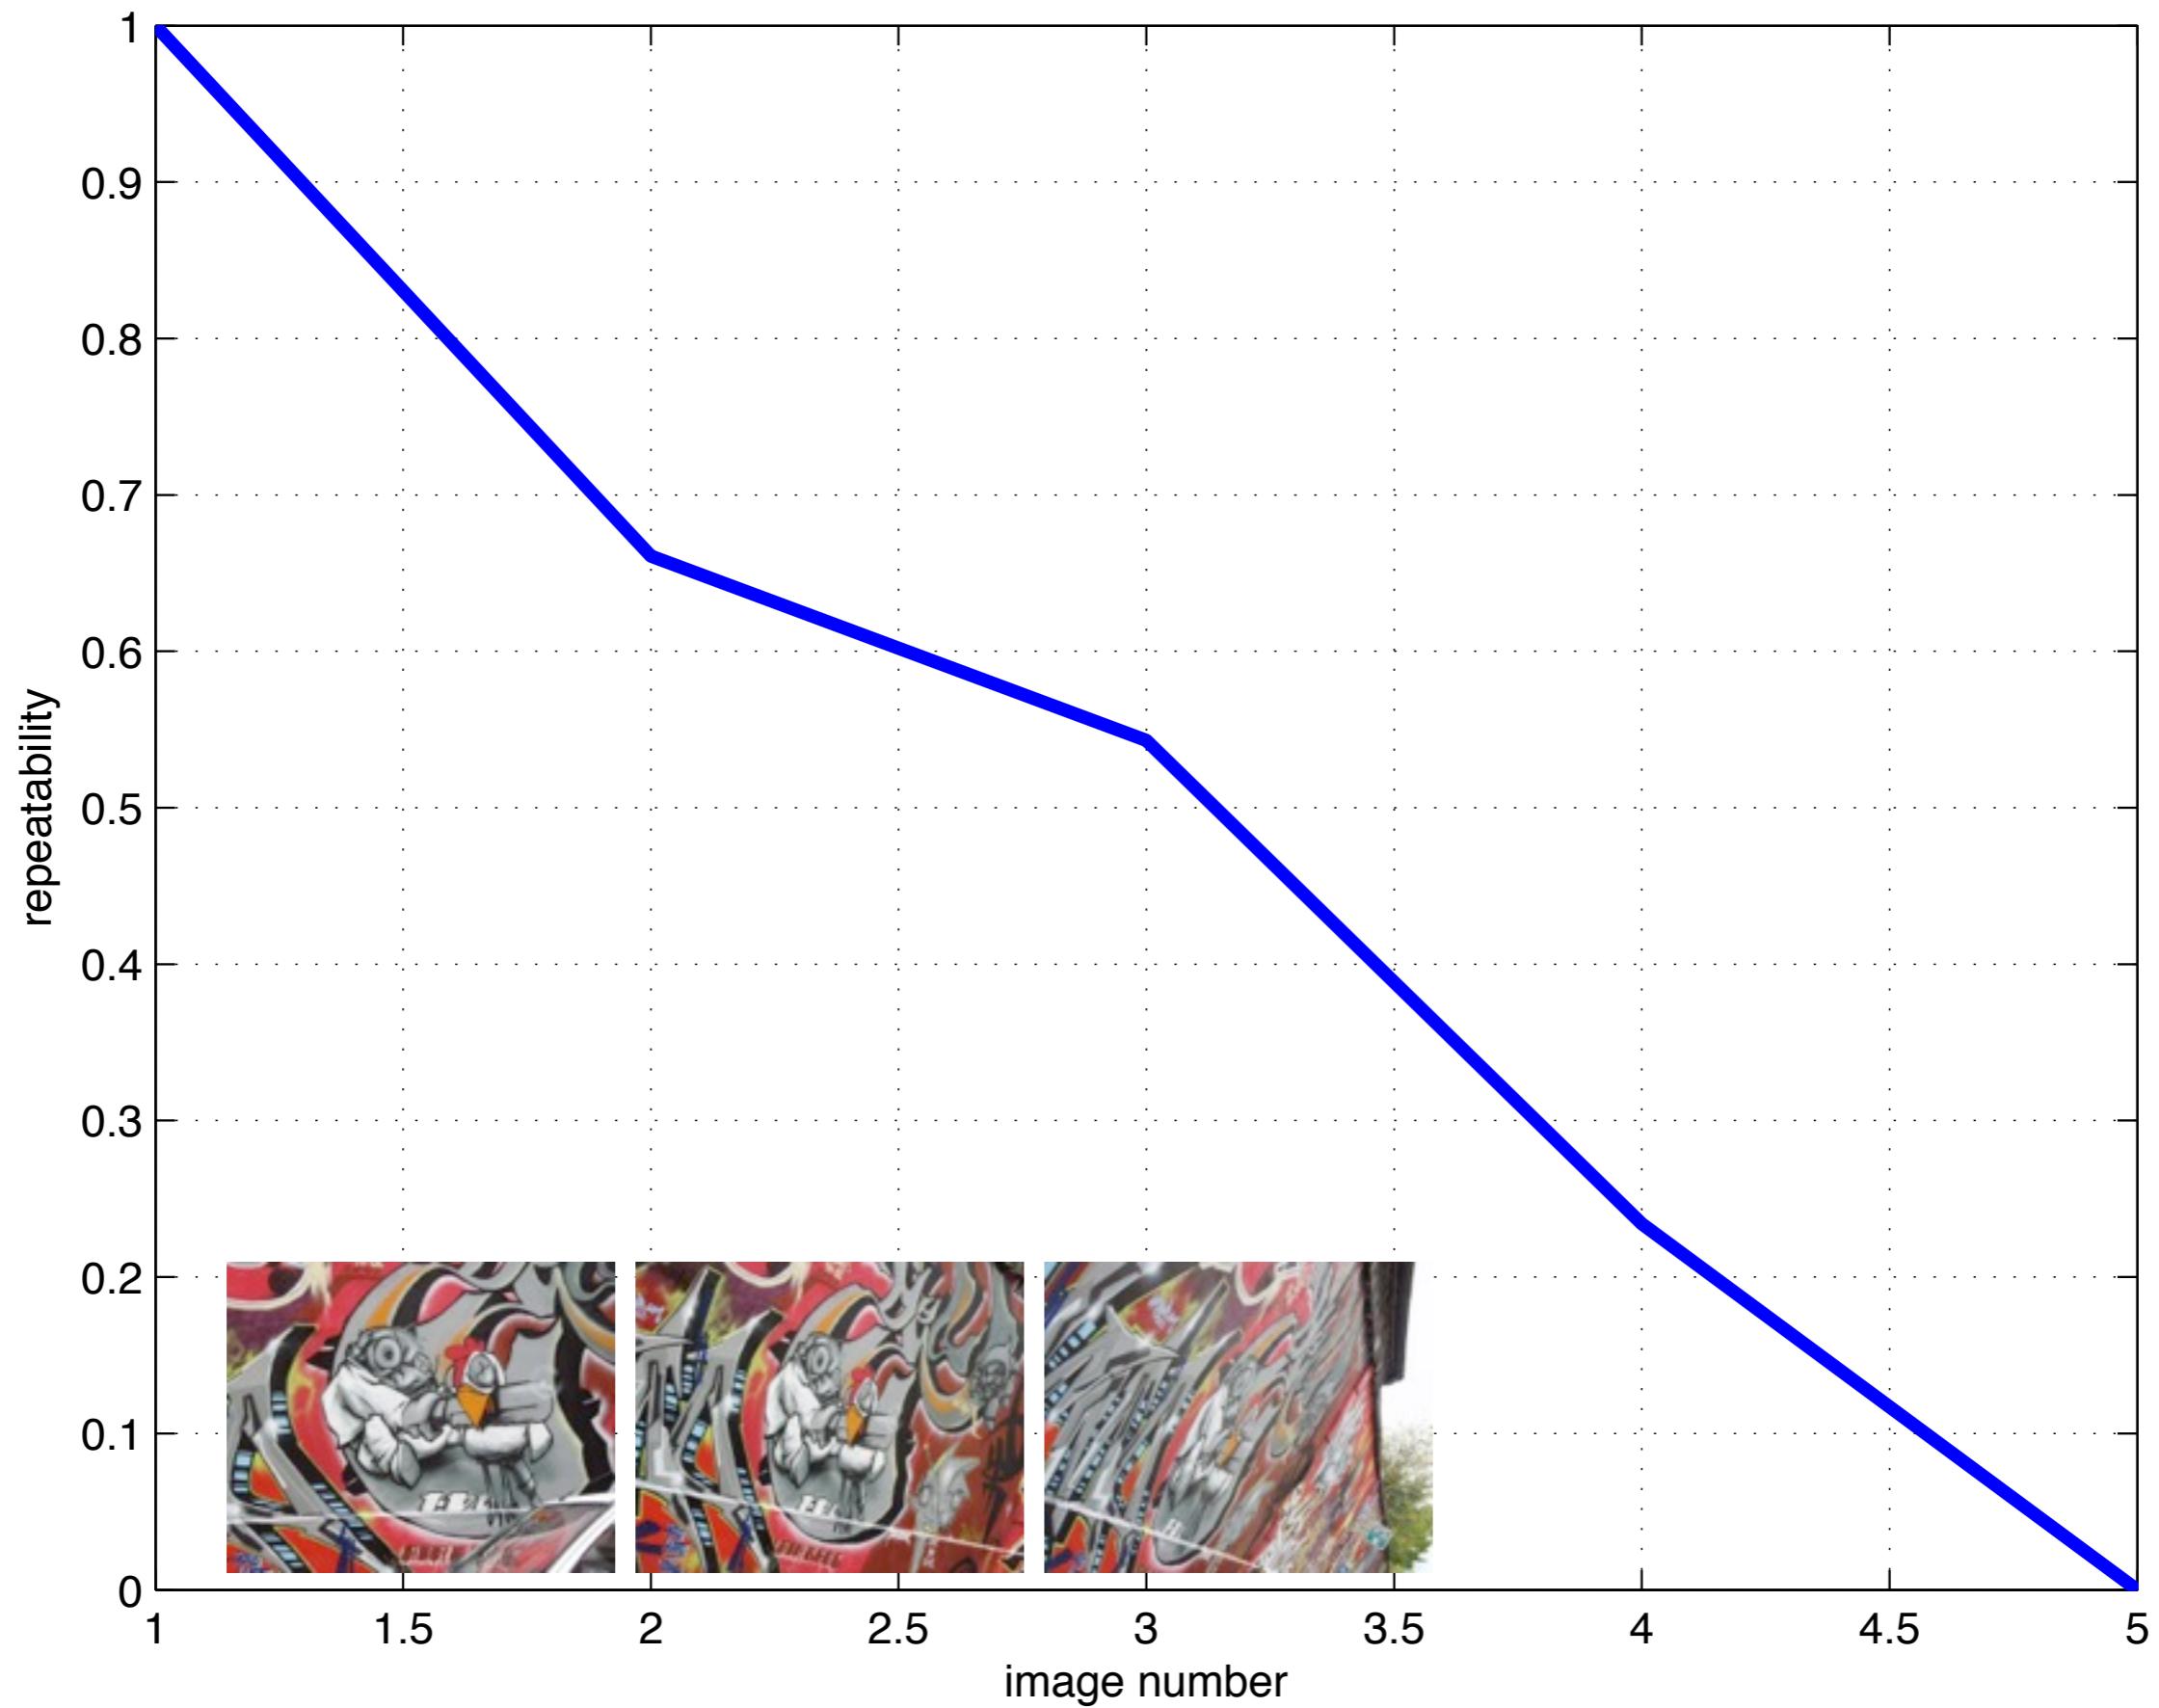
detector = VLFeatSift() ;
dataset = VggAffineDataset('category','graf') ;
benchmark = RepeatabilityBenchmark() ;

for j = 1:5
    repeatability(j) = ...
        benchmark.testFeatureExtractor(detector, ...
            dataset.getTransformation(j), ...
            dataset.getImagePath(1), ...
            dataset.getImagePath(j)) ;

end

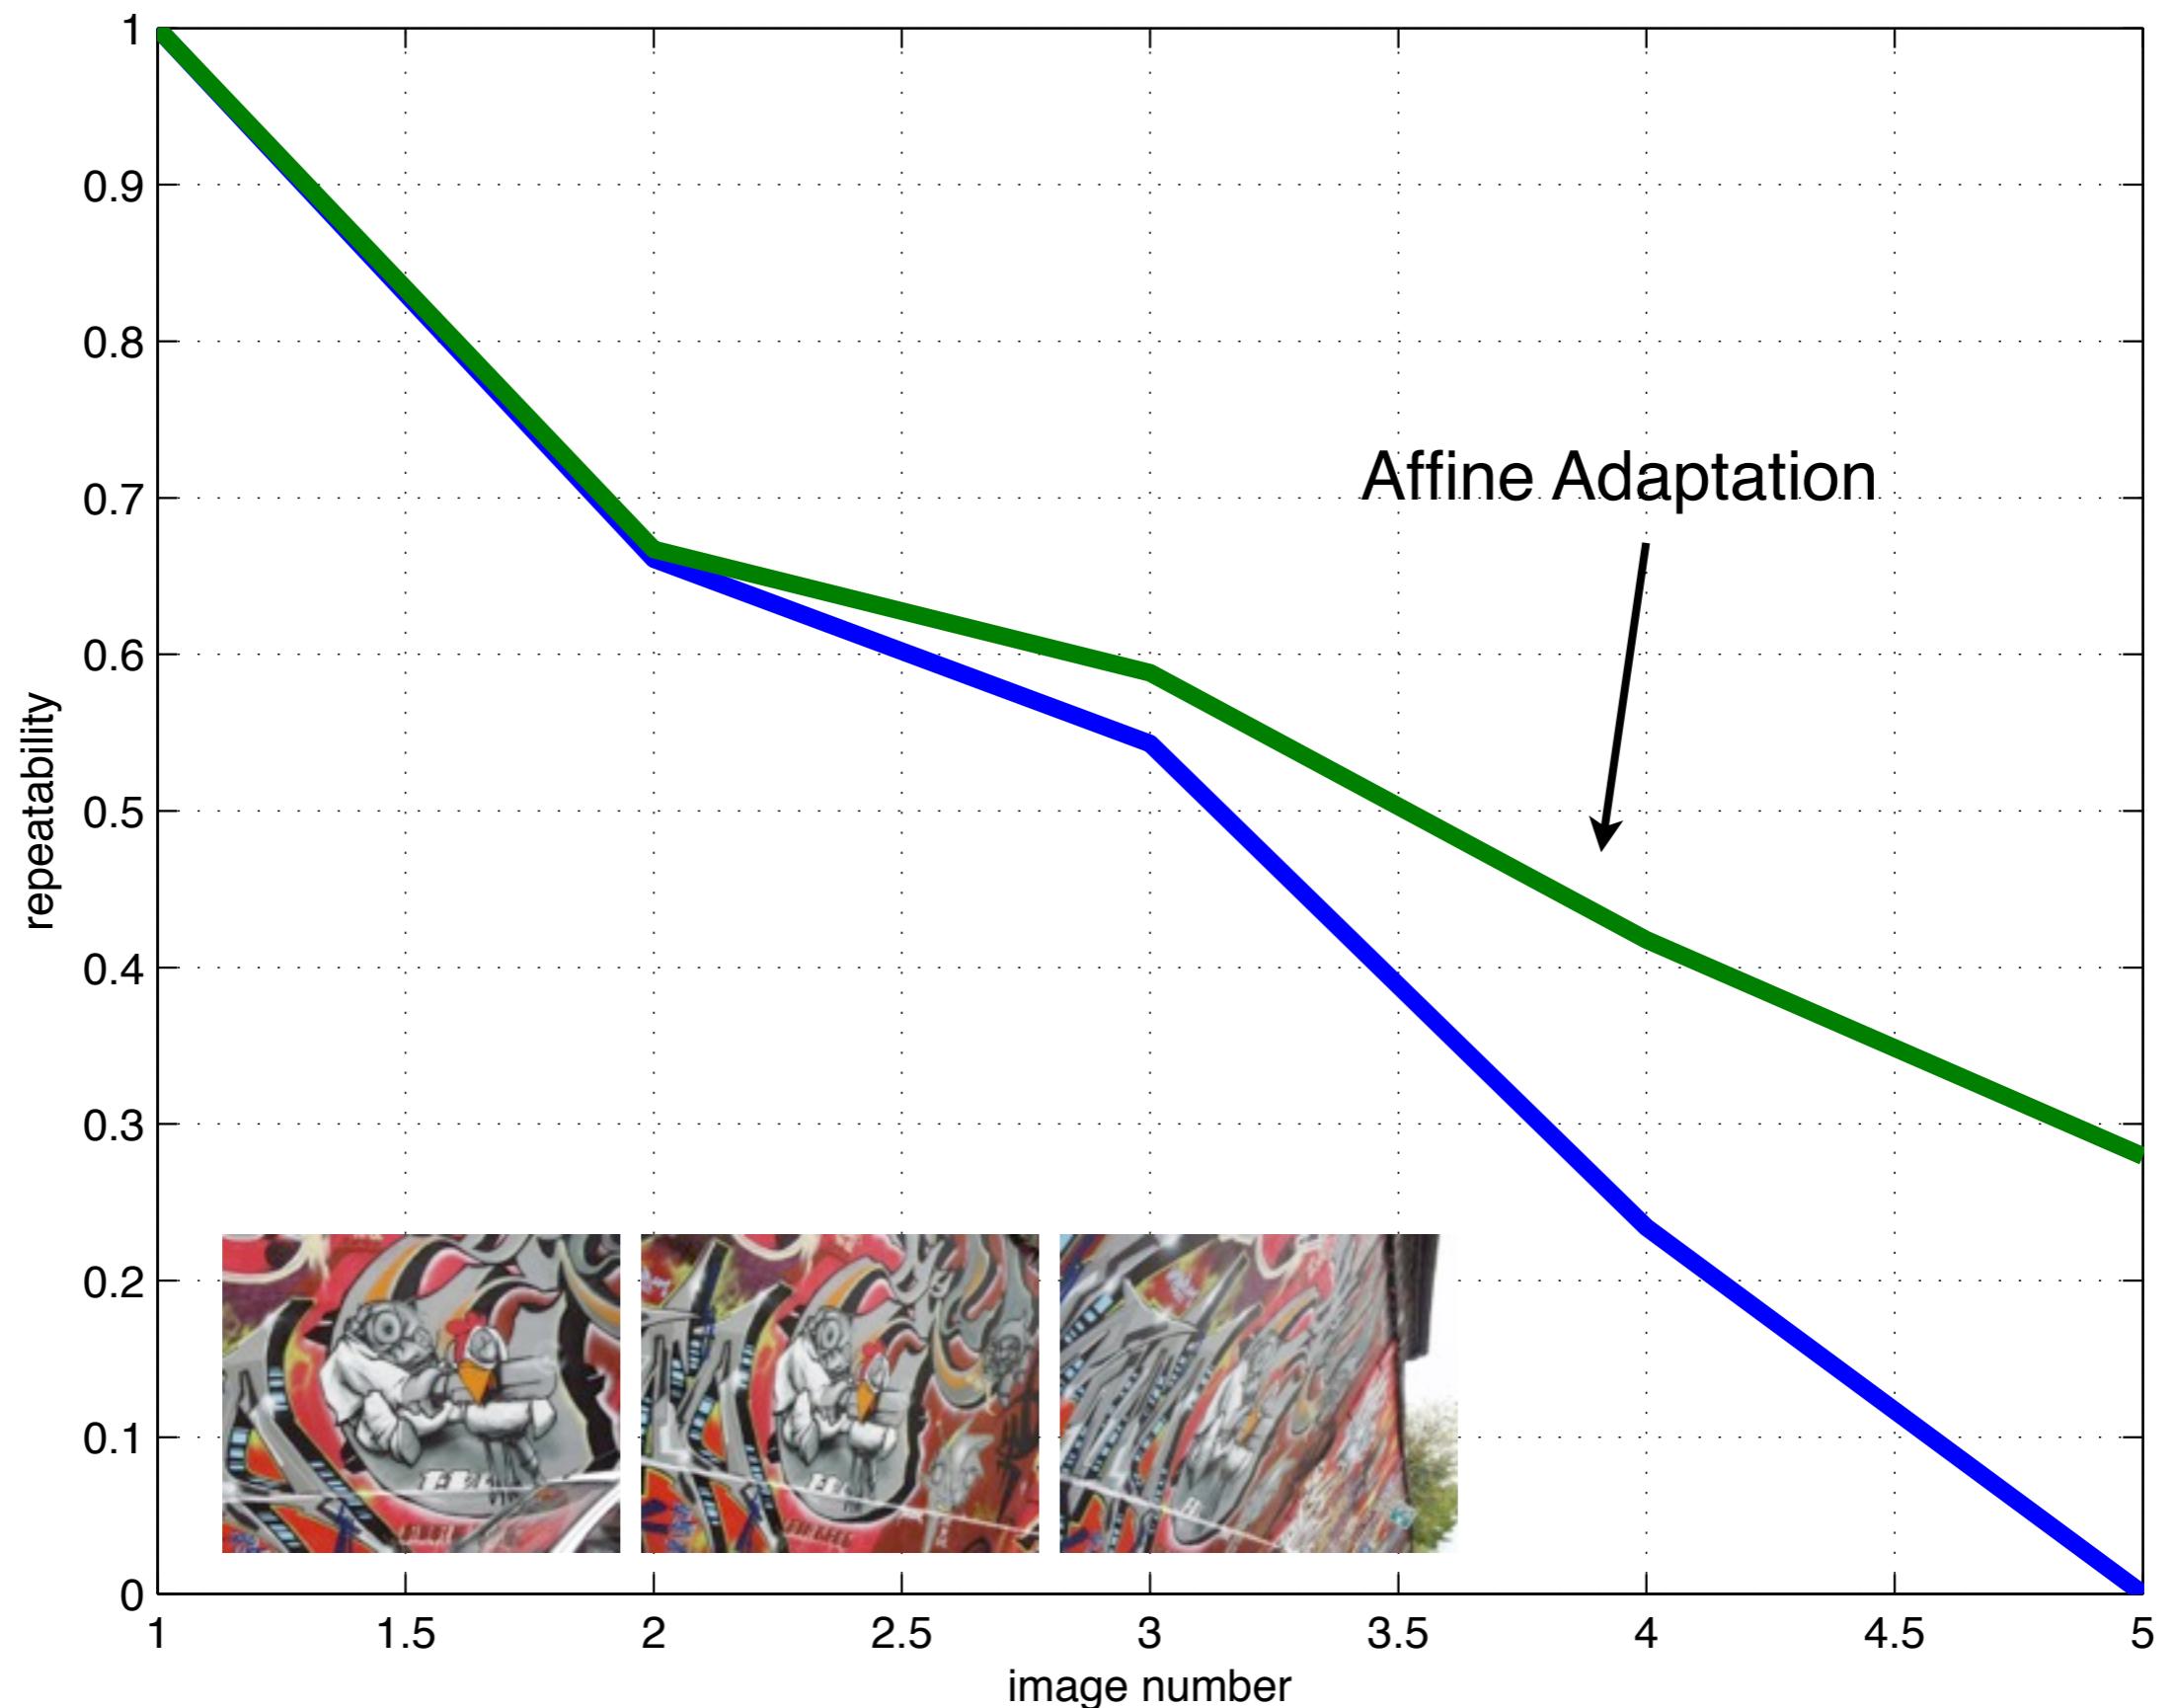
clf ; plot(repeatability, 'linewi
 xlabel('image number') ;
 ylabel('repeatability') ;
 grid on ;
```

Use **parfor** on a cluster!



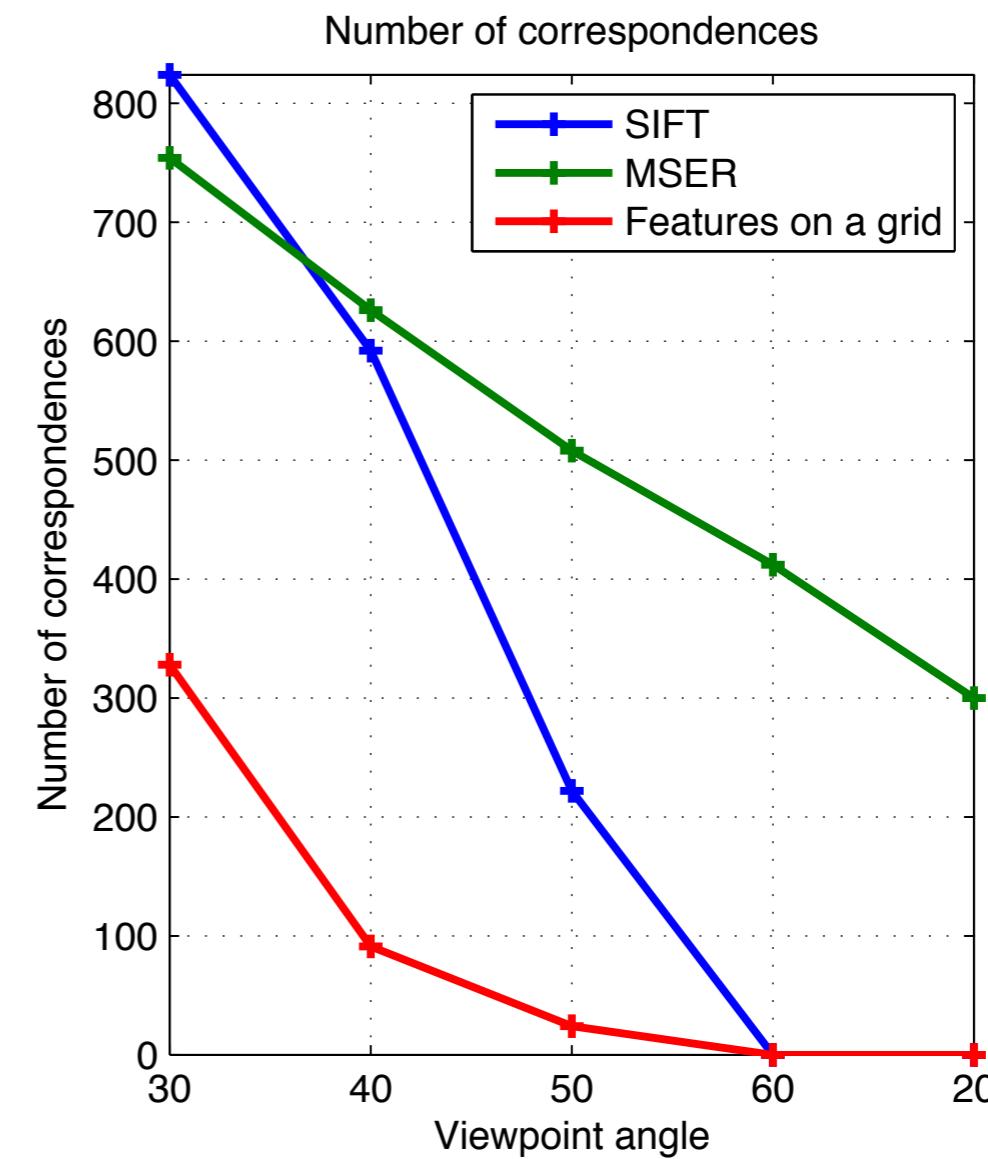
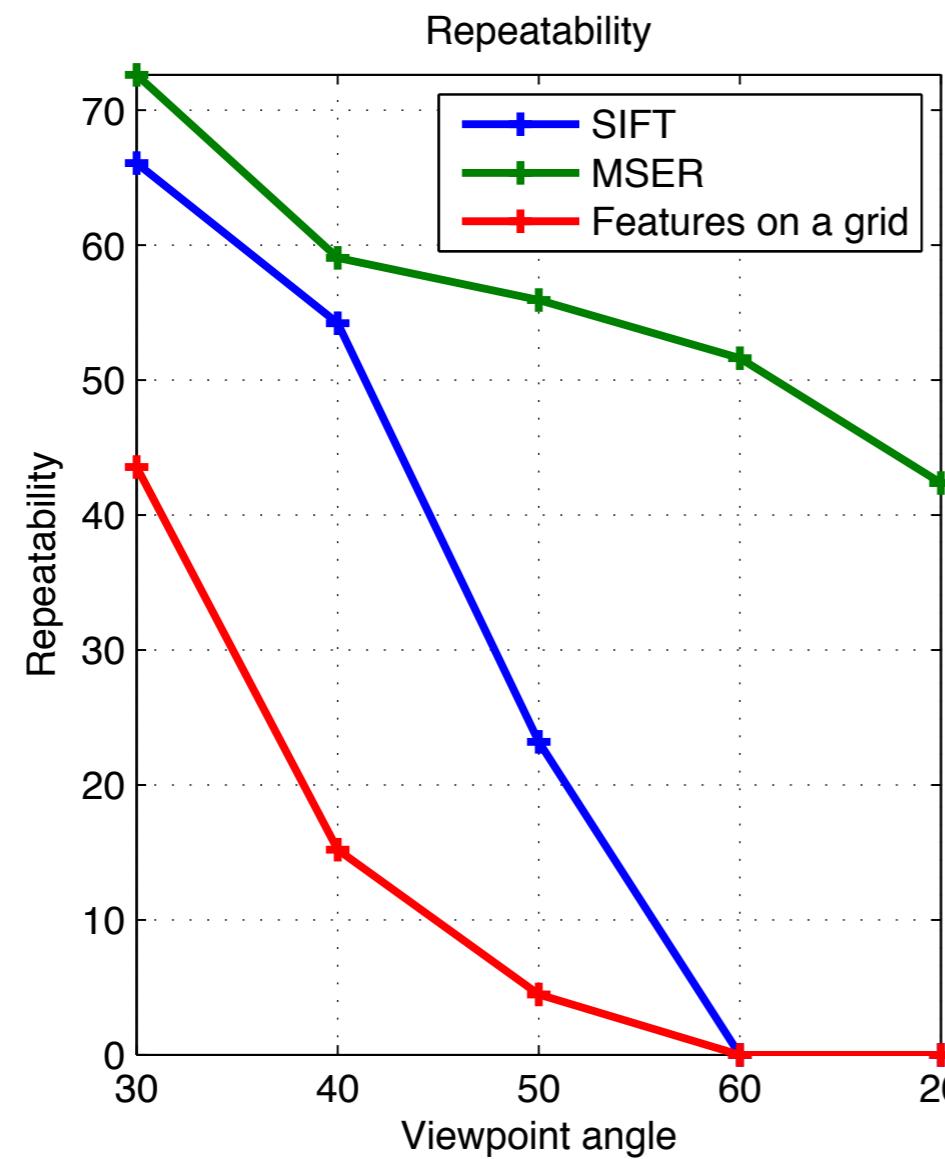
Comparing two features

```
import datasets.*; import localFeatures.*; import benchmarks.*;
detectors{1} = VlFeatSift() ;
detectors{2} = VlFeatCovdet('EstimateAffineShape', true) ;
dataset = VggAffineDataset('category','graf') ;
benchmark = RepeatabilityBenchmark() ;
for d = 1:2
    for j = 1:5
        repeatability(j,d) = ...
            benchmark. testFeatureExtractor(detectors{d}, ...
                dataset.getTransformation(j), ...
                dataset.getImagePath(1), ...
                dataset.getImagePath(j)) ;
    end
end
clf ; plot(repeatability, 'linewidth', 4) ;
xlabel('image number') ;
ylabel('repeatability') ;
grid on ;
```



Example

- Compare the following features
 - SIFT, MSER, and features on a grid
 - on the Graffiti sequence
 - for repeatability and number of correspondence



Backward compatible

- Previously published results can be easily reproduced
 - if interested, try the script **reproduceIjcv05.m**

A Comparison of Affine Region Detectors

K. MIKOLAJCZYK

University of Oxford, OX1 3PJ, Oxford, United Kingdom

km@robots.ox.ac.uk

T. TUYTELAARS

University of Leuven, Kasteelpark Arenberg 10, 3001 Leuven, Belgium

tuytelaa@esat.kuleuven.be

C. SCHMID

INRIA, GRAVIR-CNRS, 655, av. de l'Europe, 38330, Montbonnot, France

schmid@inrialpes.fr

A. ZISSEMAN

University of Oxford, OX1 3PJ, Oxford, United Kingdom

az@robots.ox.ac.uk

J. MATAS

Czech Technical University, Karlovo Namesti 13, 121 35, Prague, Czech Republic

Other useful tricks

- Compare different parameter settings

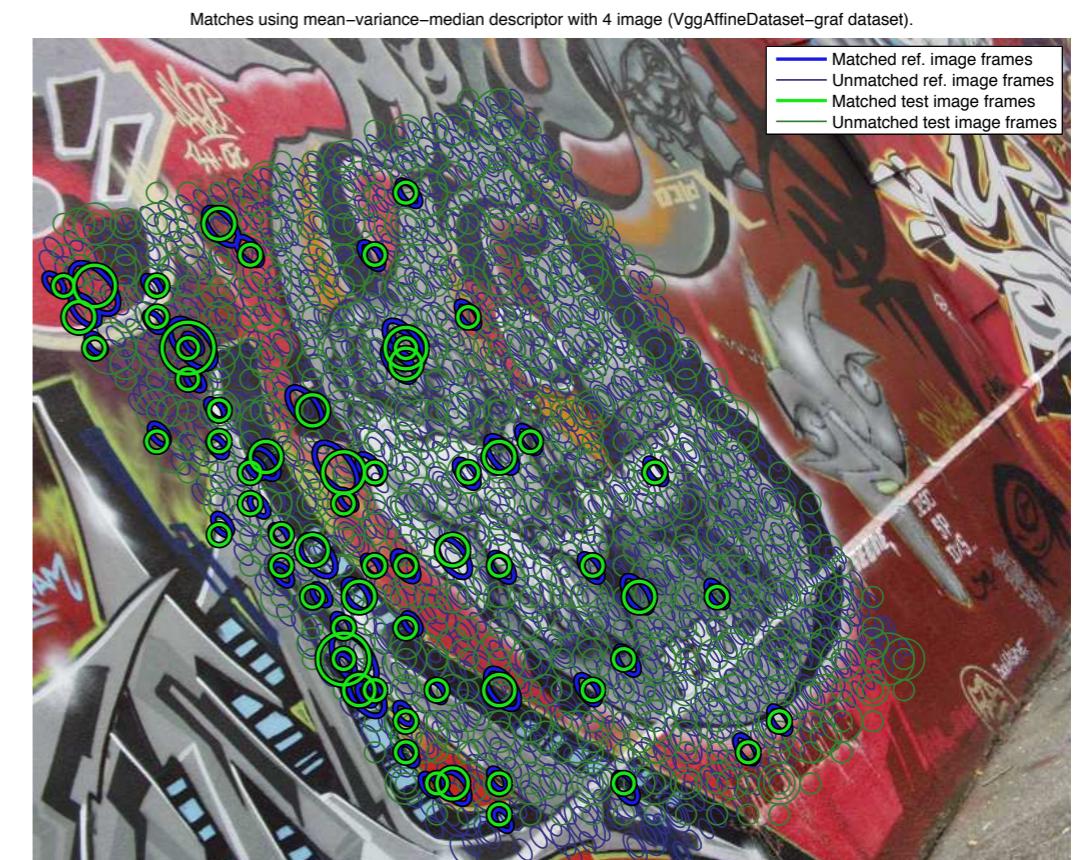
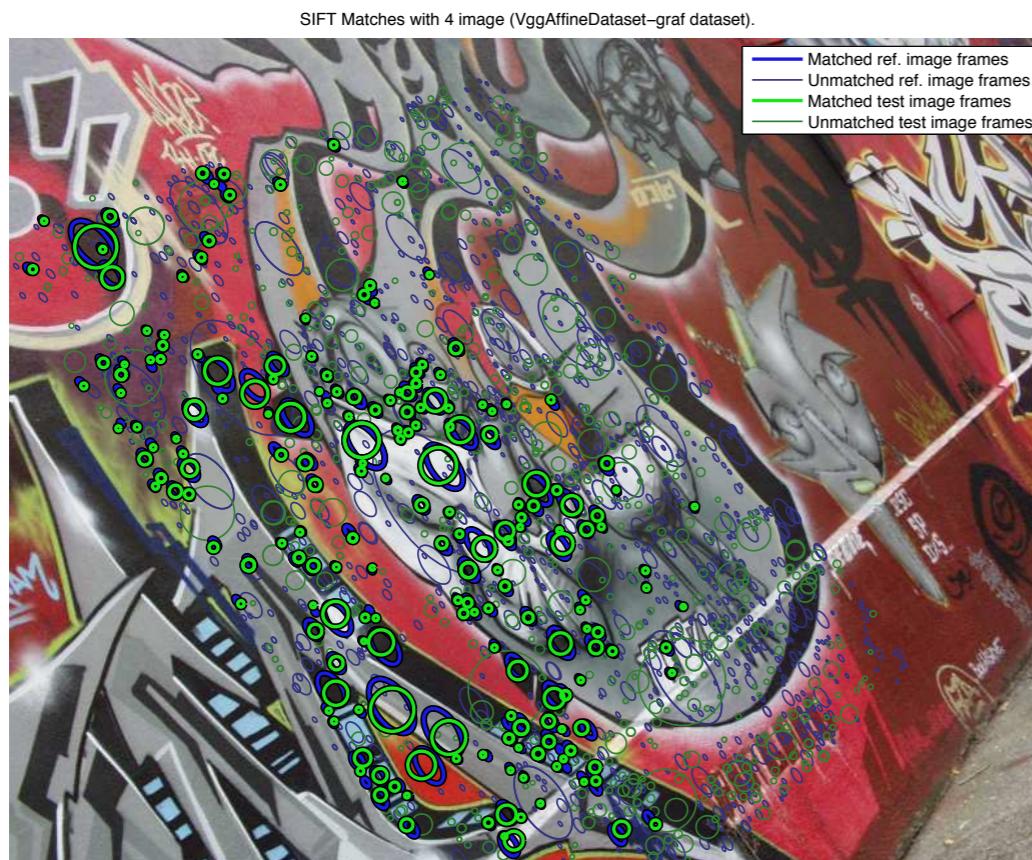
```
detectors{1} = VggAffine('Detector','haraff', 'Threshold ', 500) ;
detectors{2} = VggAffine('Detector','haraff', 'Threshold ', 1000) ;
```

- Visualising matches

```
[~ ~ matches reprojFrames] = benchmark.testFeatureExtractor( ... )
```

...

```
benchmarks.helpers.plotFrameMatches(matches, reprojFrames)
```



Other benchmarks

- **Detector matching score**

```
benchmark = RepeatabilityBenchmark('mode','MatchingScore') ;
```

- **Image retrieval**

- Example: Oxford 5K lite
- mAP evaluation

```
dataset = VggRetrievalDataset('Category','oxbuild',  
                             'BadImagesNum',100);
```

```
benchmark = RetrievalBenchmark() ;
```

```
mAP = benchmark.testFeatureExtractor(detectors{d}, dataset);
```

Summary

<http://www.vlfeat.org/benchmarks/index.html>

- **Benchmarks**
 - Indirect: repeatability and matching score
 - Direct: image retrieval
- **VLBenchmarks**
 - a simple to use MATLAB framework
 - convenient
- **The future**
 - Existing measures have many shortcomings
 - Hopefully better benchmarks will be available soon
 - And they will be added to VL Benchmarks for your convenience

Credits

87



Karel Lenc



Varun Gulshan



Krystian Mikolajczyk

Tinne Tuytelaars

Jiri Matas

Cordelia Schmid

Andrew Zisserman

Thank you for coming!

88

VLFeat

<http://www.vlfeat.org/>

VLBenchmarks

<http://www.vlfeat.org/benchmarks/>